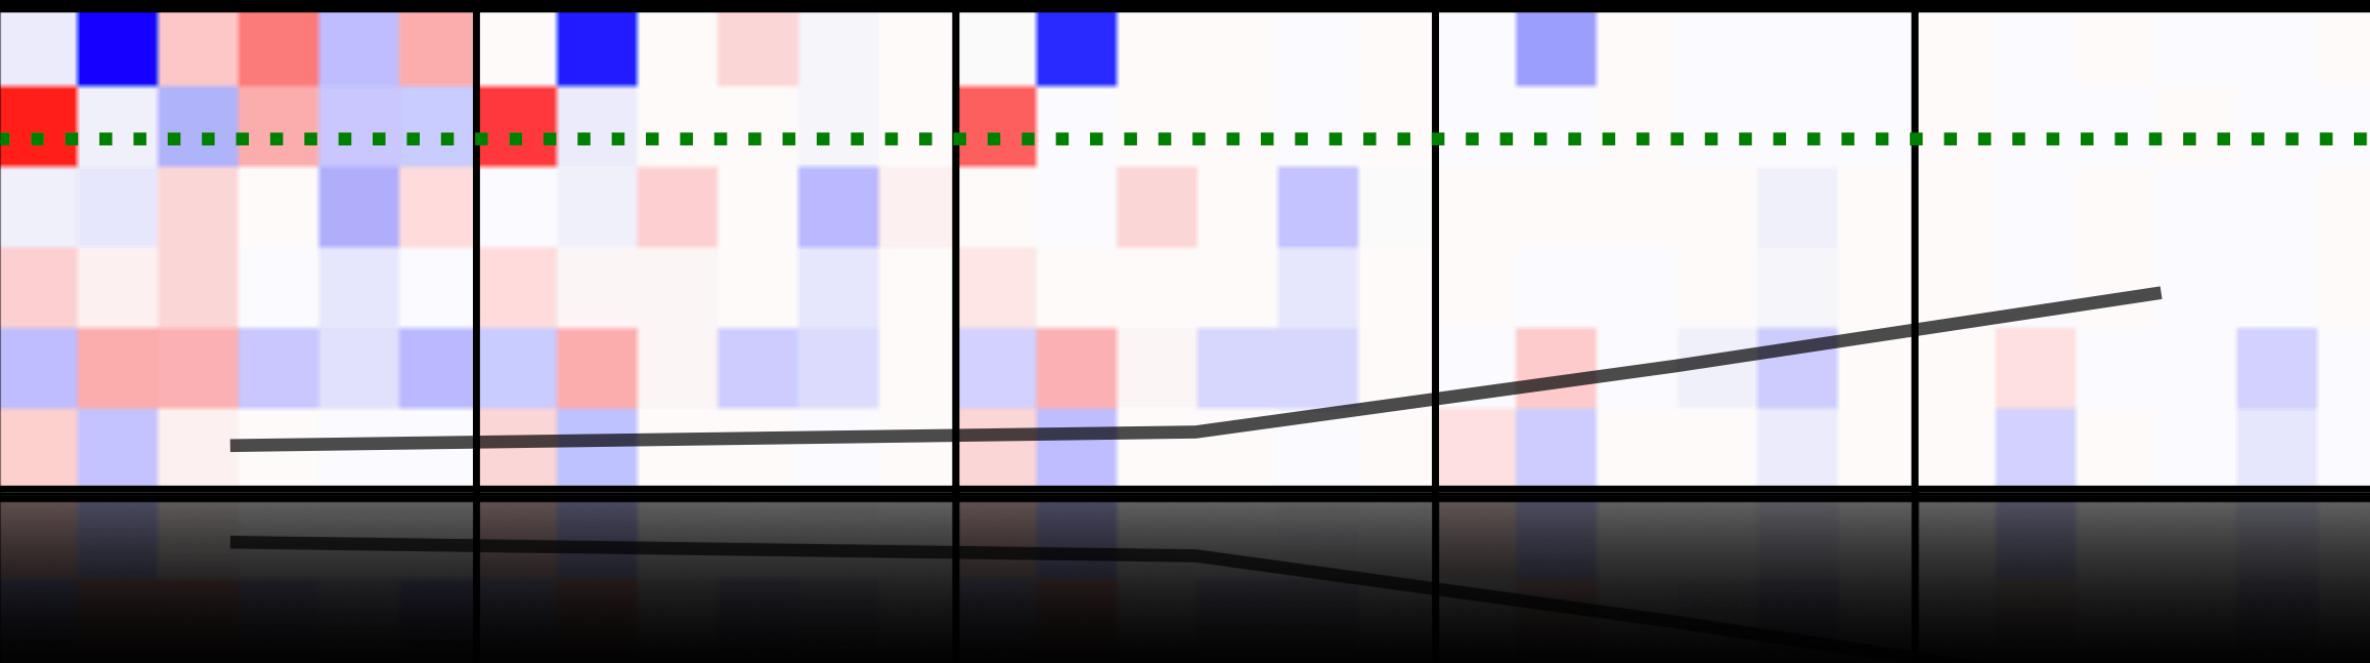


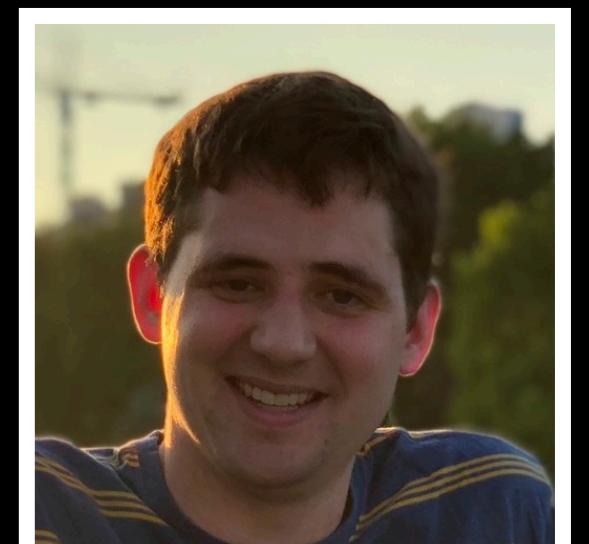
Slides: bit.ly/lcl-slides-2
Preprint: bit.ly/lcl-paper
Code: bit.ly/lcl-code
Data: bit.ly/lcl-data



Learning Interpretable Feature Context Effects in Discrete Choice

Kiran Tomlinson

@ Stanford SOAL



research with Austin R. Benson

1. Choices and context effects

Discrete choices are everywhere



amazon.com

Tervis Stanford Cardinal Logo Tumbler with Emblem and Black Lid 16oz Mug, Clear

Champion NCAA Men's Eco Powerblend Pocketless 1/4 Zip

Top of the World NCAA Mens Adjustable Vintage Team Icon hat

Includes taxes & fees

\$16⁹⁹

✓prime FREE Delivery Sun, Nov 22

\$53⁰⁰

FREE Delivery for Prime members

Expedia

Best Western University Inn
Ithaca

Black Friday / Cyber Monday Deals Now

Free Shuttle Transportation, Grab & Go Breakfast, WiFi & Parking. Pet friendly, Outdoor Pool, Fitness Center. Sanitizing Daily

Breakfast included

3.9/5 Good (999 reviews)

\$63

per night \$71 total Includes taxes & fees

Quality Inn Ithaca - University Area
Ithaca

Black Friday / Cyber Monday Deals Now

Complimentary Breakfast. Free Airport Shuttle, WiFi & parking. Close to Ithaca College & Cornell University. Pets welcome.

Breakfast included

3.6/5 Good (694 reviews)

\$59

per night \$66 total Includes taxes & fees

Member Price available

Hotel Ithaca
Ithaca

Black Friday / Cyber Monday Deals Now

Includes taxes & fees

Member Price available

\$94

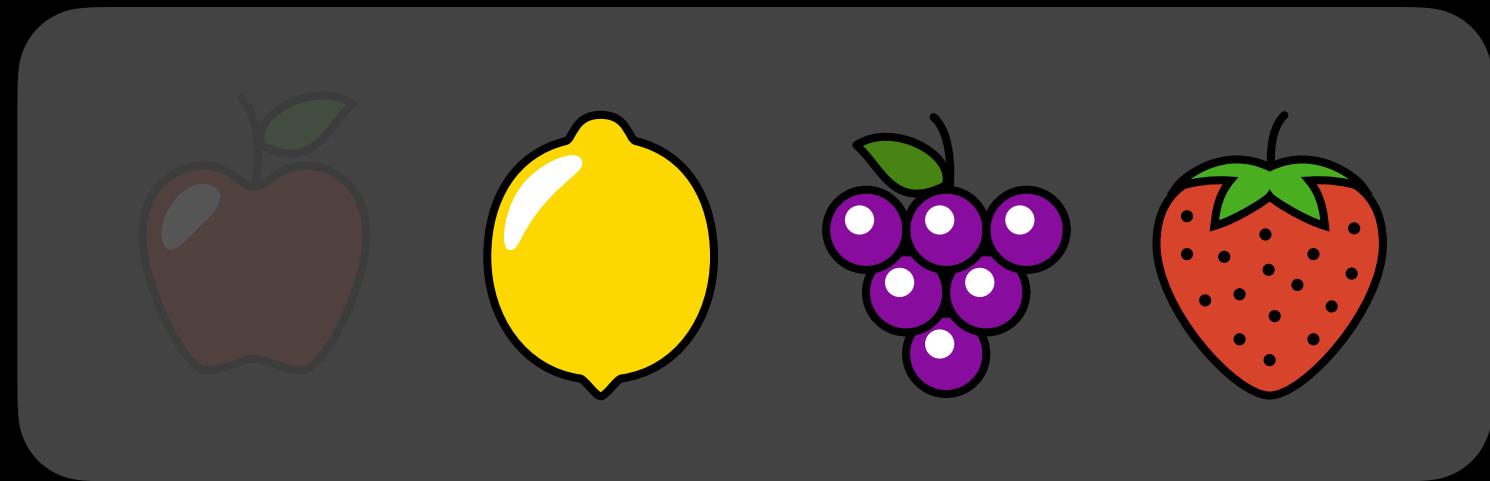
per night \$106 total Includes taxes & fees

4.0/5 Very Good (842 reviews)

“The fundamental problem of discrete choice”

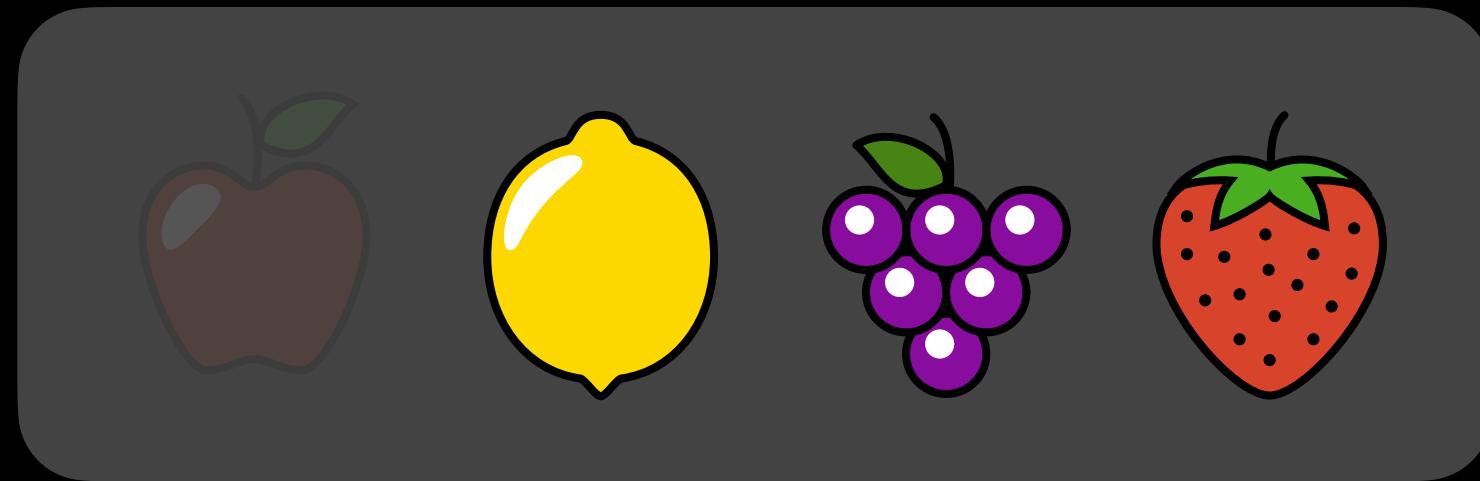
“The fundamental problem of discrete choice”

choice set

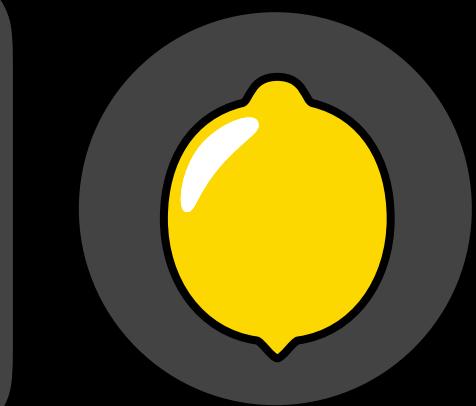


“The fundamental problem of discrete choice”

choice set

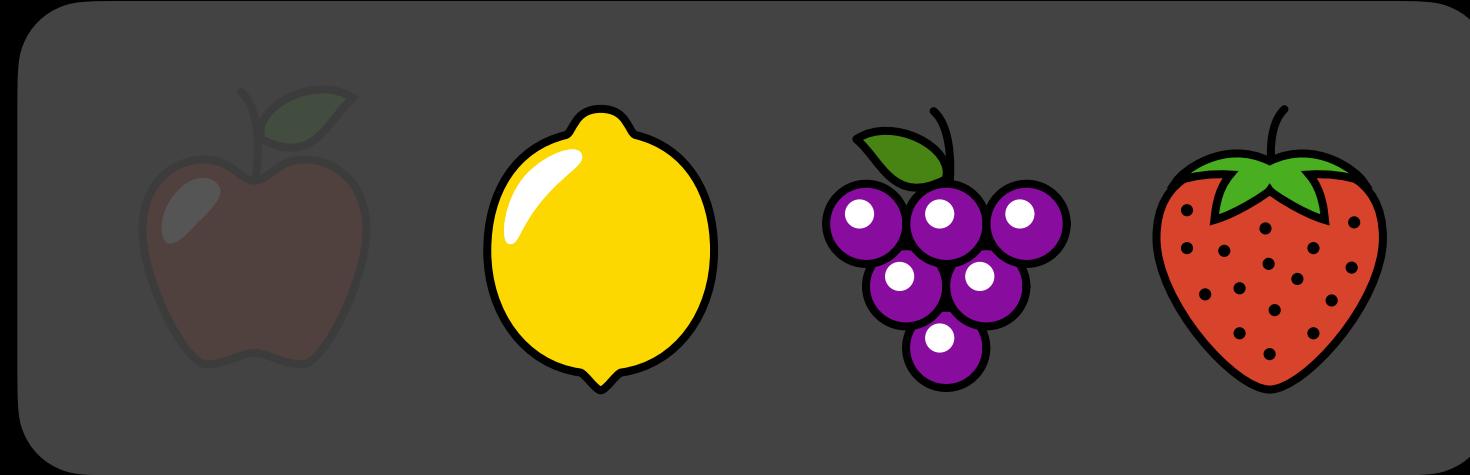


choice

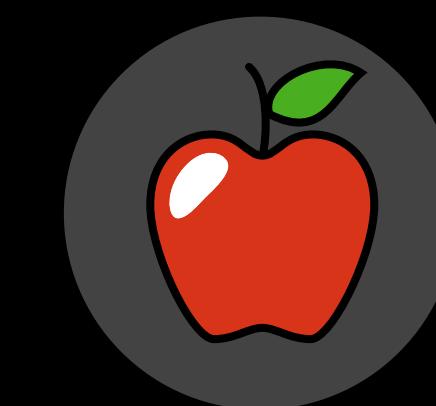
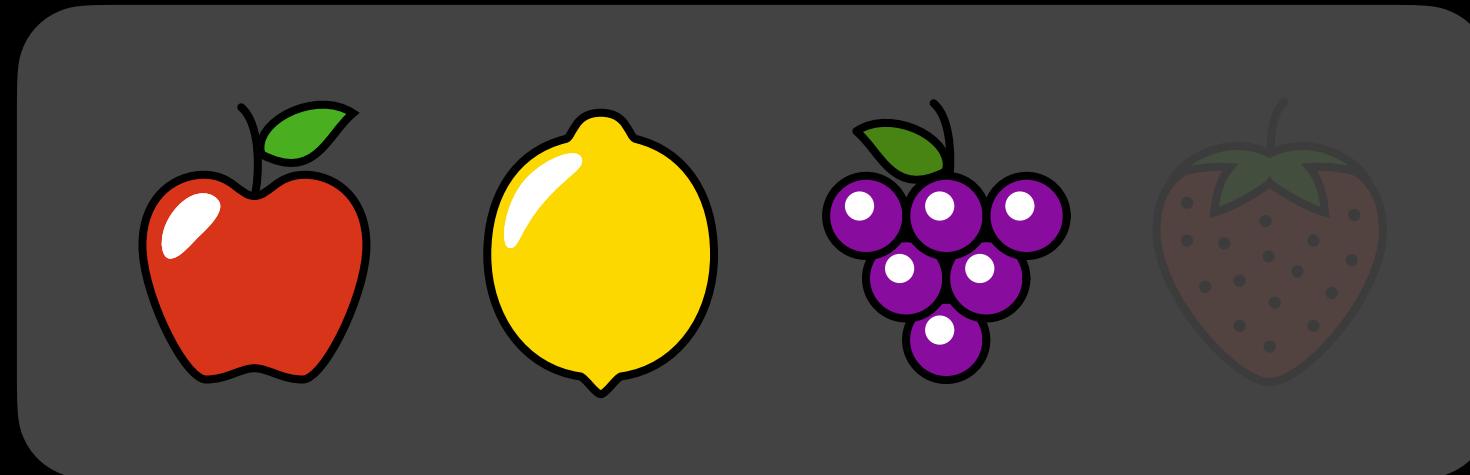
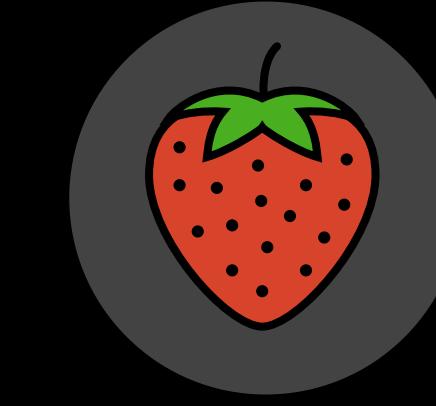
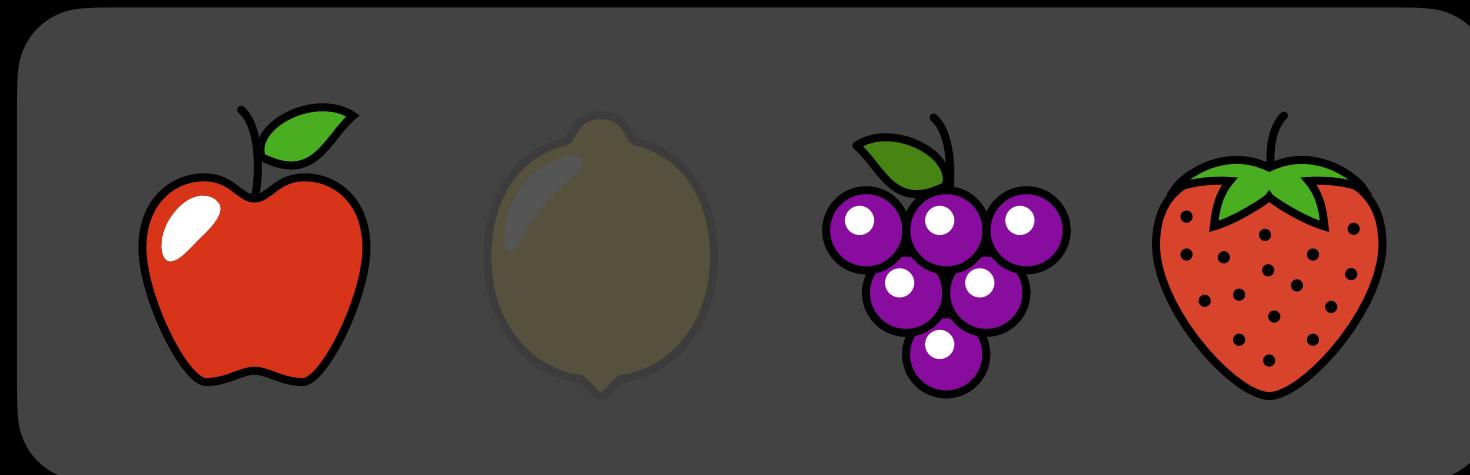
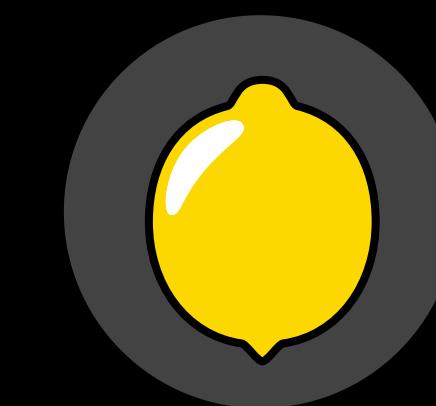


“The fundamental problem of discrete choice”

choice set

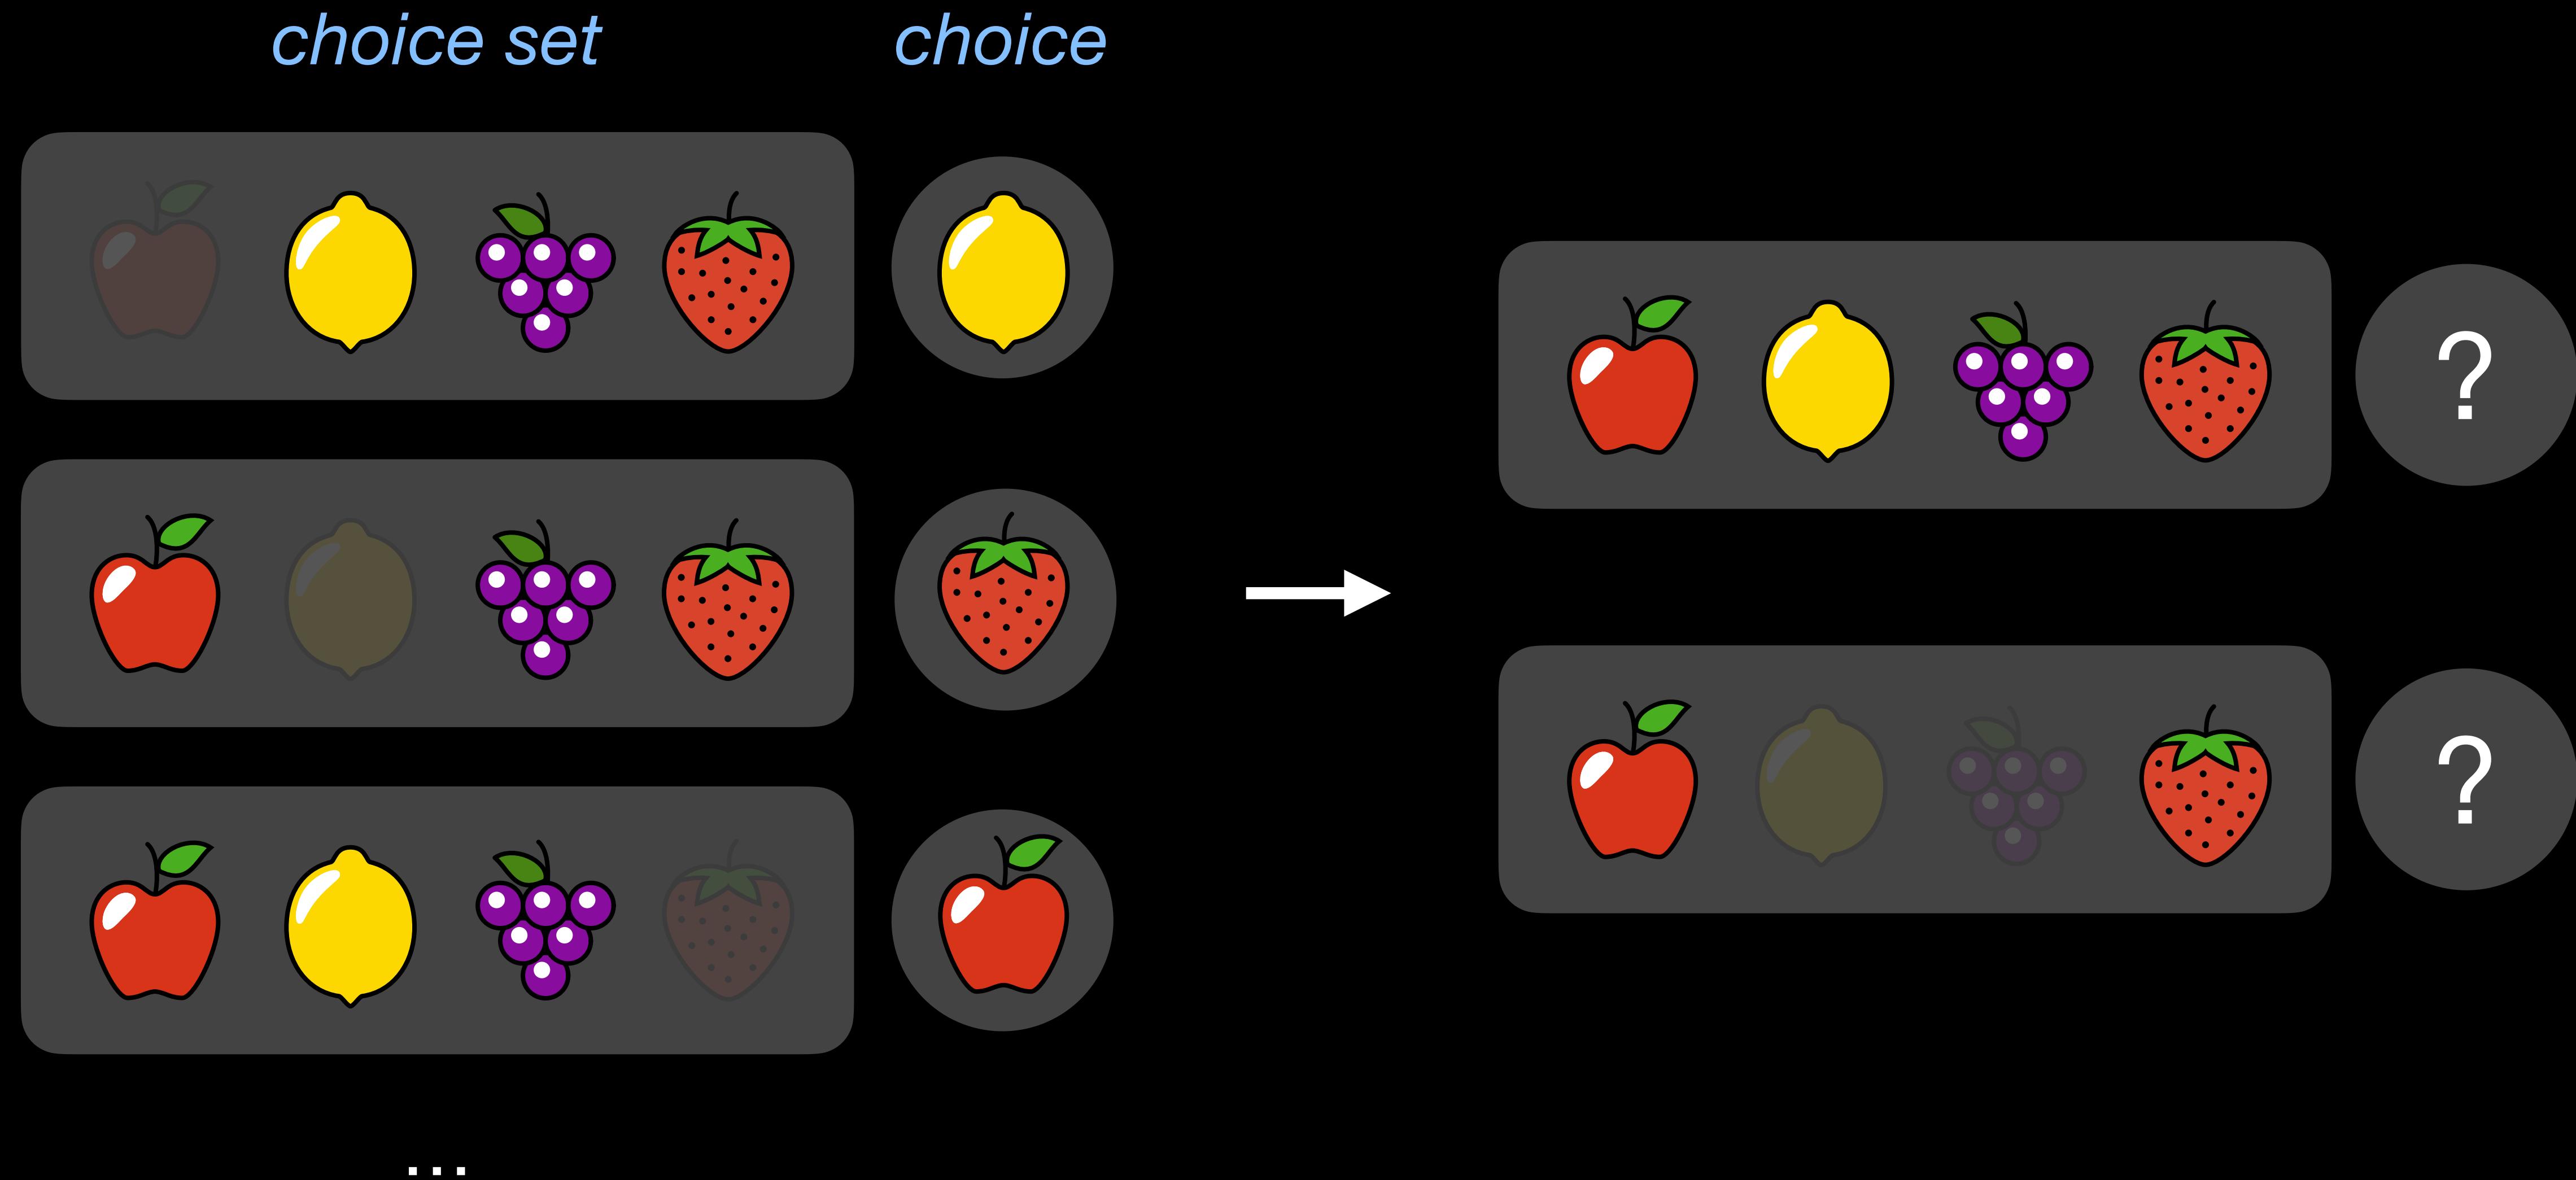


choice



...

“The fundamental problem of discrete choice”



The classic model: *multinomial logit (MNL)*

(McFadden, *Frontiers in Econometrics* 1973)

The classic model: *multinomial logit (MNL)*

(McFadden, *Frontiers in Econometrics* 1973)

Assume *item i* has *utility* u_i

$$\Pr(i \mid C) = \frac{\exp(u_i)}{\sum_{j \in C} \exp(u_j)}$$

The classic model: *multinomial logit (MNL)*

(McFadden, *Frontiers in Econometrics* 1973)

Assume *item i* has *utility u_i*

$$\Pr(i \mid C) = \frac{\exp(u_i)}{\sum_{j \in C} \exp(u_j)}$$

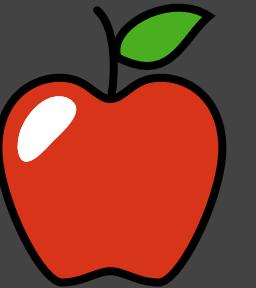
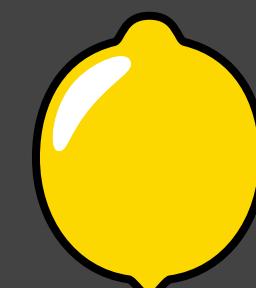
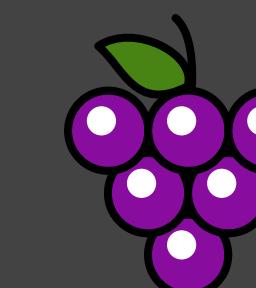
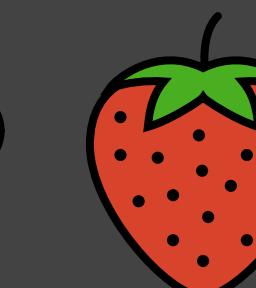


The classic model: *multinomial logit (MNL)*

(McFadden, *Frontiers in Econometrics* 1973)

Assume *item i* has *utility u_i*

$$\Pr(i \mid C) = \frac{\exp(u_i)}{\sum_{j \in C} \exp(u_j)}$$

C				
u_i	1	-1	0	2
$\Pr(i \mid C)$.24	.03	.09	.64

Econometric derivation:

- Draw random utilities $u_i + \epsilon_i$ for each item (ϵ_i i.i.d. Gumbel)
- Rational agent picks $\operatorname{argmax}_{i \in C} u_i + \epsilon_i$

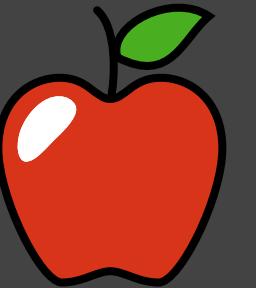
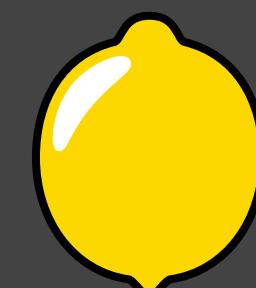
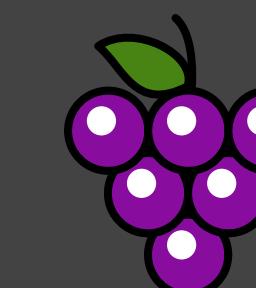
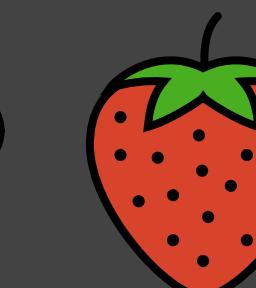
→ MNL

The classic model: *multinomial logit (MNL)*

(McFadden, *Frontiers in Econometrics* 1973)

Assume *item i* has *utility u_i*

$$\Pr(i \mid C) = \frac{\exp(u_i)}{\sum_{j \in C} \exp(u_j)}$$

C					u_i	1	-1	0	2	$\Pr(i \mid C)$.24	.03	.09	.64
-----	---	---	---	---	-------	---	----	---	---	-----------------	-----	-----	-----	-----

Econometric derivation:

- Draw random utilities $u_i + \epsilon_i$ for each item (ϵ_i i.i.d. Gumbel)
- Rational agent picks $\operatorname{argmax}_{i \in C} u_i + \epsilon_i$

→ MNL

Unique choice model satisfying
independence of irrelevant alternatives (IIA):

(Luce, *Individual Choice Behavior* 1959)

$$\frac{\Pr(i \mid C)}{\Pr(j \mid C)} = \frac{\Pr(i \mid C')}{\Pr(j \mid C')}$$

Learning an MNL from choice data

Learning an MNL from choice data

Likelihood of utilities \mathbf{u} given dataset \mathcal{D} :
$$\mathcal{L}(\mathbf{u}; \mathcal{D}) = \prod_{(i,C) \in \mathcal{D}} \frac{\exp(u_i)}{\sum_{j \in C} \exp(u_j)}$$

Learning an MNL from choice data

Likelihood of utilities \mathbf{u} given dataset \mathcal{D} :
$$\mathcal{L}(\mathbf{u}; \mathcal{D}) = \prod_{(i,C) \in \mathcal{D}} \frac{\exp(u_i)}{\sum_{j \in C} \exp(u_j)}$$

Log-likelihood is concave:
$$\ell(\mathbf{u}; \mathcal{D}) = \sum_{(i,C) \in \mathcal{D}} u_i - \log \sum_{j \in C} \exp(u_j)$$

Learning an MNL from choice data

Likelihood of utilities \mathbf{u} given dataset \mathcal{D} :
$$\mathcal{L}(\mathbf{u}; \mathcal{D}) = \prod_{(i,C) \in \mathcal{D}} \frac{\exp(u_i)}{\sum_{j \in C} \exp(u_j)}$$

Log-likelihood is concave:
$$\ell(\mathbf{u}; \mathcal{D}) = \sum_{(i,C) \in \mathcal{D}} u_i - \log \sum_{j \in C} \exp(u_j)$$

→ gradient descent to learn \mathbf{u}
(SGD, Adam, ...)

Problem for MNL: *context effects*

Problem for MNL: *context effects*

The choice set influences preferences.

Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)

Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)



Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)

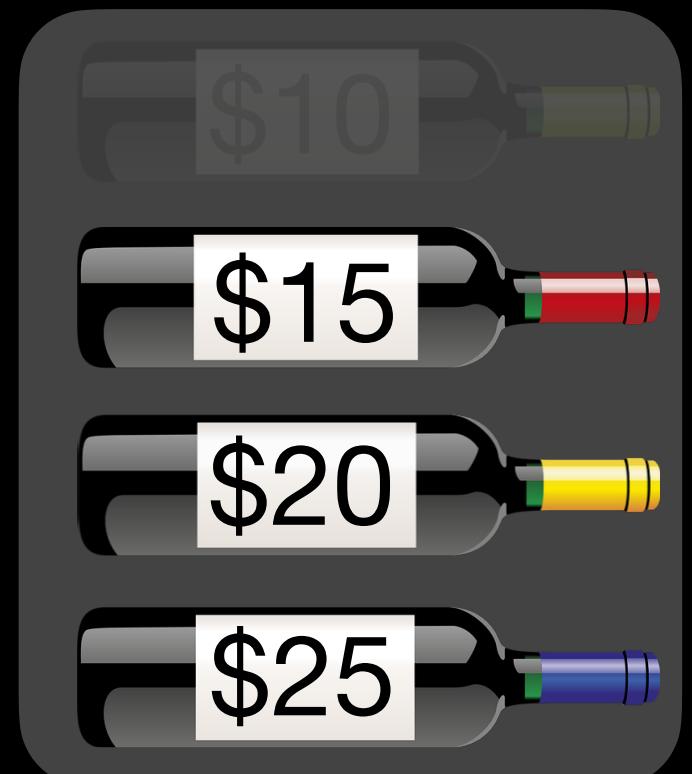


Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)



Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)



Problem for MNL: *context effects*

The choice set influences preferences.

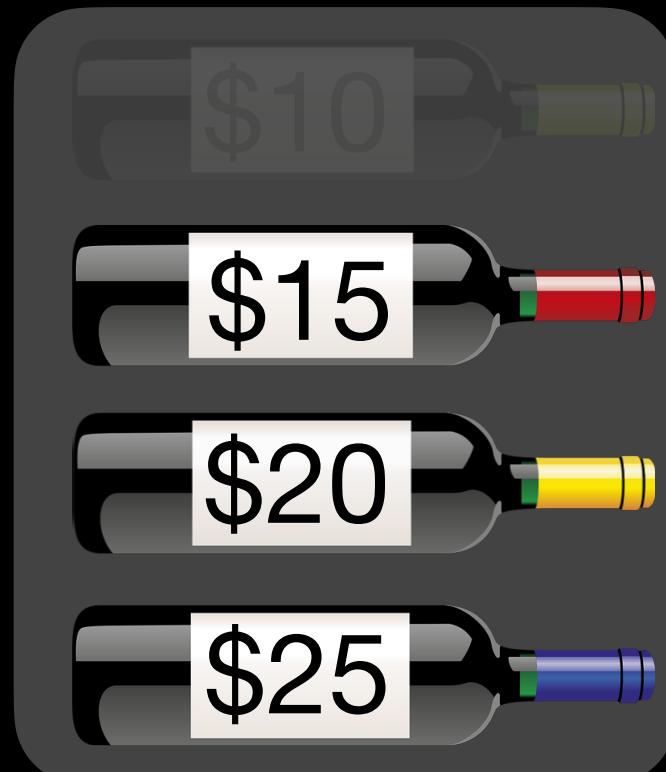
Compromise

(Simonson, 1989)



Asymmetric dominance

(Huber et al., 1982)



Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)



Asymmetric dominance

(Huber et al., 1982)

Economist.com

SUBSCRIPTIONS

Welcome to
The Economist Subscription Centre

Pick the type of subscription you want to buy or renew.

Economist.com subscription - US \$59.00
One-year subscription to Economist.com.
Includes online access to all articles from *The Economist* since 1997.

Print subscription - US \$125.00
One-year subscription to the print edition of *The Economist*.

Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist* since 1997.



Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)



Asymmetric dominance

(Huber et al., 1982)

Economist.com **SUBSCRIPTIONS**
Welcome to The Economist Subscription Centre
Pick the type of subscription you want to buy or renew.
 Economist.com subscription - US \$59.00
One-year subscription to Economist.com.
Includes online access to all articles from *The Economist* since 1997.
 Print subscription - US \$125.00
One-year subscription to the print edition of *The Economist*.
 Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist* since 1997.



Economist.com **Welcome to The Economist Subscription Centre**
Pick the type of subscription you want to buy or renew.
 Economist.com subscription - US \$59.00
One-year subscription to Economist.com.
Includes online access to all articles from *The Economist* since 1997. **16%** → **68%**
 Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist* since 1997. **84%** → **32%**

(Ariely, 2008)

Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)



Asymmetric dominance

(Huber et al., 1982)

Economist.com **SUBSCRIPTIONS**
Welcome to The Economist Subscription Centre
Pick the type of subscription you want to buy or renew.
 Economist.com subscription - US \$59.00
One-year subscription to Economist.com.
Includes online access to all articles from *The Economist* since 1997.
 Print subscription - US \$125.00
One-year subscription to the print edition of *The Economist*.
 Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist* since 1997.



Similarity

(Tversky, 1972)



(Ariely, 2008)

Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)



Asymmetric dominance

(Huber et al., 1982)

Economist.com **SUBSCRIPTIONS**
Welcome to The Economist Subscription Centre
Pick the type of subscription you want to buy or renew.
 Economist.com subscription - US \$59.00
One-year subscription to Economist.com. Includes online access to all articles from *The Economist* since 1997.
 Print subscription - US \$125.00
One-year subscription to the print edition of *The Economist*.
 Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist* since 1997.



Economist.com **Welcome to The Economist Subscription Centre**
Pick the type of subscription you want to buy or renew.
 Economist.com subscription - US \$59.00
One-year subscription to Economist.com. Includes online access to all articles from *The Economist* since 1997. **16%** → **68%**
 Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist* since 1997. **84%** → **32%**

Similarity

(Tversky, 1972)

	David Perdue*	✓	Rep.	2,462,617	49.7%
	Jon Ossoff	✓	Dem.	2,374,519	47.9

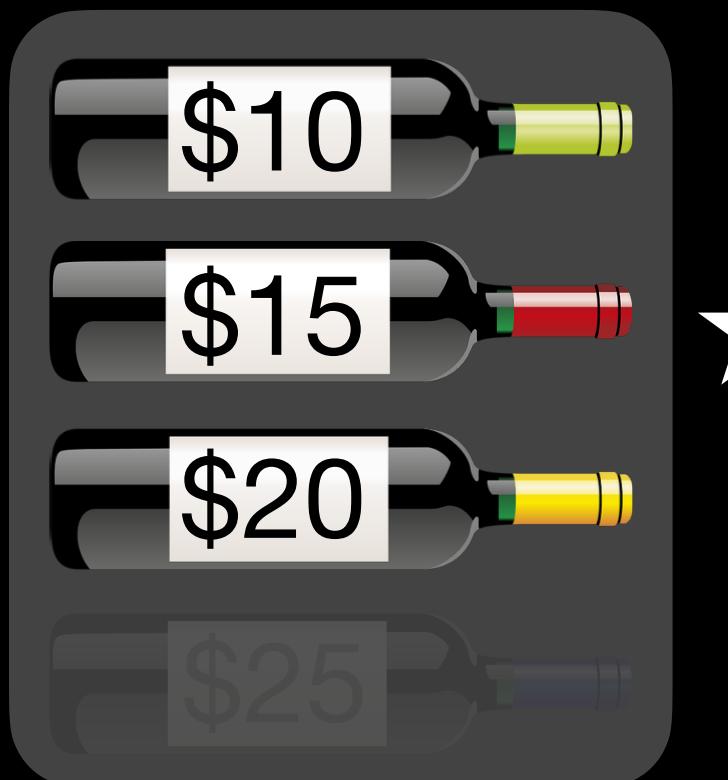
(Ariely, 2008)

Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)



Asymmetric dominance

(Huber et al., 1982)

Economist.com

SUBSCRIPTIONS

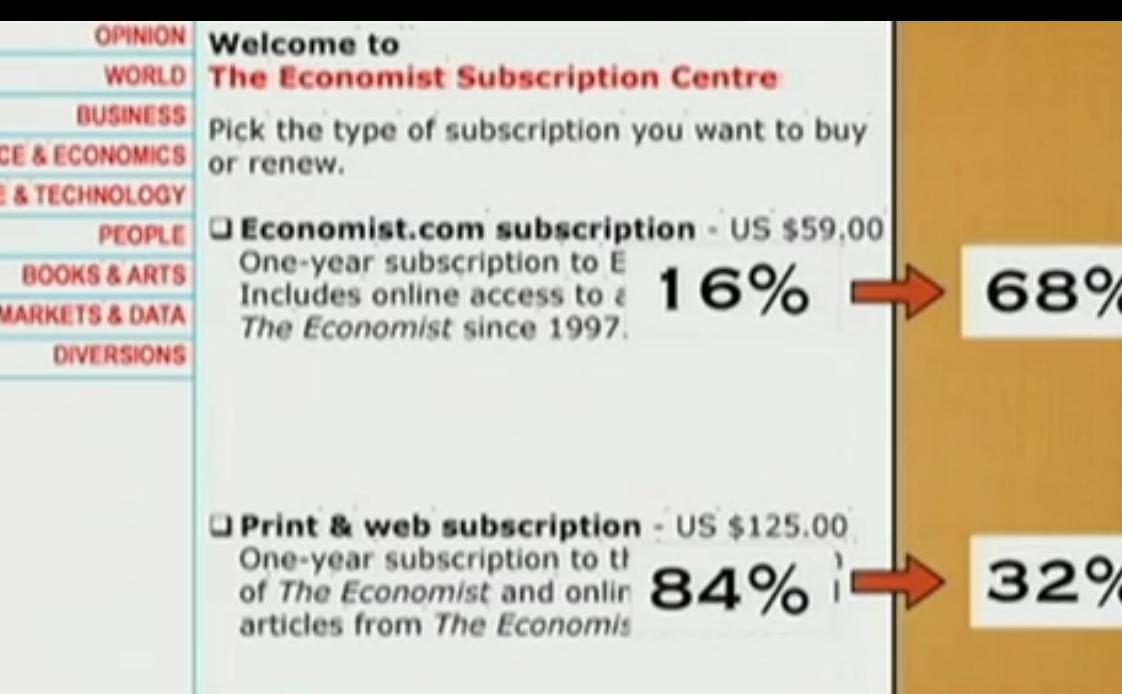
Welcome to The Economist Subscription Centre

Pick the type of subscription you want to buy or renew.

Economist.com subscription - US \$59.00
One-year subscription to Economist.com. Includes online access to all articles from *The Economist* since 1997.

Print subscription - US \$125.00
One-year subscription to the print edition of *The Economist*.

Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist* since 1997.



Similarity

(Tversky, 1972)

	David Perdue*	✓	Rep.	2,462,617	49.7%
	Jon Ossoff	✓	Dem.	2,374,519	47.9
	Raphael Warnock	✓	Dem.	1,617,035	32.9%
	Kelly Loeffler*	✓	Rep.	1,273,214	25.9
	Doug Collins		Rep.	980,454	20.0

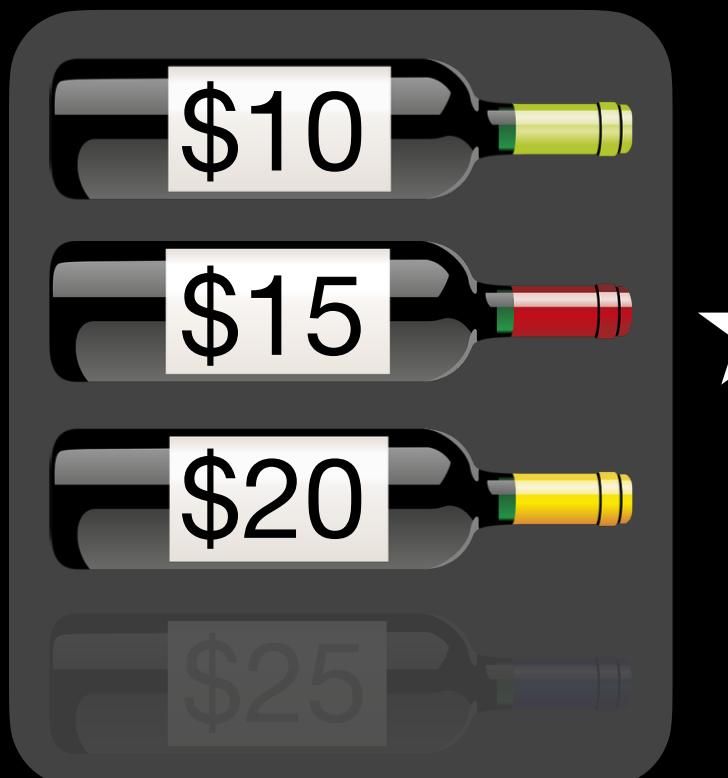
(Ariely, 2008)

Problem for MNL: *context effects*

The choice set influences preferences.

Compromise

(Simonson, 1989)



Asymmetric dominance

(Huber et al., 1982)

Economist.com

SUBSCRIPTIONS

Welcome to The Economist Subscription Centre

Pick the type of subscription you want to buy or renew.

Economist.com subscription - US \$59.00
One-year subscription to Economist.com. Includes online access to all articles from *The Economist* since 1997.

Print subscription - US \$125.00
One-year subscription to the print edition of *The Economist*.

Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist* since 1997.



Welcome to The Economist Subscription Centre

Pick the type of subscription you want to buy or renew.

Economist.com subscription - US \$59.00
One-year subscription to Economist.com. Includes online access to all articles from *The Economist* since 1997. **16%** → **68%**

Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist*. **84%** → **32%**

(Ariely, 2008)

Similarity

(Tversky, 1972)

	David Perdue*	✓	Rep.	2,462,617	49.7%
	Jon Ossoff	✓	Dem.	2,374,519	47.9
	Raphael Warnock	✓	Dem.	1,617,035	32.9%
	Kelly Loeffler*	✓	Rep.	1,273,214	25.9
	Doug Collins		Rep.	980,454	20.0

IIA violations:

$$\frac{\Pr(i \mid C)}{\Pr(j \mid C)} \neq \frac{\Pr(i \mid C')}{\Pr(j \mid C')}$$

Natural context effect model: *CDM*

(Seshadri, Peysakhovich, & Ugander, ICML 2019)

Natural context effect model: *CDM*

(Seshadri, Peysakhovich, & Ugander, ICML 2019)

Item j exerts *pull* u_{ij} on item i , item utility is sum of pulls:

$$\Pr(i \mid C) = \frac{\exp\left(\sum_{k \in C \setminus i} u_{ik}\right)}{\sum_{j \in C} \exp\left(\sum_{k \in C \setminus i} u_{jk}\right)}$$

Natural context effect model: *CDM*

(Seshadri, Peysakhovich, & Ugander, ICML 2019)

Item j exerts *pull* u_{ij} on item i , item utility is sum of pulls:

$$\Pr(i \mid C) = \frac{\exp\left(\sum_{k \in C \setminus i} u_{ik}\right)}{\sum_{j \in C} \exp\left(\sum_{k \in C \setminus i} u_{jk}\right)}$$

	Raphael Warnock ✓	Dem. 1,617,035 32.9%
	Kelly Loeffler* ✓	Rep. 1,273,214 25.9
	Doug Collins	Rep. 980,454 20.0

$$u_{\text{Loeffler, Collins}} < 0$$
$$u_{\text{Collins, Loeffler}} < 0$$

Natural context effect model: *CDM*

(Seshadri, Peysakhovich, & Ugander, ICML 2019)

Item j exerts *pull* u_{ij} on item i , item utility is sum of pulls:

$$\Pr(i \mid C) = \frac{\exp\left(\sum_{k \in C \setminus i} u_{ik}\right)}{\sum_{j \in C} \exp\left(\sum_{k \in C \setminus i} u_{jk}\right)}$$

	Raphael Warnock ✓	Dem. 1,617,035 32.9%
	Kelly Loeffler* ✓	Rep. 1,273,214 25.9
	Doug Collins	Rep. 980,454 20.0

$$u_{\text{Loeffler, Collins}} < 0$$
$$u_{\text{Collins, Loeffler}} < 0$$

Economist.com

SUBSCRIPTIONS

Welcome to
The Economist Subscription Centre

Pick the type of subscription you want to buy or renew.

Economist.com subscription - US \$59.00
One-year subscription to Economist.com.
Includes online access to all articles from *The Economist* since 1997.

Print subscription - US \$125.00
One-year subscription to the print edition of *The Economist*.

Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist* since 1997.

$$u_{\text{Print\&Web, Print}} > 0$$

Natural context effect model: *CDM*

(Seshadri, Peysakhovich, & Ugander, ICML 2019)

Item j exerts *pull* u_{ij} on item i , item utility is sum of pulls:

$$\Pr(i \mid C) = \frac{\exp\left(\sum_{k \in C \setminus i} u_{ik}\right)}{\sum_{j \in C} \exp\left(\sum_{k \in C \setminus i} u_{jk}\right)}$$

	Raphael Warnock ✓	Dem. 1,617,035 32.9%
	Kelly Loeffler* ✓	Rep. 1,273,214 25.9
	Doug Collins	Rep. 980,454 20.0

$$u_{\text{Loeffler, Collins}} < 0$$
$$u_{\text{Collins, Loeffler}} < 0$$

Economist.com

SUBSCRIPTIONS

Welcome to
The Economist Subscription Centre

Pick the type of subscription you want to buy or renew.

Economist.com subscription - US \$59.00
One-year subscription to Economist.com.
Includes online access to all articles from *The Economist* since 1997.

Print subscription - US \$125.00
One-year subscription to the print edition of *The Economist*.

Print & web subscription - US \$125.00
One-year subscription to the print edition of *The Economist* and online access to all articles from *The Economist* since 1997.

$$u_{\text{Print\&Web, Print}} > 0$$

Assumes no higher-order effects
(i.e., context effects decompose additively into effects of items)

2. Item features and the LCL

Choice models with *item features*

Choice models with *item features*

So far, models have per-item parameters

Choice models with *item features*

So far, models have per-item parameters

→ can't generalize to new items not in training set

Choice models with *item features*

So far, models have per-item parameters

- can't generalize to new items not in training set
- hard to learn utilities for rare items

Choice models with *item features*

So far, models have per-item parameters

- can't generalize to new items not in training set
- hard to learn utilities for rare items
- too many parameters with many items

Choice models with *item features*

So far, models have per-item parameters

- can't generalize to new items not in training set
- hard to learn utilities for rare items
- too many parameters with many items

Use item features:



genre: drama,
in_top_10: True,
has_new_episodes: True,
producer: Netflix



genre: comedy,
in_top_10: False,
has_new_episodes: False,
producer: NBC



genre: drama,
in_top_10: True,
has_new_episodes: False,
producer: Netflix



genre: reality,
in_top_10: True,
has_new_episodes: False,
producer: Banijay

MNL with item features: *conditional logit*

MNL with item features: *conditional logit*

Feature vector $x_i \in \mathbb{R}^d$ for each item i

Preference vector $\theta \in \mathbb{R}^d$

MNL with item features: *conditional logit*

Feature vector $x_i \in \mathbb{R}^d$ for each item i

Preference vector $\theta \in \mathbb{R}^d$

MNL:

$$\Pr(i \mid C) = \frac{\exp(u_i)}{\sum_{j \in C} \exp(u_j)}$$

MNL with item features: *conditional logit*

Feature vector $x_i \in \mathbb{R}^d$ for each item i

Preference vector $\theta \in \mathbb{R}^d$

MNL:

$$\Pr(i \mid C) = \frac{\exp(u_i)}{\sum_{j \in C} \exp(u_j)} \rightarrow$$

Conditional logit:

$$\Pr(i \mid C) = \frac{\exp(\theta^T x_i)}{\sum_{j \in C} \exp(\theta^T x_j)}$$

MNL with item features: *conditional logit*

Feature vector $x_i \in \mathbb{R}^d$ for each item i

Preference vector $\theta \in \mathbb{R}^d$

MNL:

$$\Pr(i | C) = \frac{\exp(u_i)}{\sum_{j \in C} \exp(u_j)} \rightarrow$$

Conditional logit:

$$\Pr(i | C) = \frac{\exp(\theta^T x_i)}{\sum_{j \in C} \exp(\theta^T x_j)}$$

Preference coefficient θ_k is easy to interpret: importance of the k^{th} feature

Incorporating *feature context effects* into conditional logit

Incorporating *feature context effects* into conditional logit

Conditional logit utility: $u_i = \theta^T x_i$

Incorporating *feature context effects* into conditional logit

Conditional logit utility: $u_i = \theta^T x_i$ → *Contextual utility:* $u_{i,C} = [\theta + F(C)]^T x_i$

Incorporating *feature context effects* into conditional logit

Conditional logit utility: $u_i = \theta^T x_i$ \rightarrow Contextual utility: $u_{i,C} = [\theta + F(C)]^T x_i$

Simplifying assumptions on $F(C)$:

Incorporating *feature context effects* into conditional logit

Conditional logit utility: $u_i = \theta^T x_i$ \rightarrow Contextual utility: $u_{i,C} = [\theta + F(C)]^T x_i$

Simplifying assumptions on $F(C)$:

1. *Additivity*: $F(C) \propto \sum_{j \in C} f(x_j)$ for some function f

Incorporating *feature context effects* into conditional logit

Conditional logit utility: $u_i = \theta^T x_i \rightarrow$ *Contextual utility:* $u_{i,C} = [\theta + F(C)]^T x_i$

Simplifying assumptions on $F(C)$:

1. *Additivity:* $F(C) \propto \sum_{j \in C} f(x_j)$ for some function f

2. *Dilution:* $F(C) = \frac{1}{|C|} \sum_{j \in C} f(x_j)$

Incorporating *feature context effects* into conditional logit

Conditional logit utility: $u_i = \theta^T x_i \rightarrow$ Contextual utility: $u_{i,C} = [\theta + F(C)]^T x_i$

Simplifying assumptions on $F(C)$:

1. *Additivity*: $F(C) \propto \sum_{j \in C} f(x_j)$ for some function f
2. *Dilution*: $F(C) = \frac{1}{|C|} \sum_{j \in C} f(x_j)$
3. *Linearity*: $f(x_j) = Ax_j$ for some matrix $A \in \mathbb{R}^{d \times d}$

Incorporating *feature context effects* into conditional logit

Conditional logit utility: $u_i = \theta^T x_i \rightarrow$ Contextual utility: $u_{i,C} = [\theta + F(C)]^T x_i$

Simplifying assumptions on $F(C)$:

1. *Additivity*: $F(C) \propto \sum_{j \in C} f(x_j)$ for some function f

2. *Dilution*: $F(C) = \frac{1}{|C|} \sum_{j \in C} f(x_j)$

3. *Linearity*: $f(x_j) = Ax_j$ for some matrix $A \in \mathbb{R}^{d \times d}$

$\rightarrow u_{i,C} = (\theta + Ax_C)^T x_i$ ($x_C = \frac{1}{|C|} \sum_{j \in C} x_j$ is the *mean feature vector*)

The *Linear Context Logit (LCL)*

The *Linear Context Logit (LCL)*

$$\Pr(i \mid C) = \frac{\exp([\theta + Ax_C]^T x_i)}{\sum_{j \in C} \exp([\theta + Ax_C]^T x_j)}$$

The *Linear Context Logit (LCL)*

$$\Pr(i \mid C) = \frac{\exp([\theta + Ax_C]^T x_i)}{\sum_{j \in C} \exp([\theta + Ax_C]^T x_j)}$$

→ convex negative log-likelihood

The *Linear Context Logit (LCL)*

$$\Pr(i \mid C) = \frac{\exp([\theta + Ax_C]^T x_i)}{\sum_{j \in C} \exp([\theta + Ax_C]^T x_j)}$$

- convex negative log-likelihood
- θ : *base preference coefficients*

The *Linear Context Logit (LCL)*

$$\Pr(i \mid C) = \frac{\exp([\theta + Ax_C]^T x_i)}{\sum_{j \in C} \exp([\theta + Ax_C]^T x_j)}$$

- convex negative log-likelihood
- θ : *base preference coefficients*
- $A_{pq} > 0$: when q is *high* in the choice set, p is *more preferred*

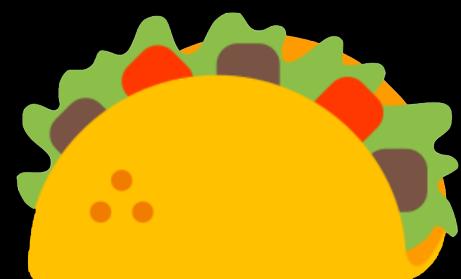
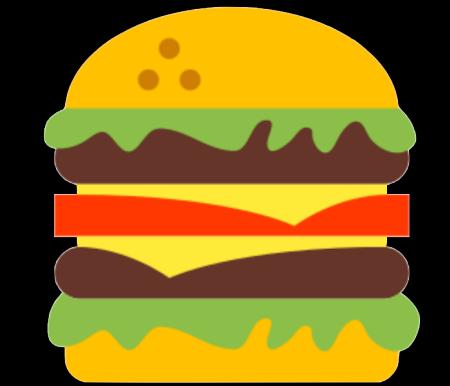
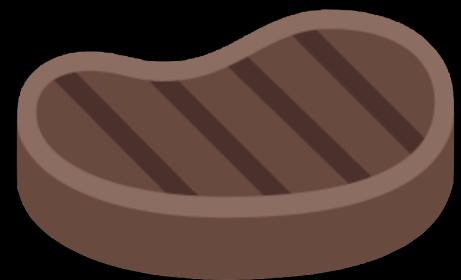
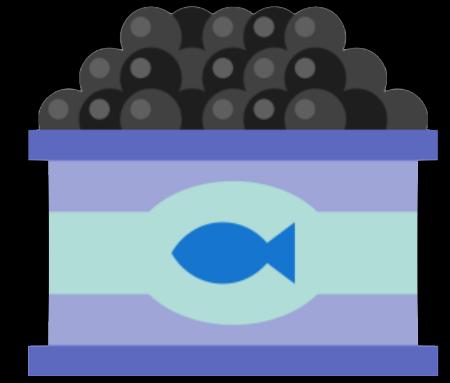
The *Linear Context Logit (LCL)*

$$\Pr(i \mid C) = \frac{\exp([\theta + Ax_C]^T x_i)}{\sum_{j \in C} \exp([\theta + Ax_C]^T x_j)}$$

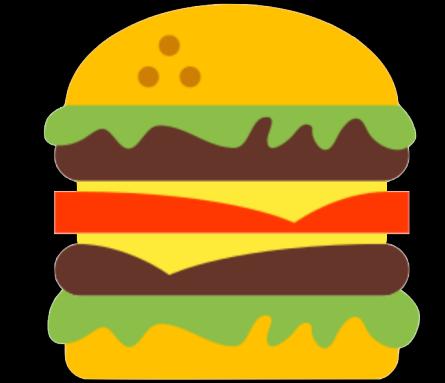
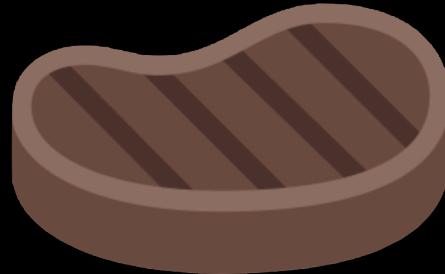
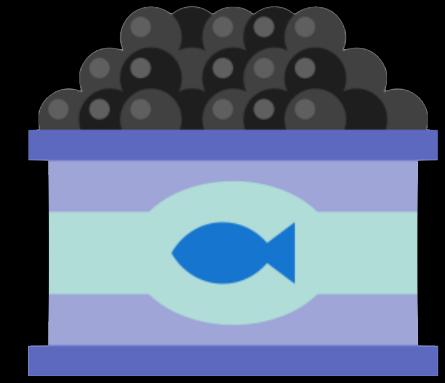
- convex negative log-likelihood
- θ : *base preference coefficients*
- $A_{pq} > 0$: when q is *high* in the choice set, p is *more* preferred
- $A_{pq} < 0$: when q is *high* in the choice set, p is *less* preferred

LCL example: restaurant selection

LCL example: restaurant selection



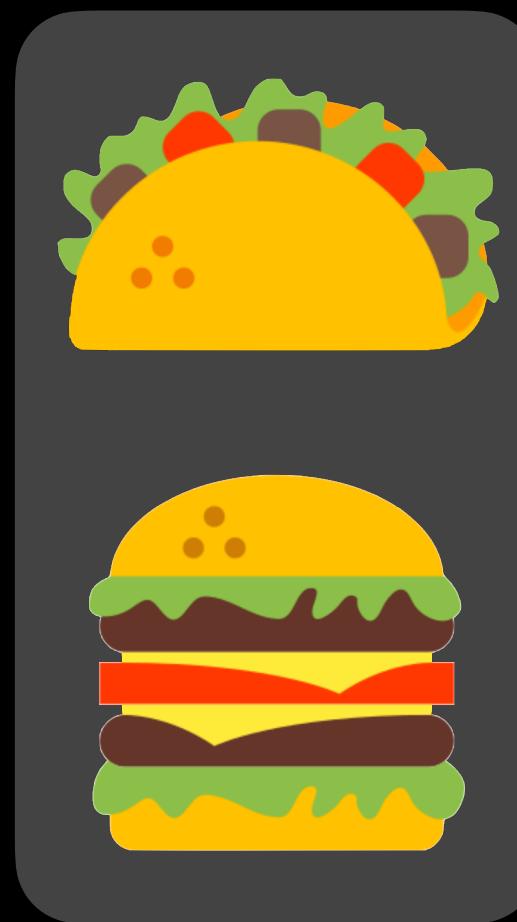
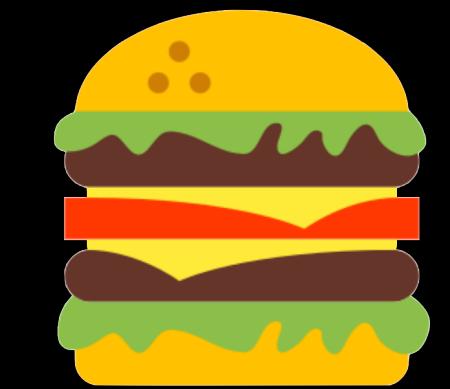
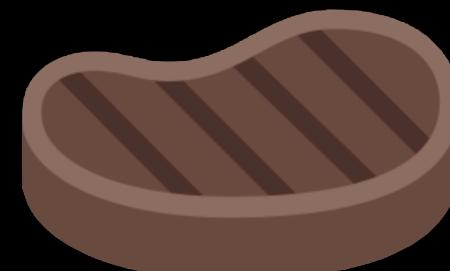
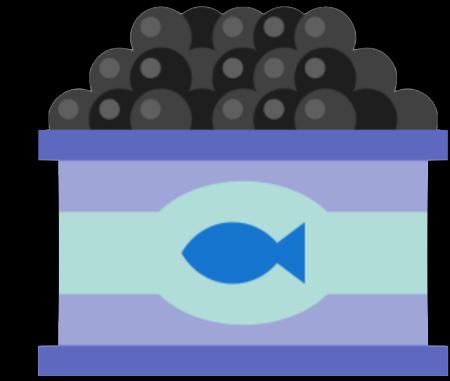
LCL example: restaurant selection



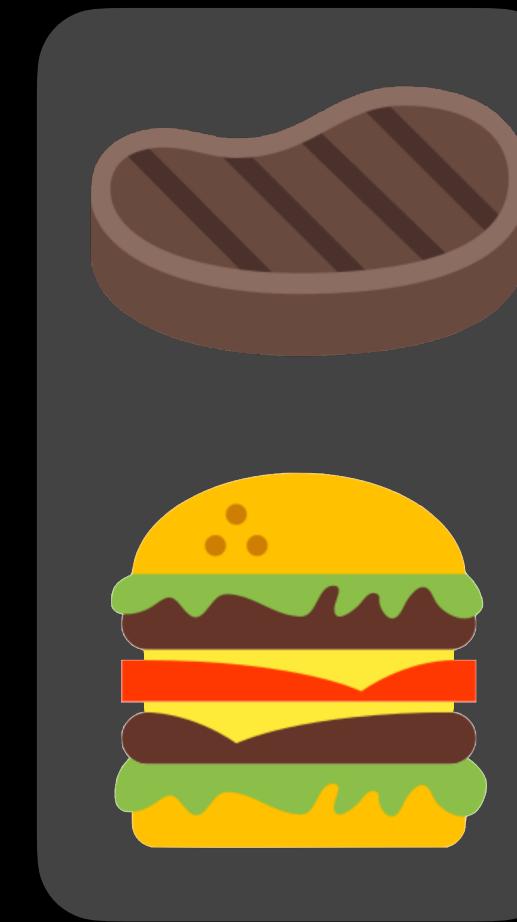
item features:

- price
- service speed
- wine selection

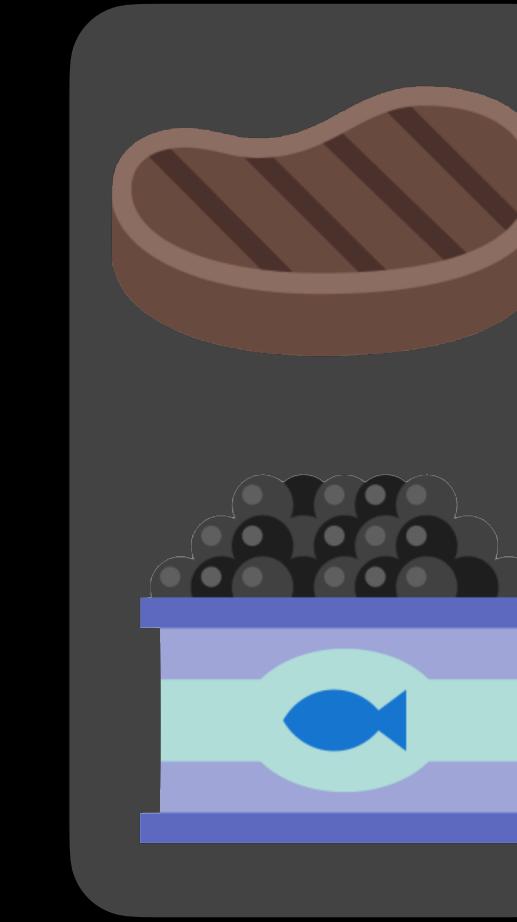
LCL example: restaurant selection



C_1



C_2

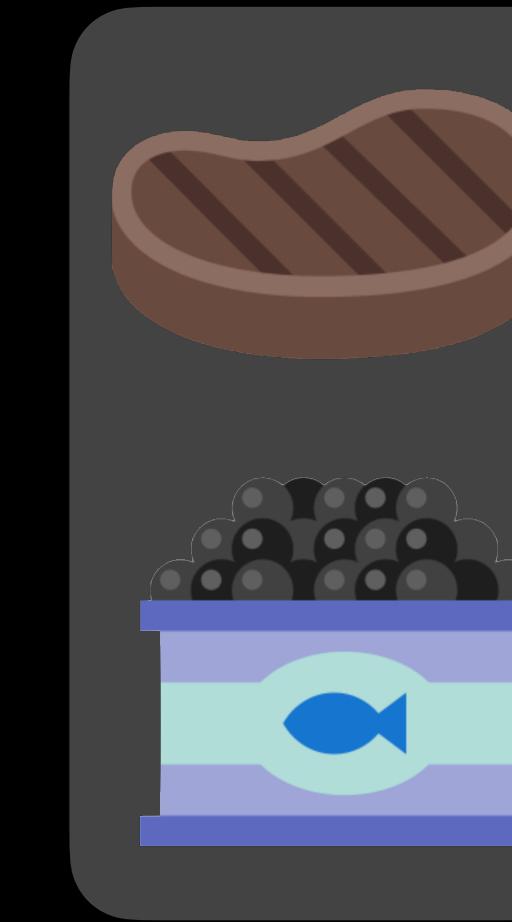
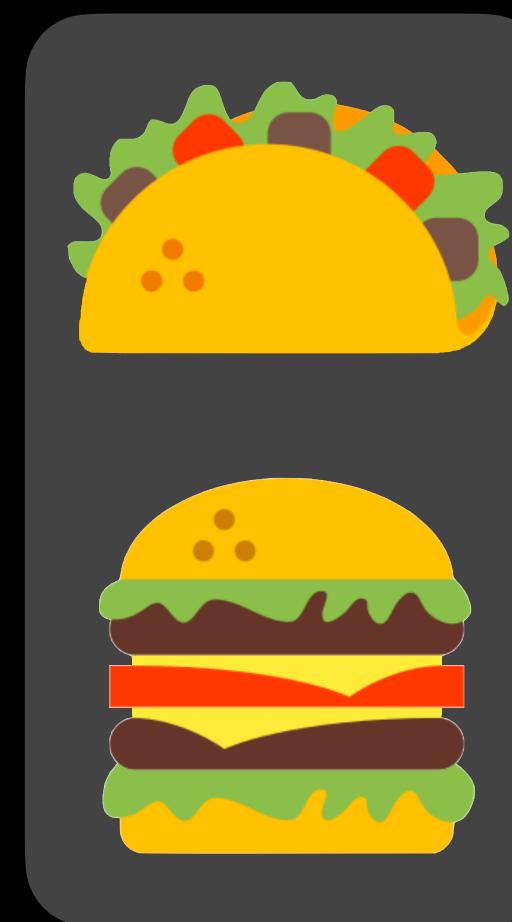
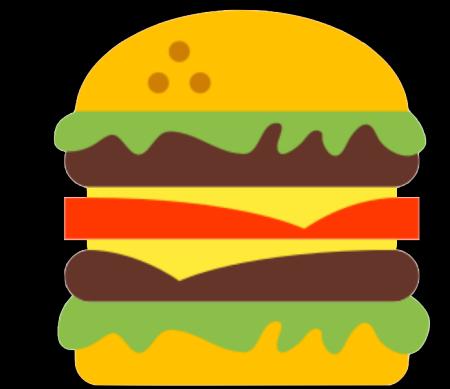
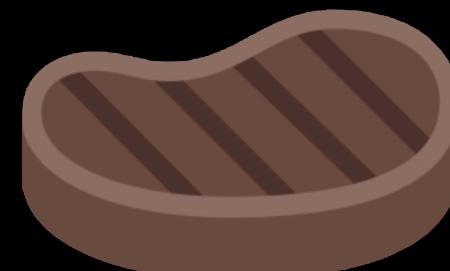
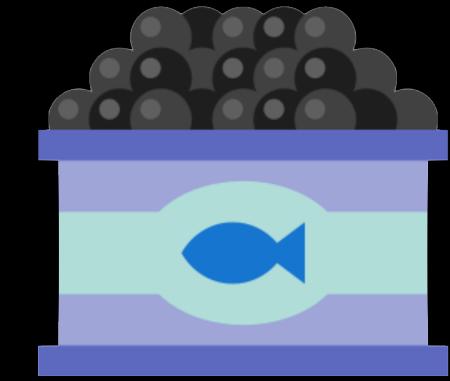


C_3

item features:

- price
- service speed
- wine selection

LCL example: restaurant selection



item features:

- price
- service speed
- wine selection

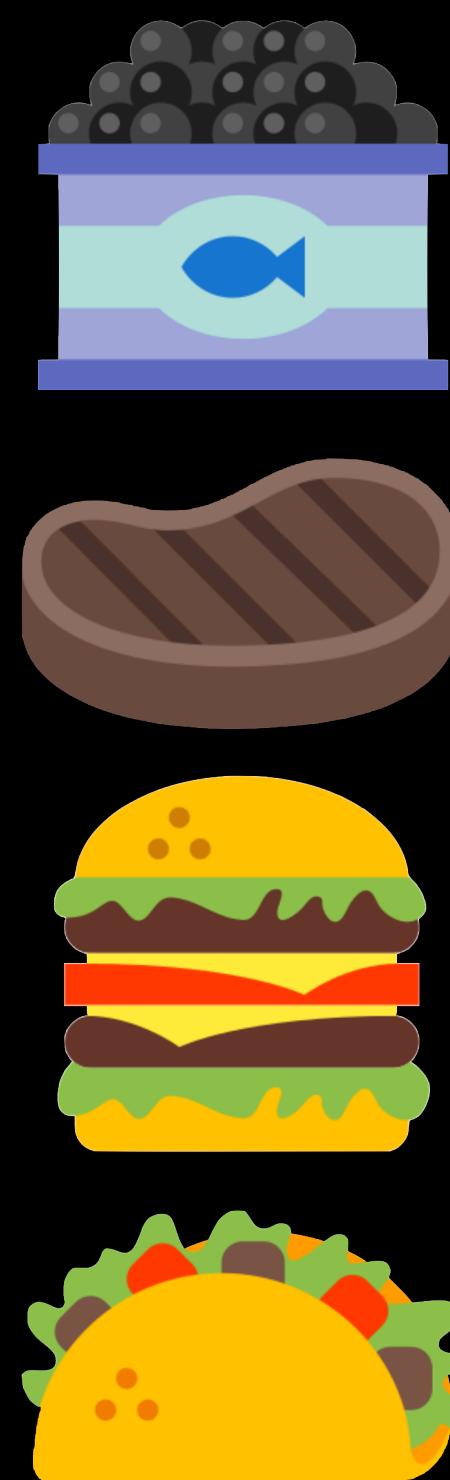
C_1

C_2

C_3

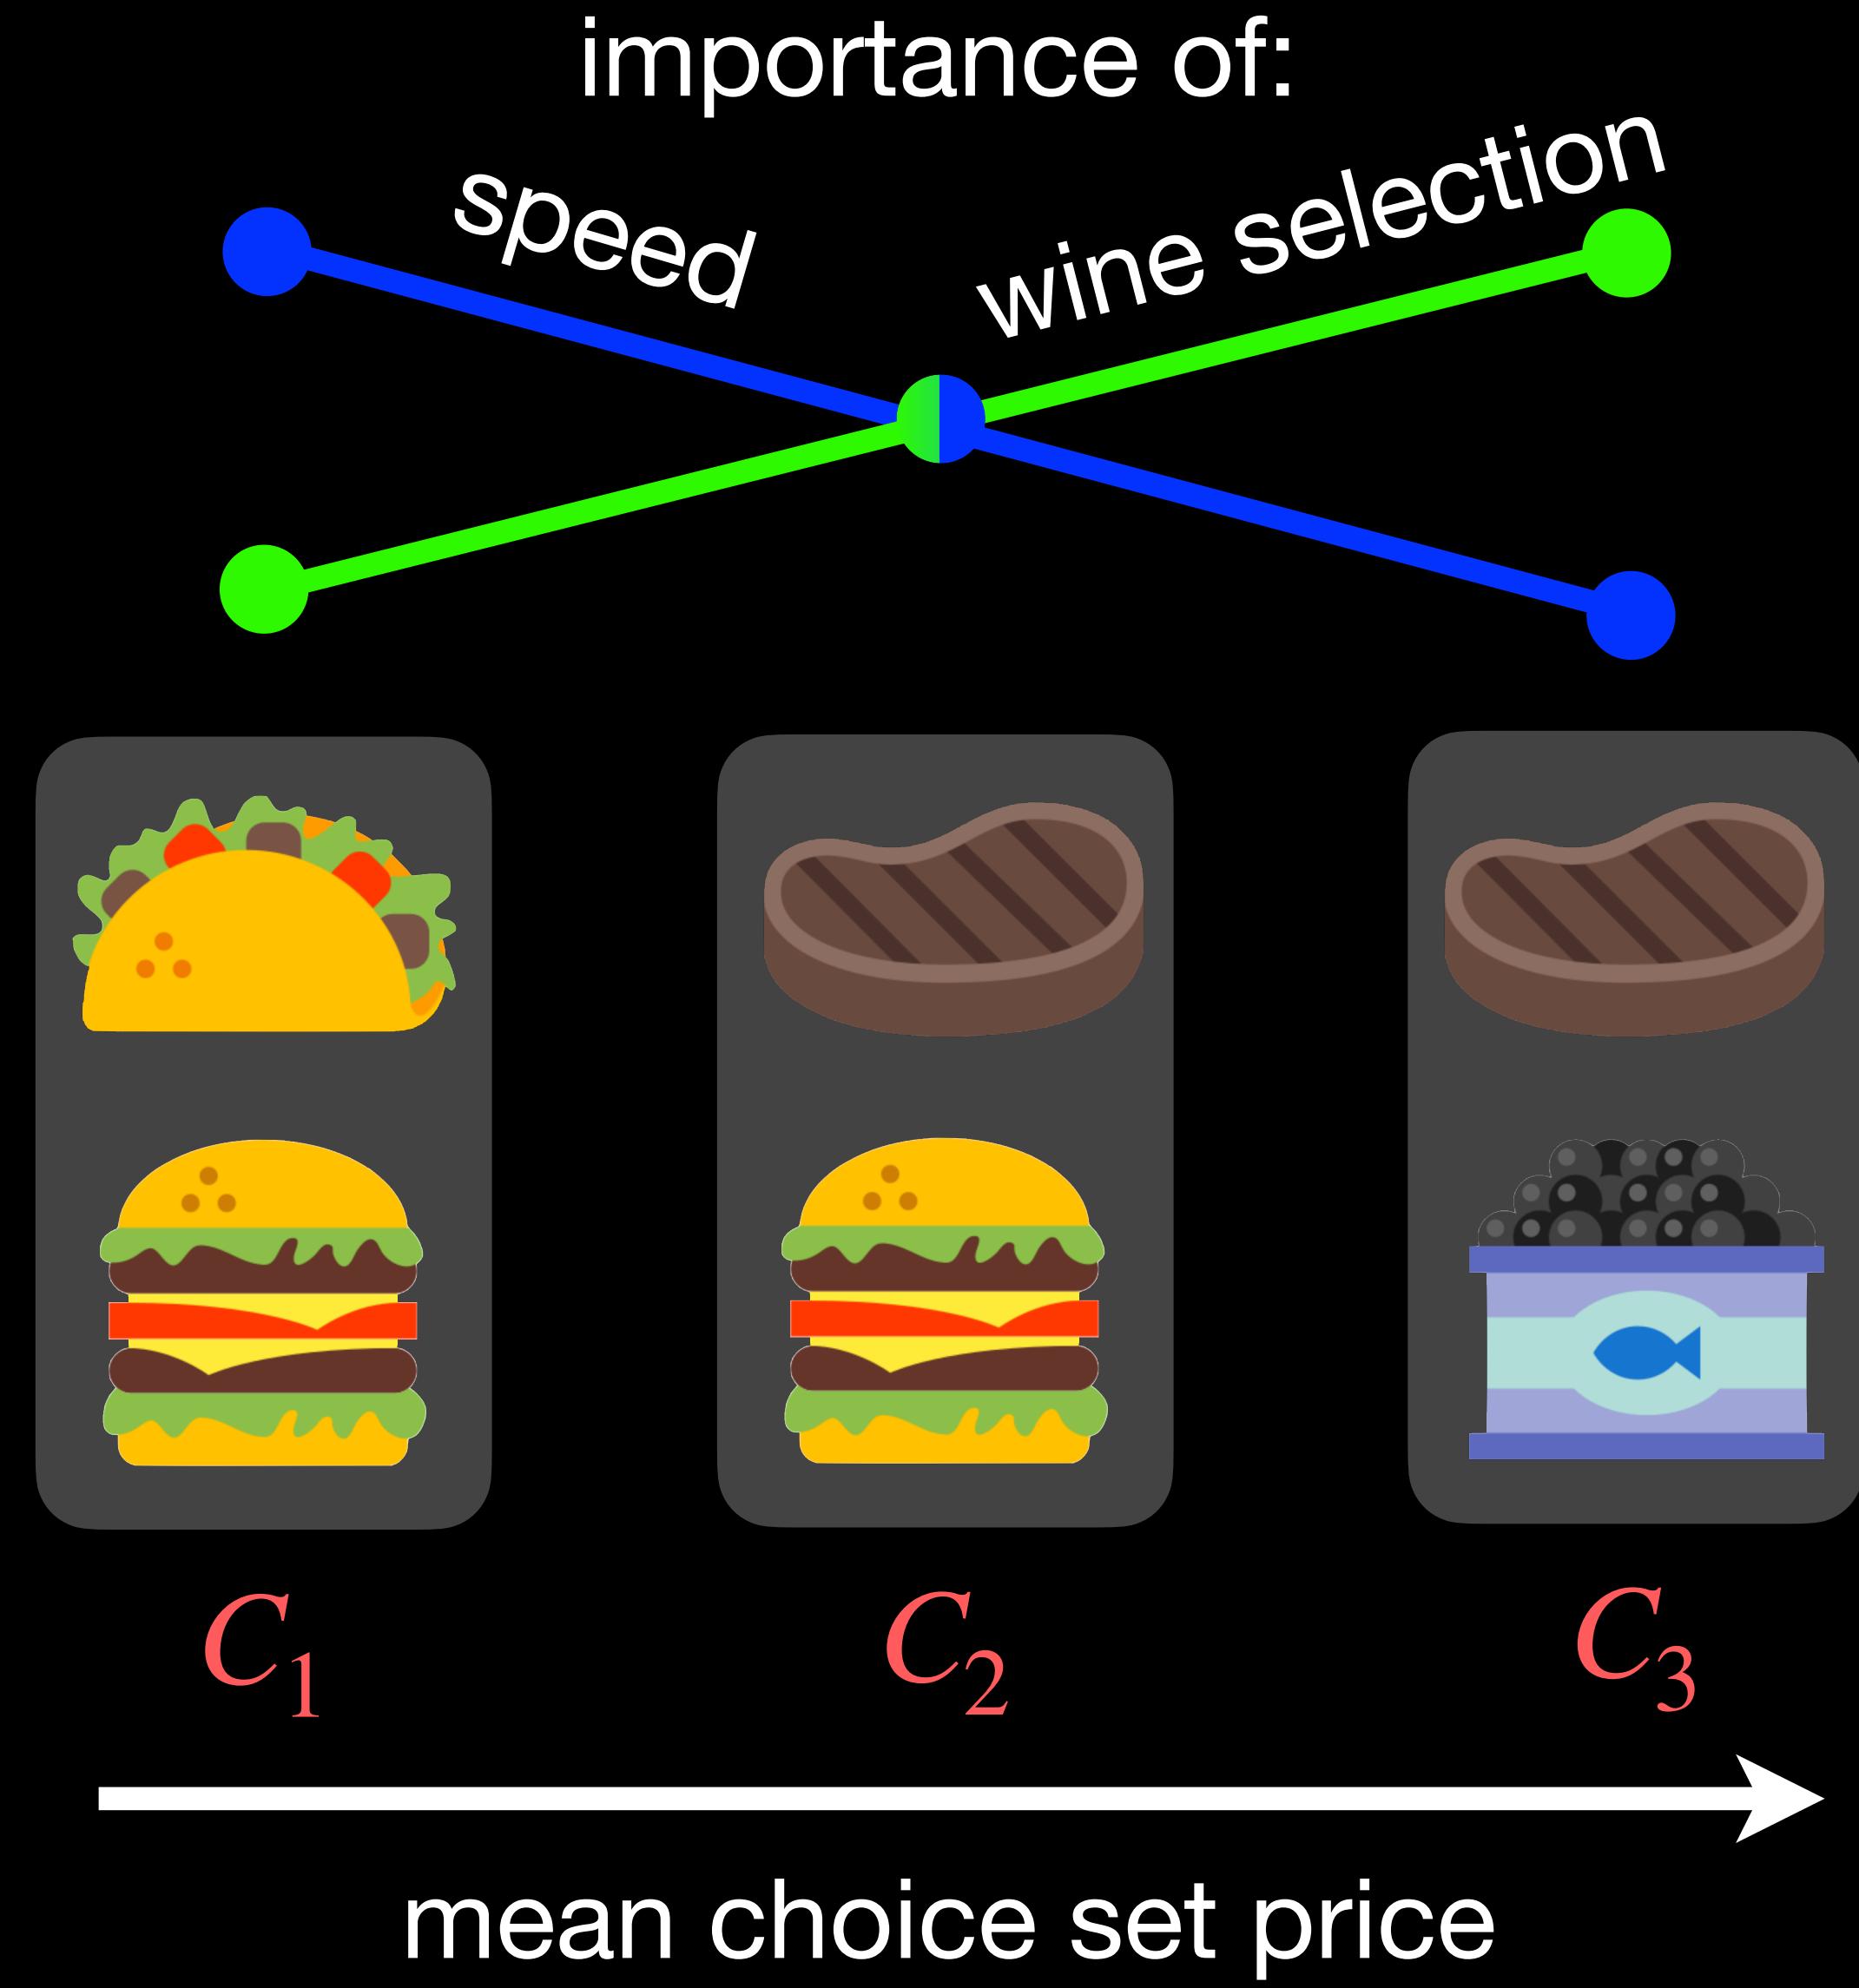
mean choice set price

LCL example: restaurant selection

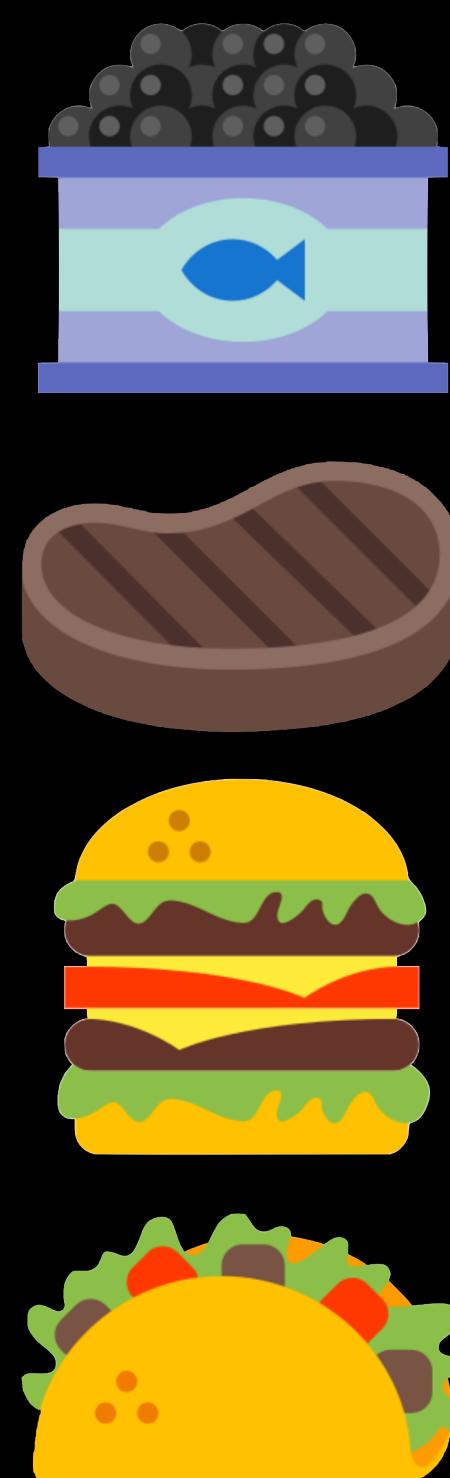


item features:

- price
- service speed
- wine selection

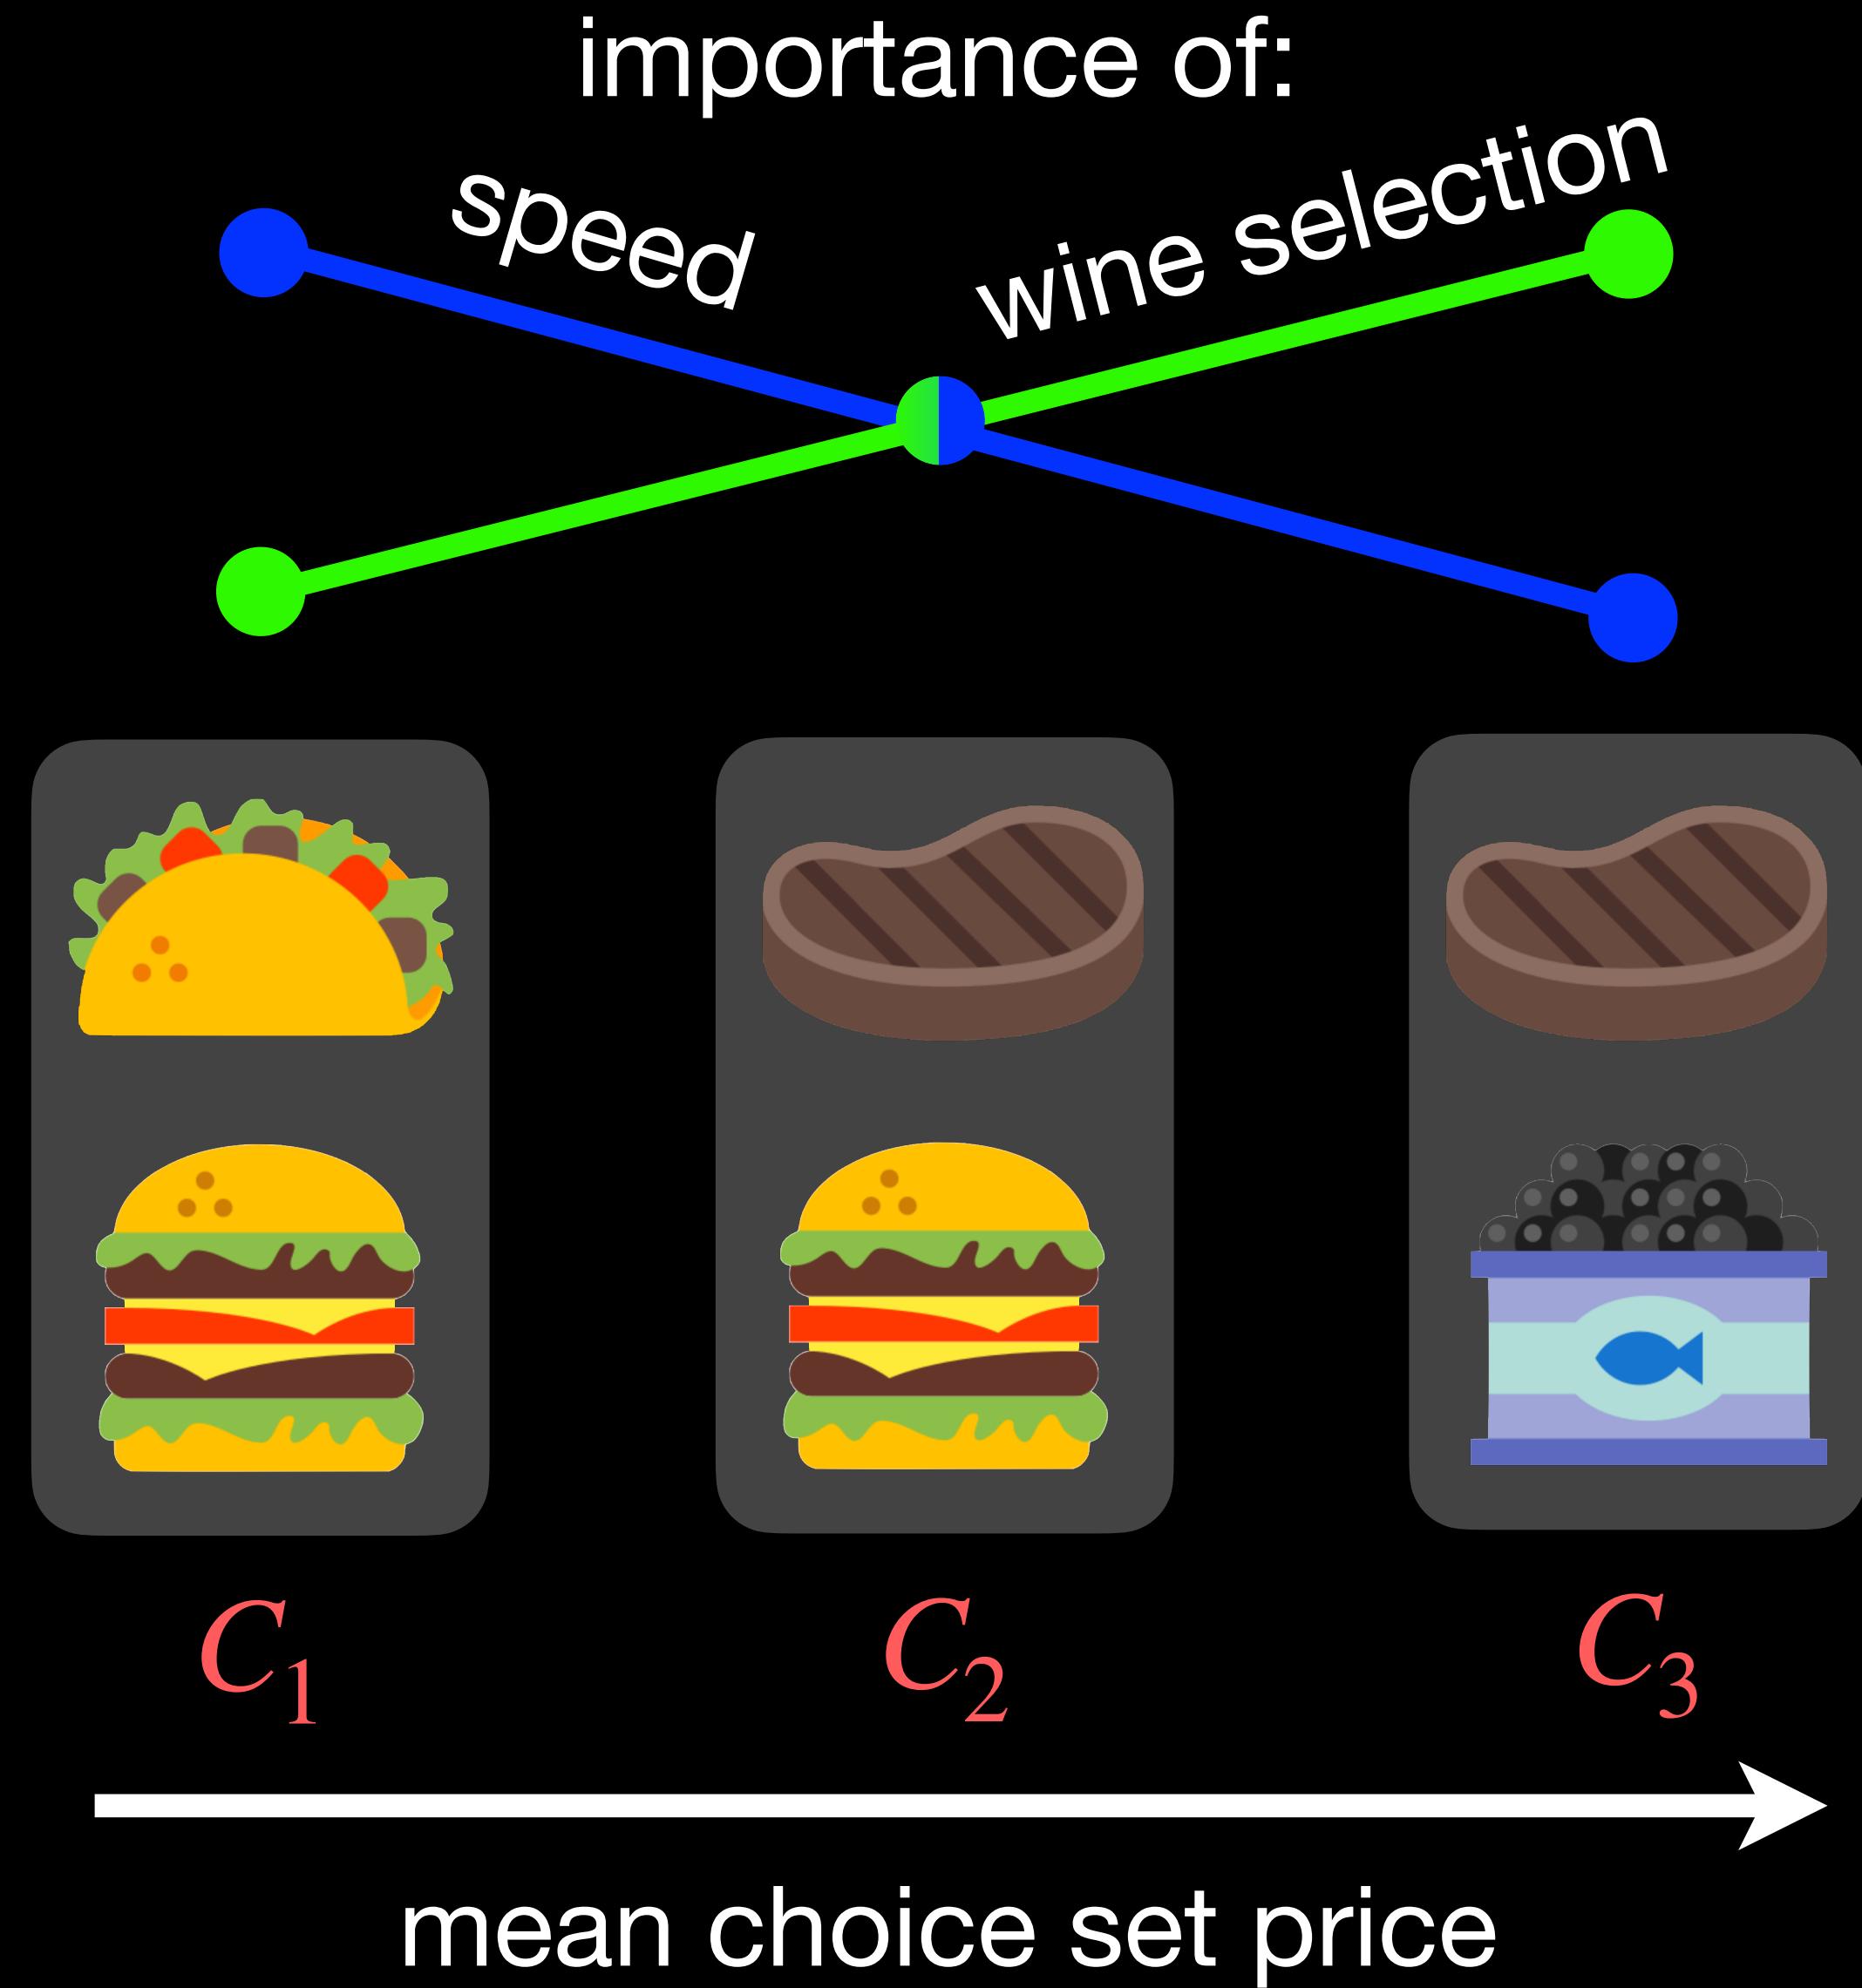


LCL example: restaurant selection



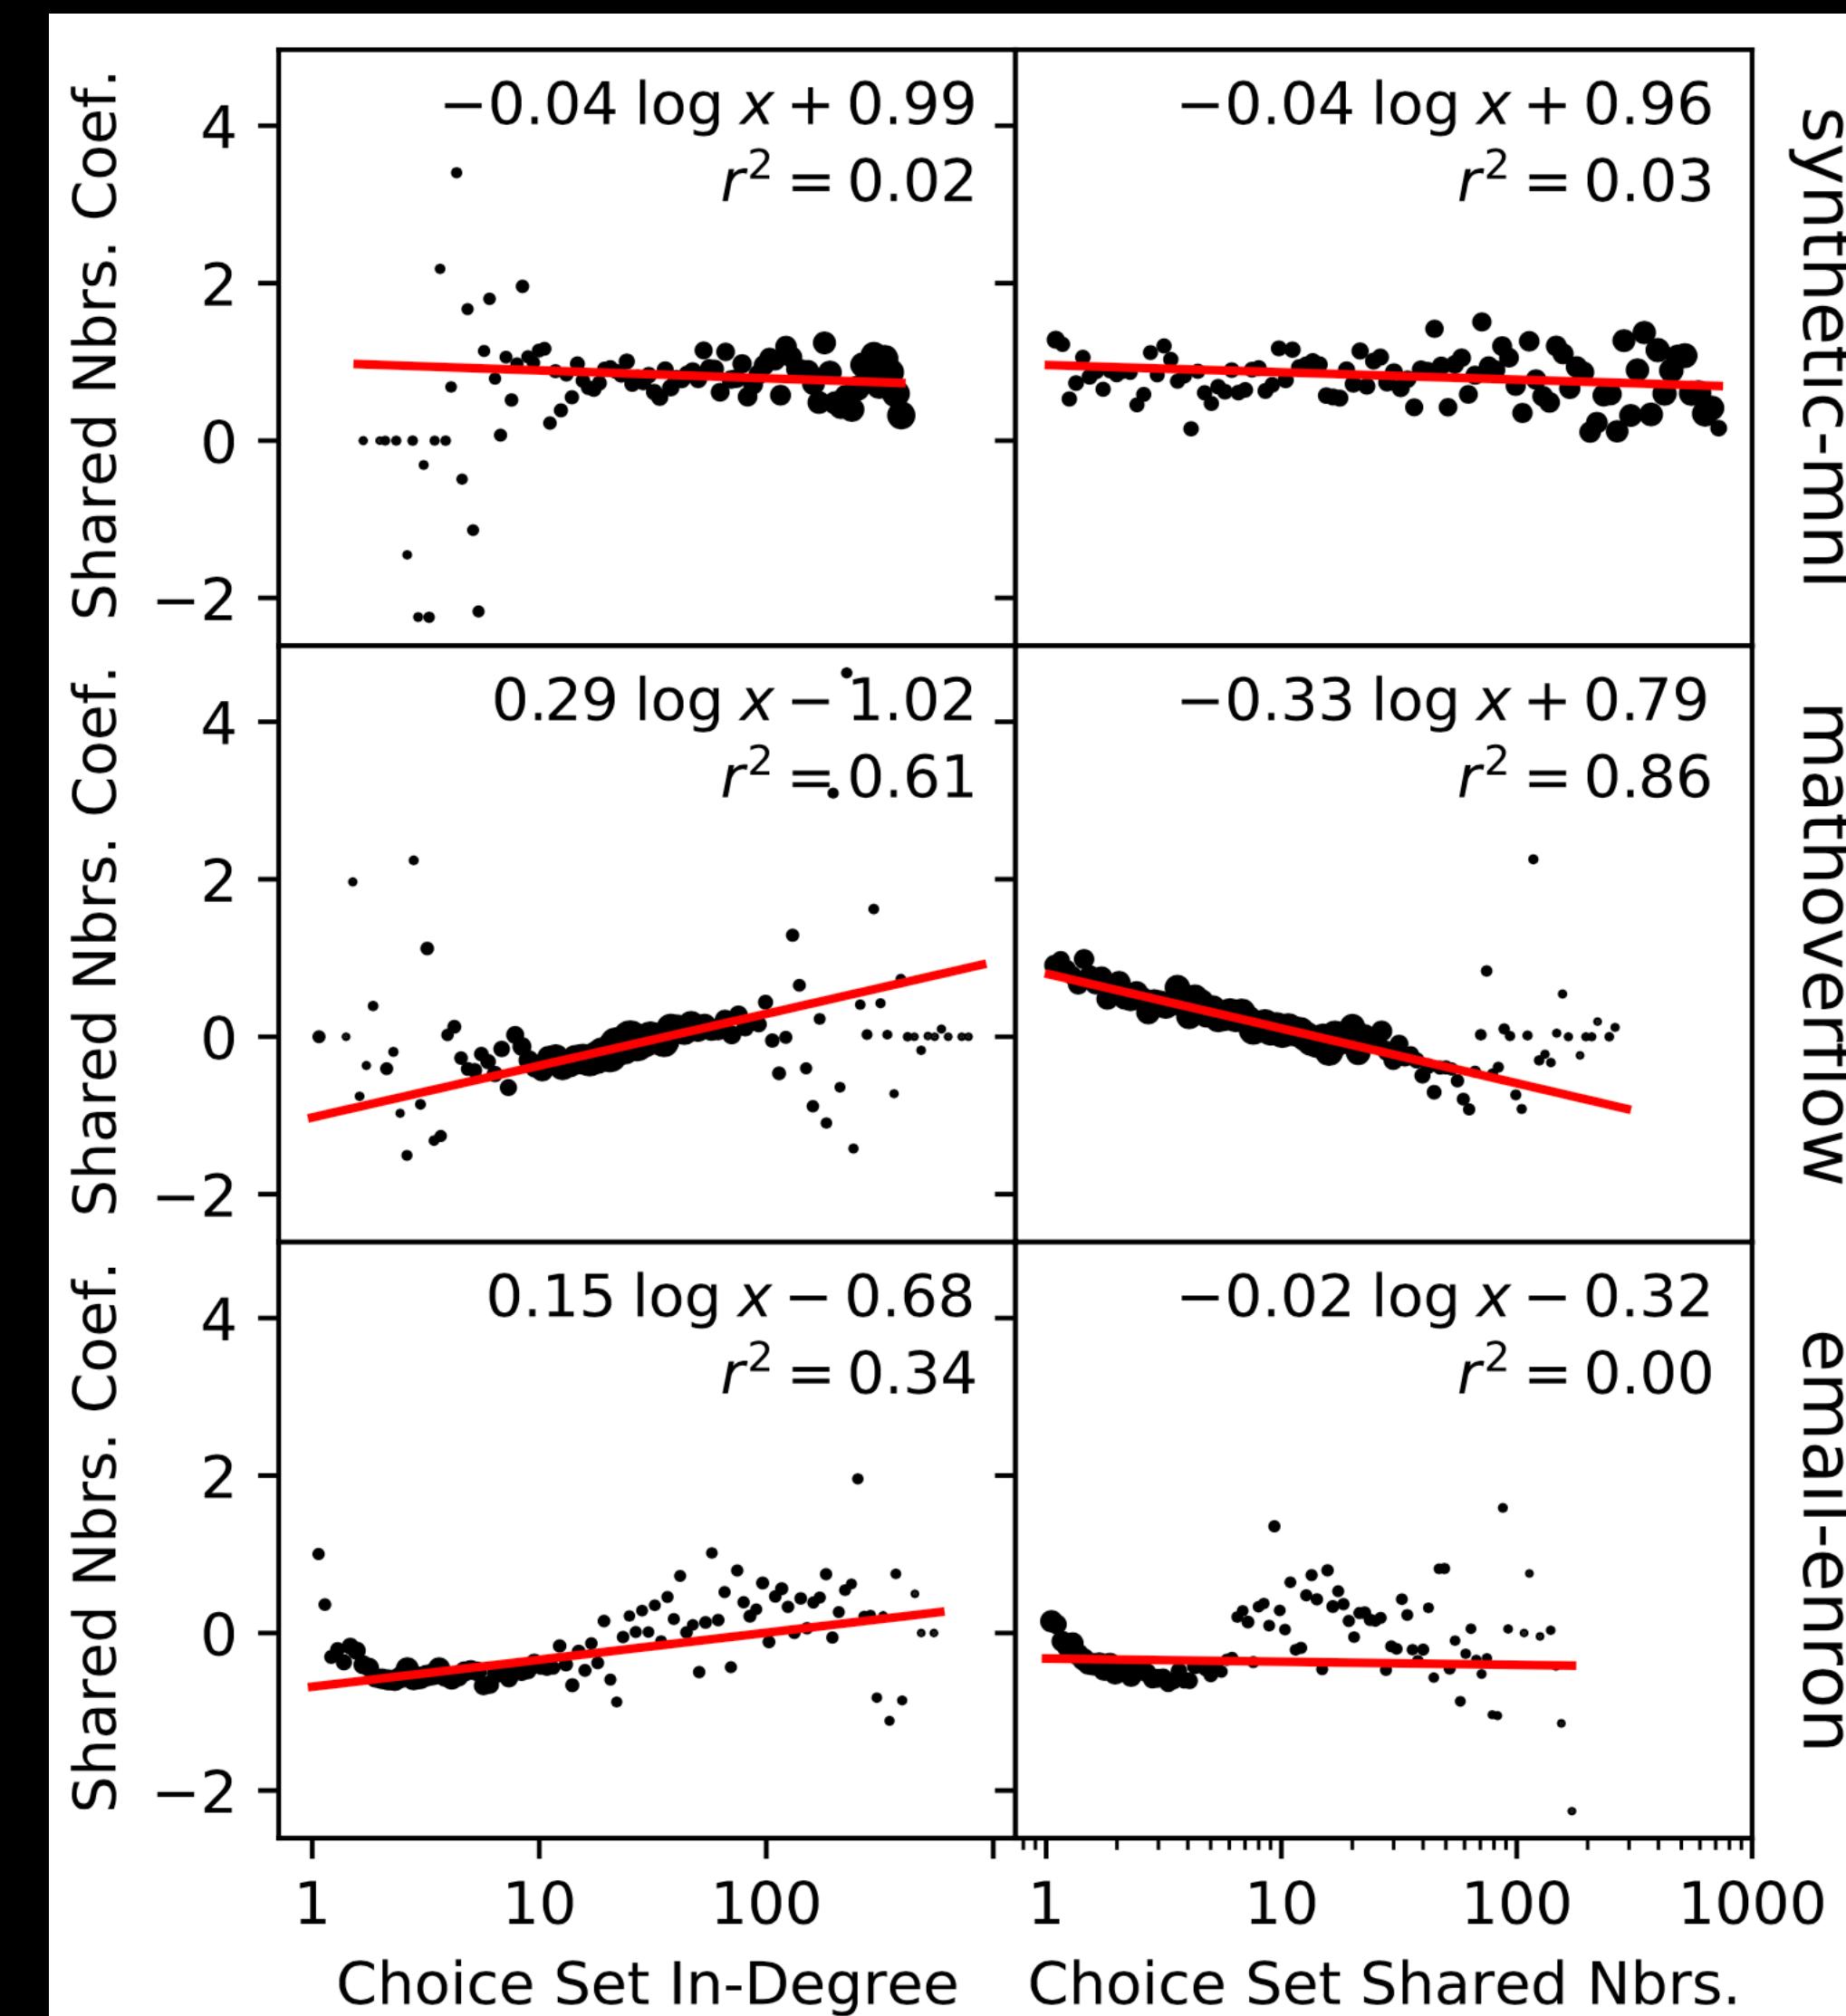
item features:

- price
- service speed
- wine selection



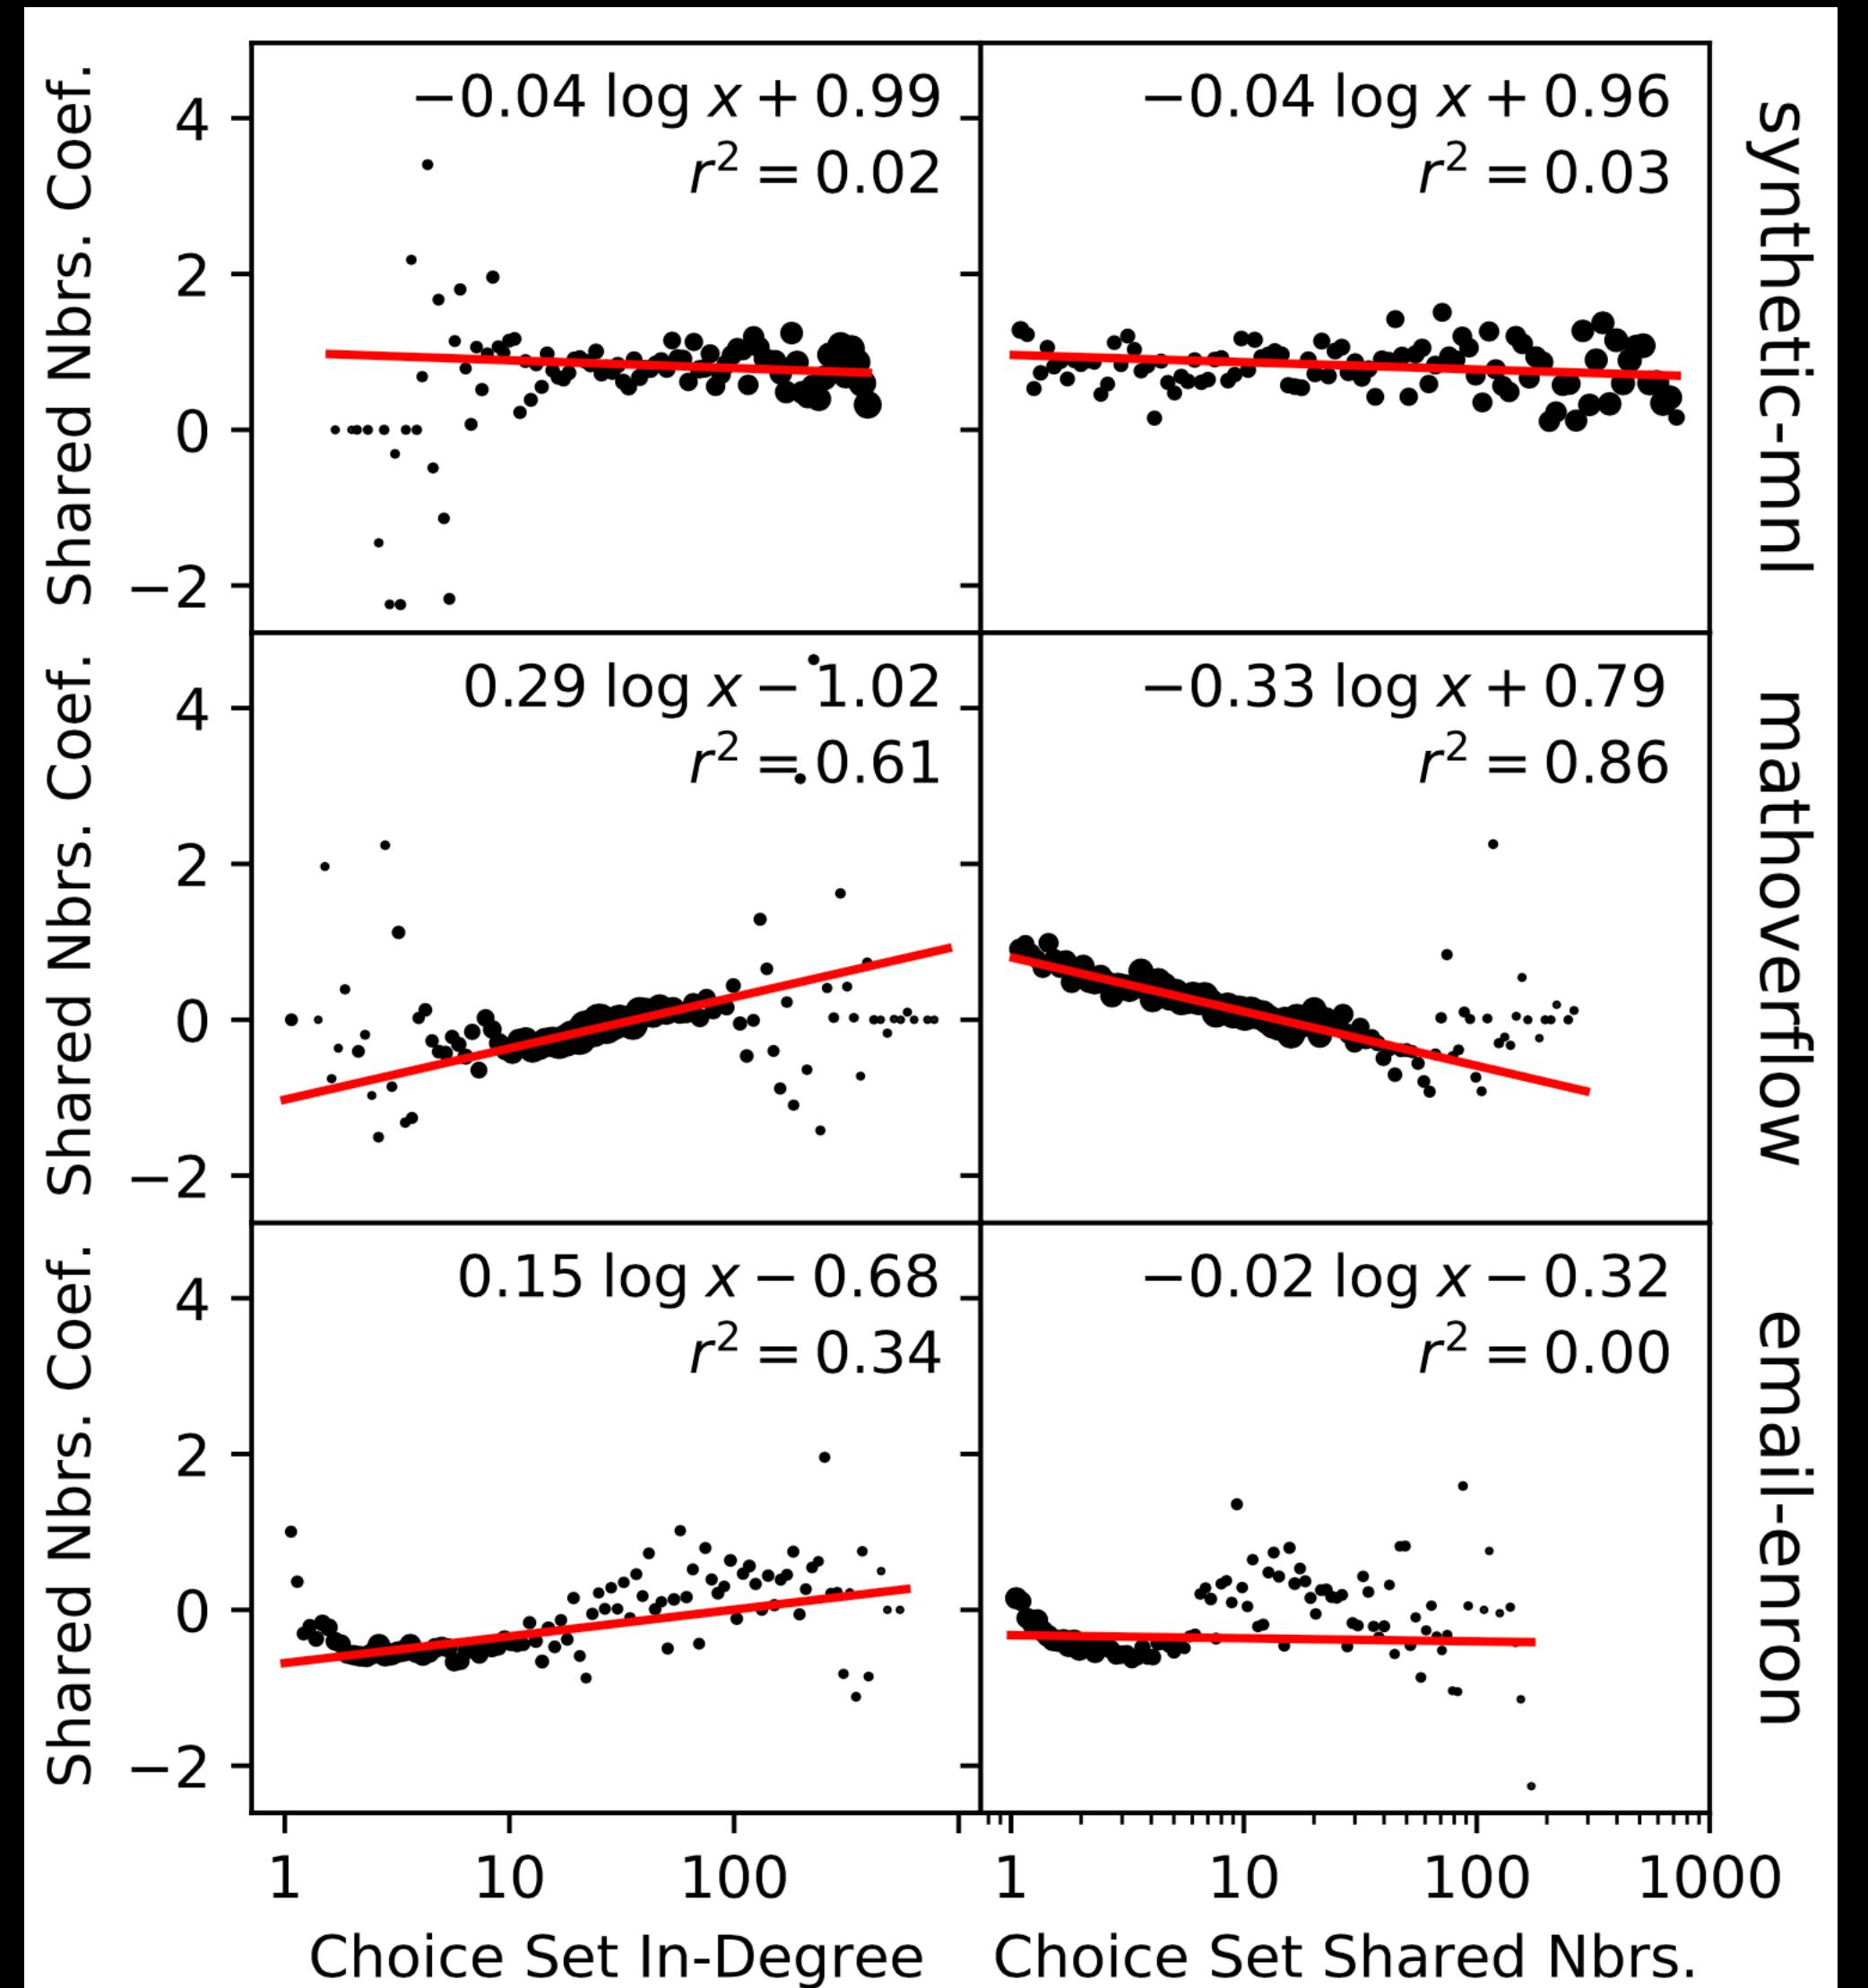
$$A = \begin{bmatrix} 0 & 0 & 0 \\ -1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix}$$

Linear feature context effects appear in real data

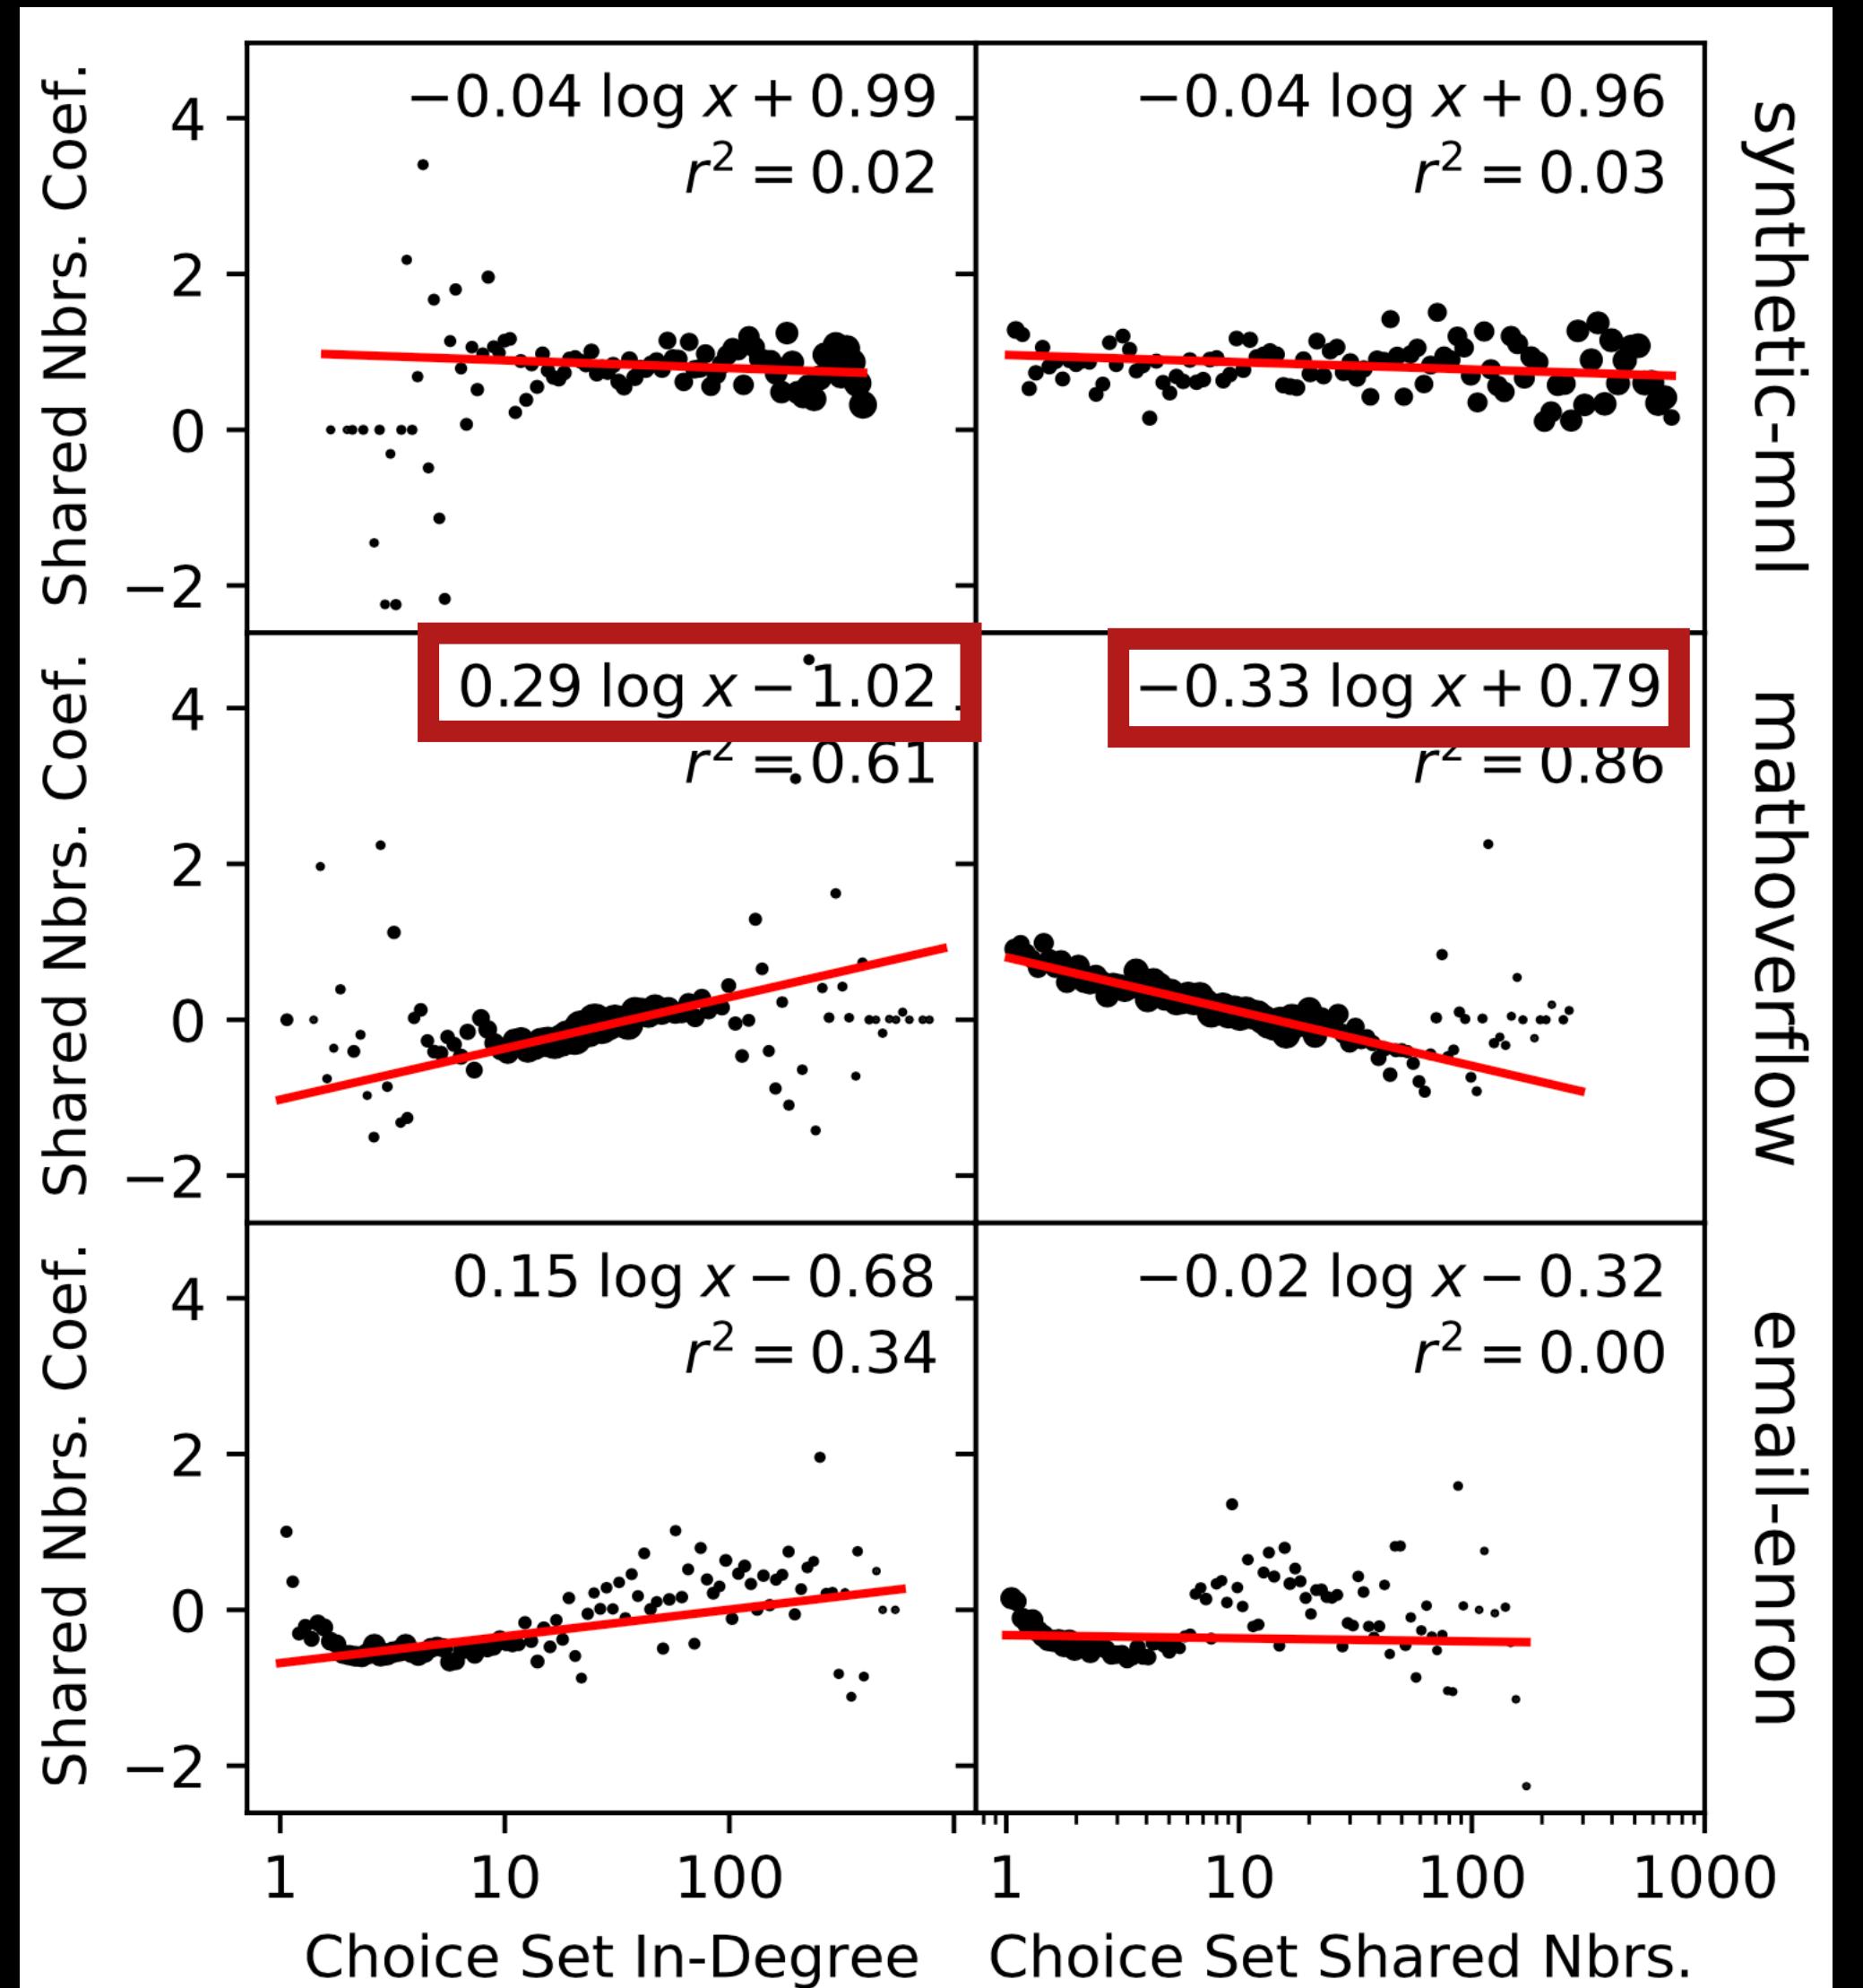


3. Decomposed LCL

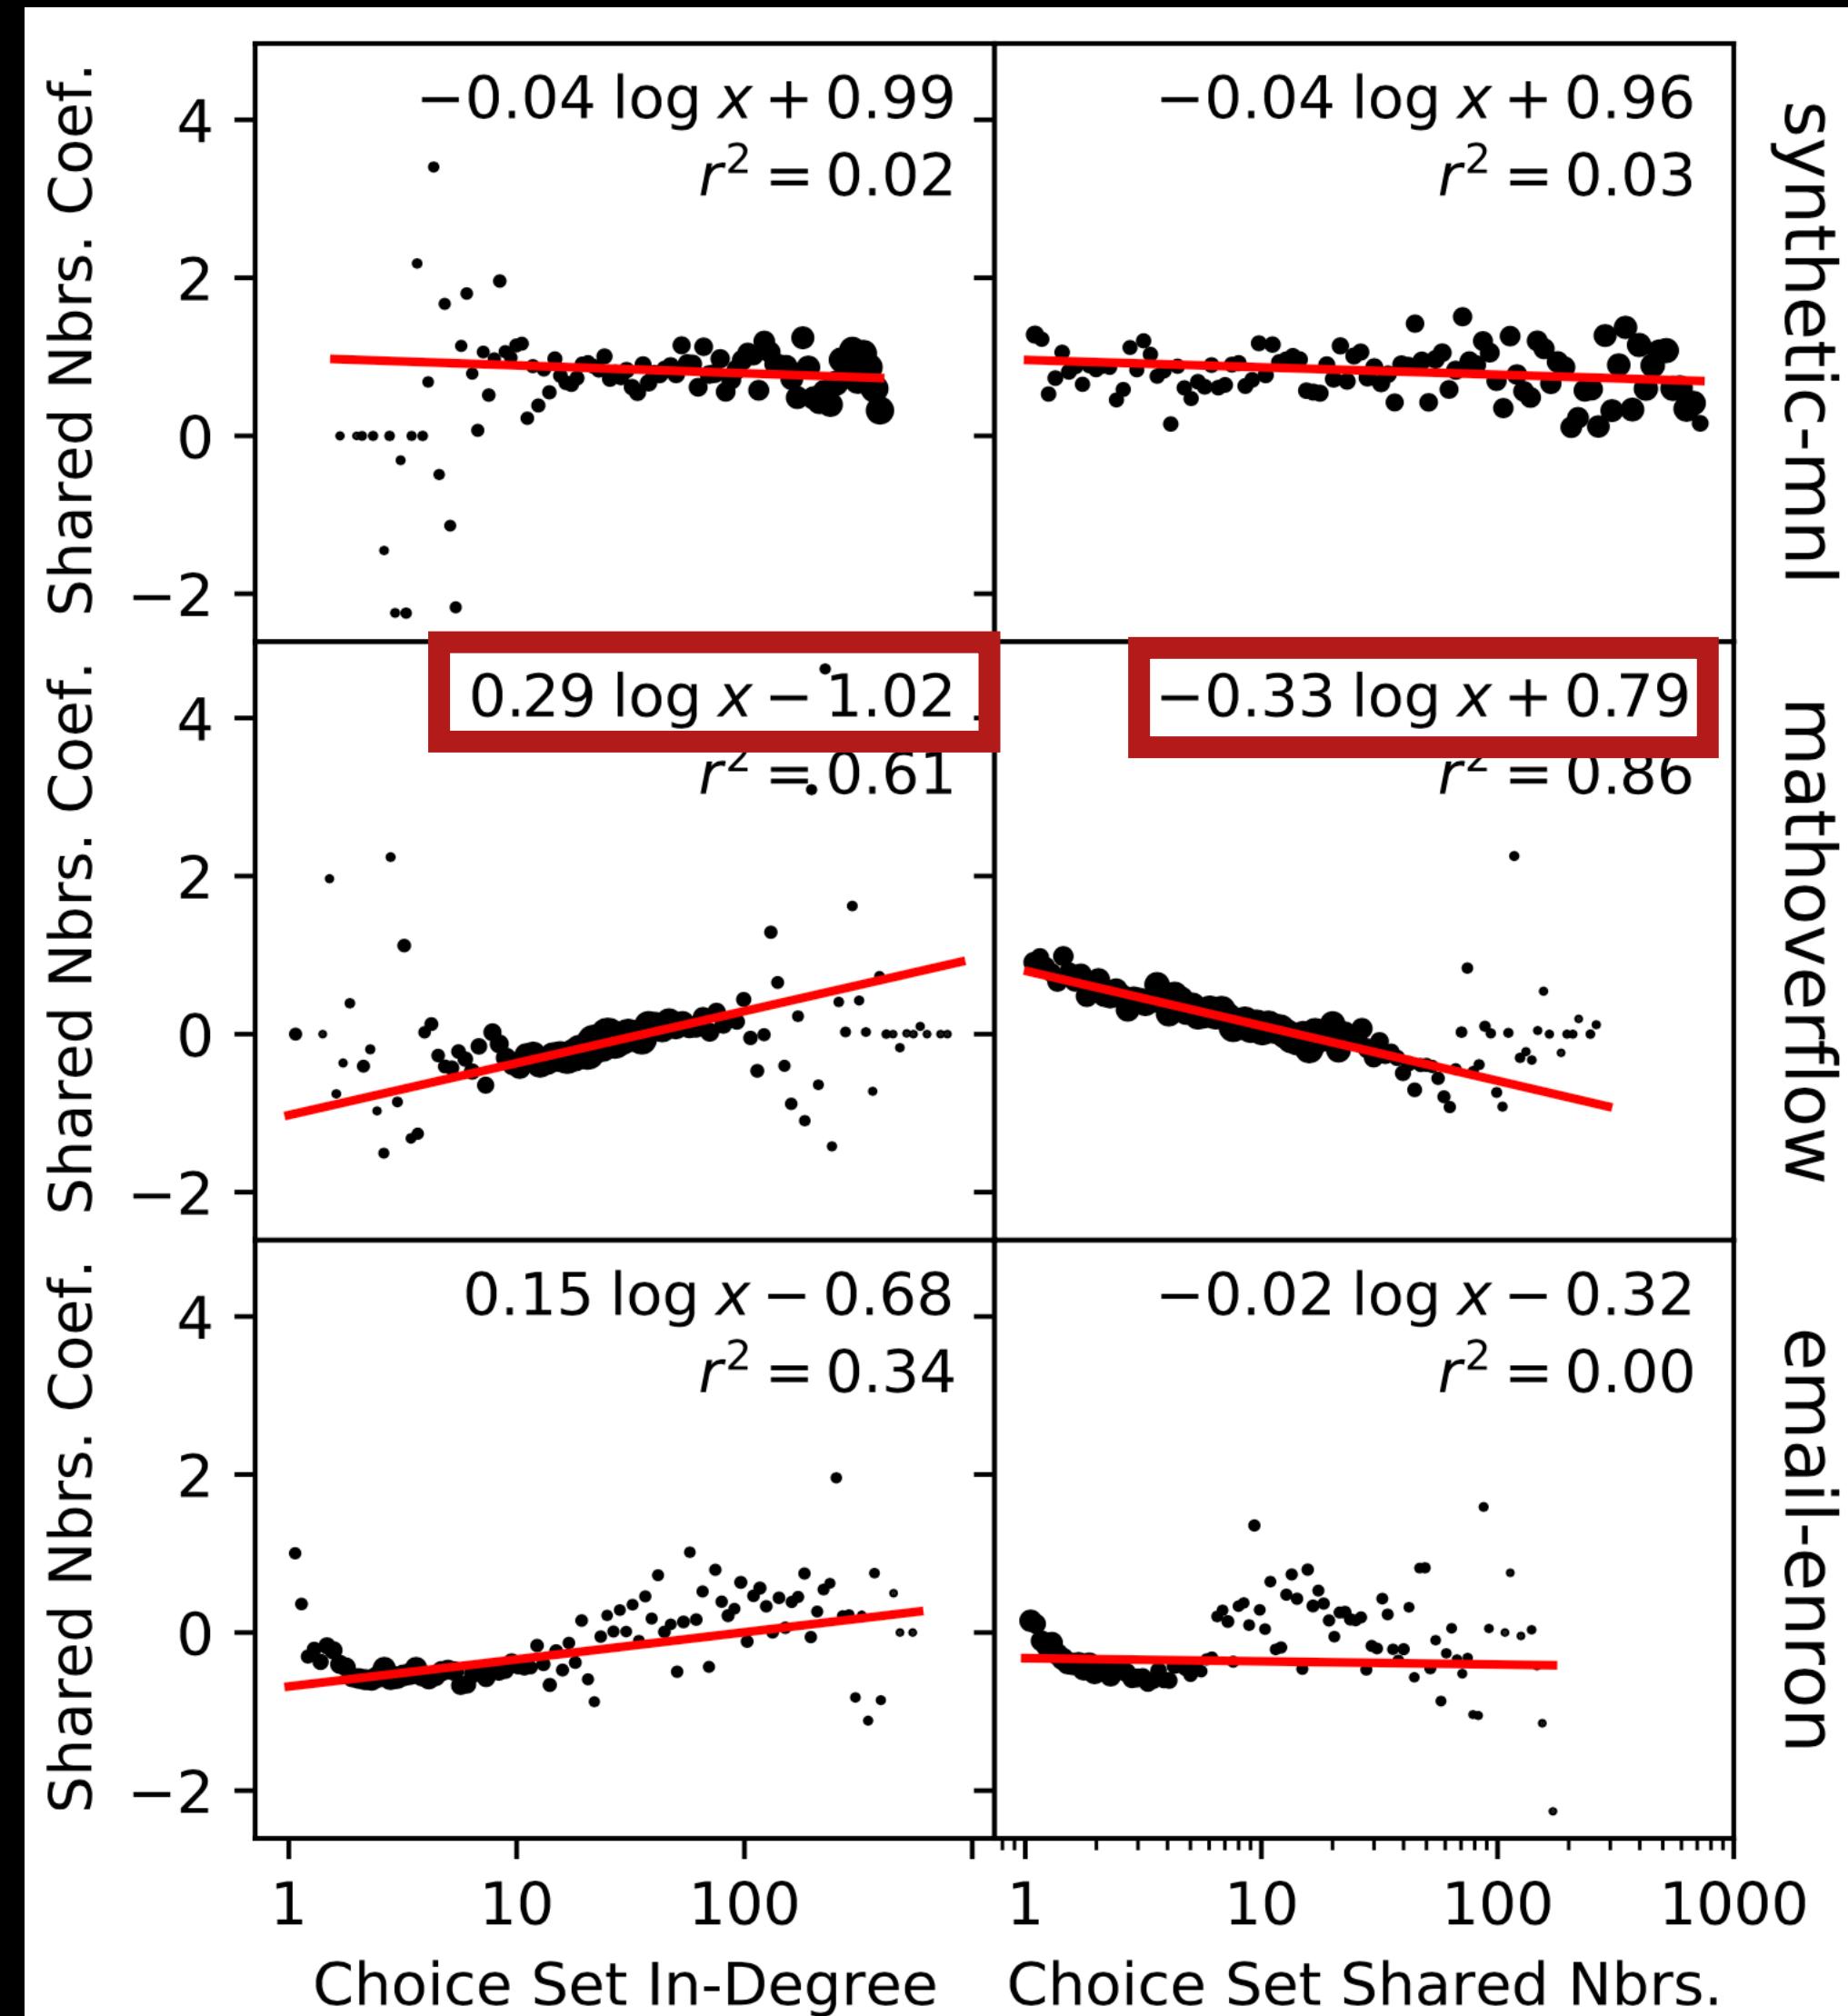
Motivation: the problem of intercepts



Motivation: the problem of intercepts



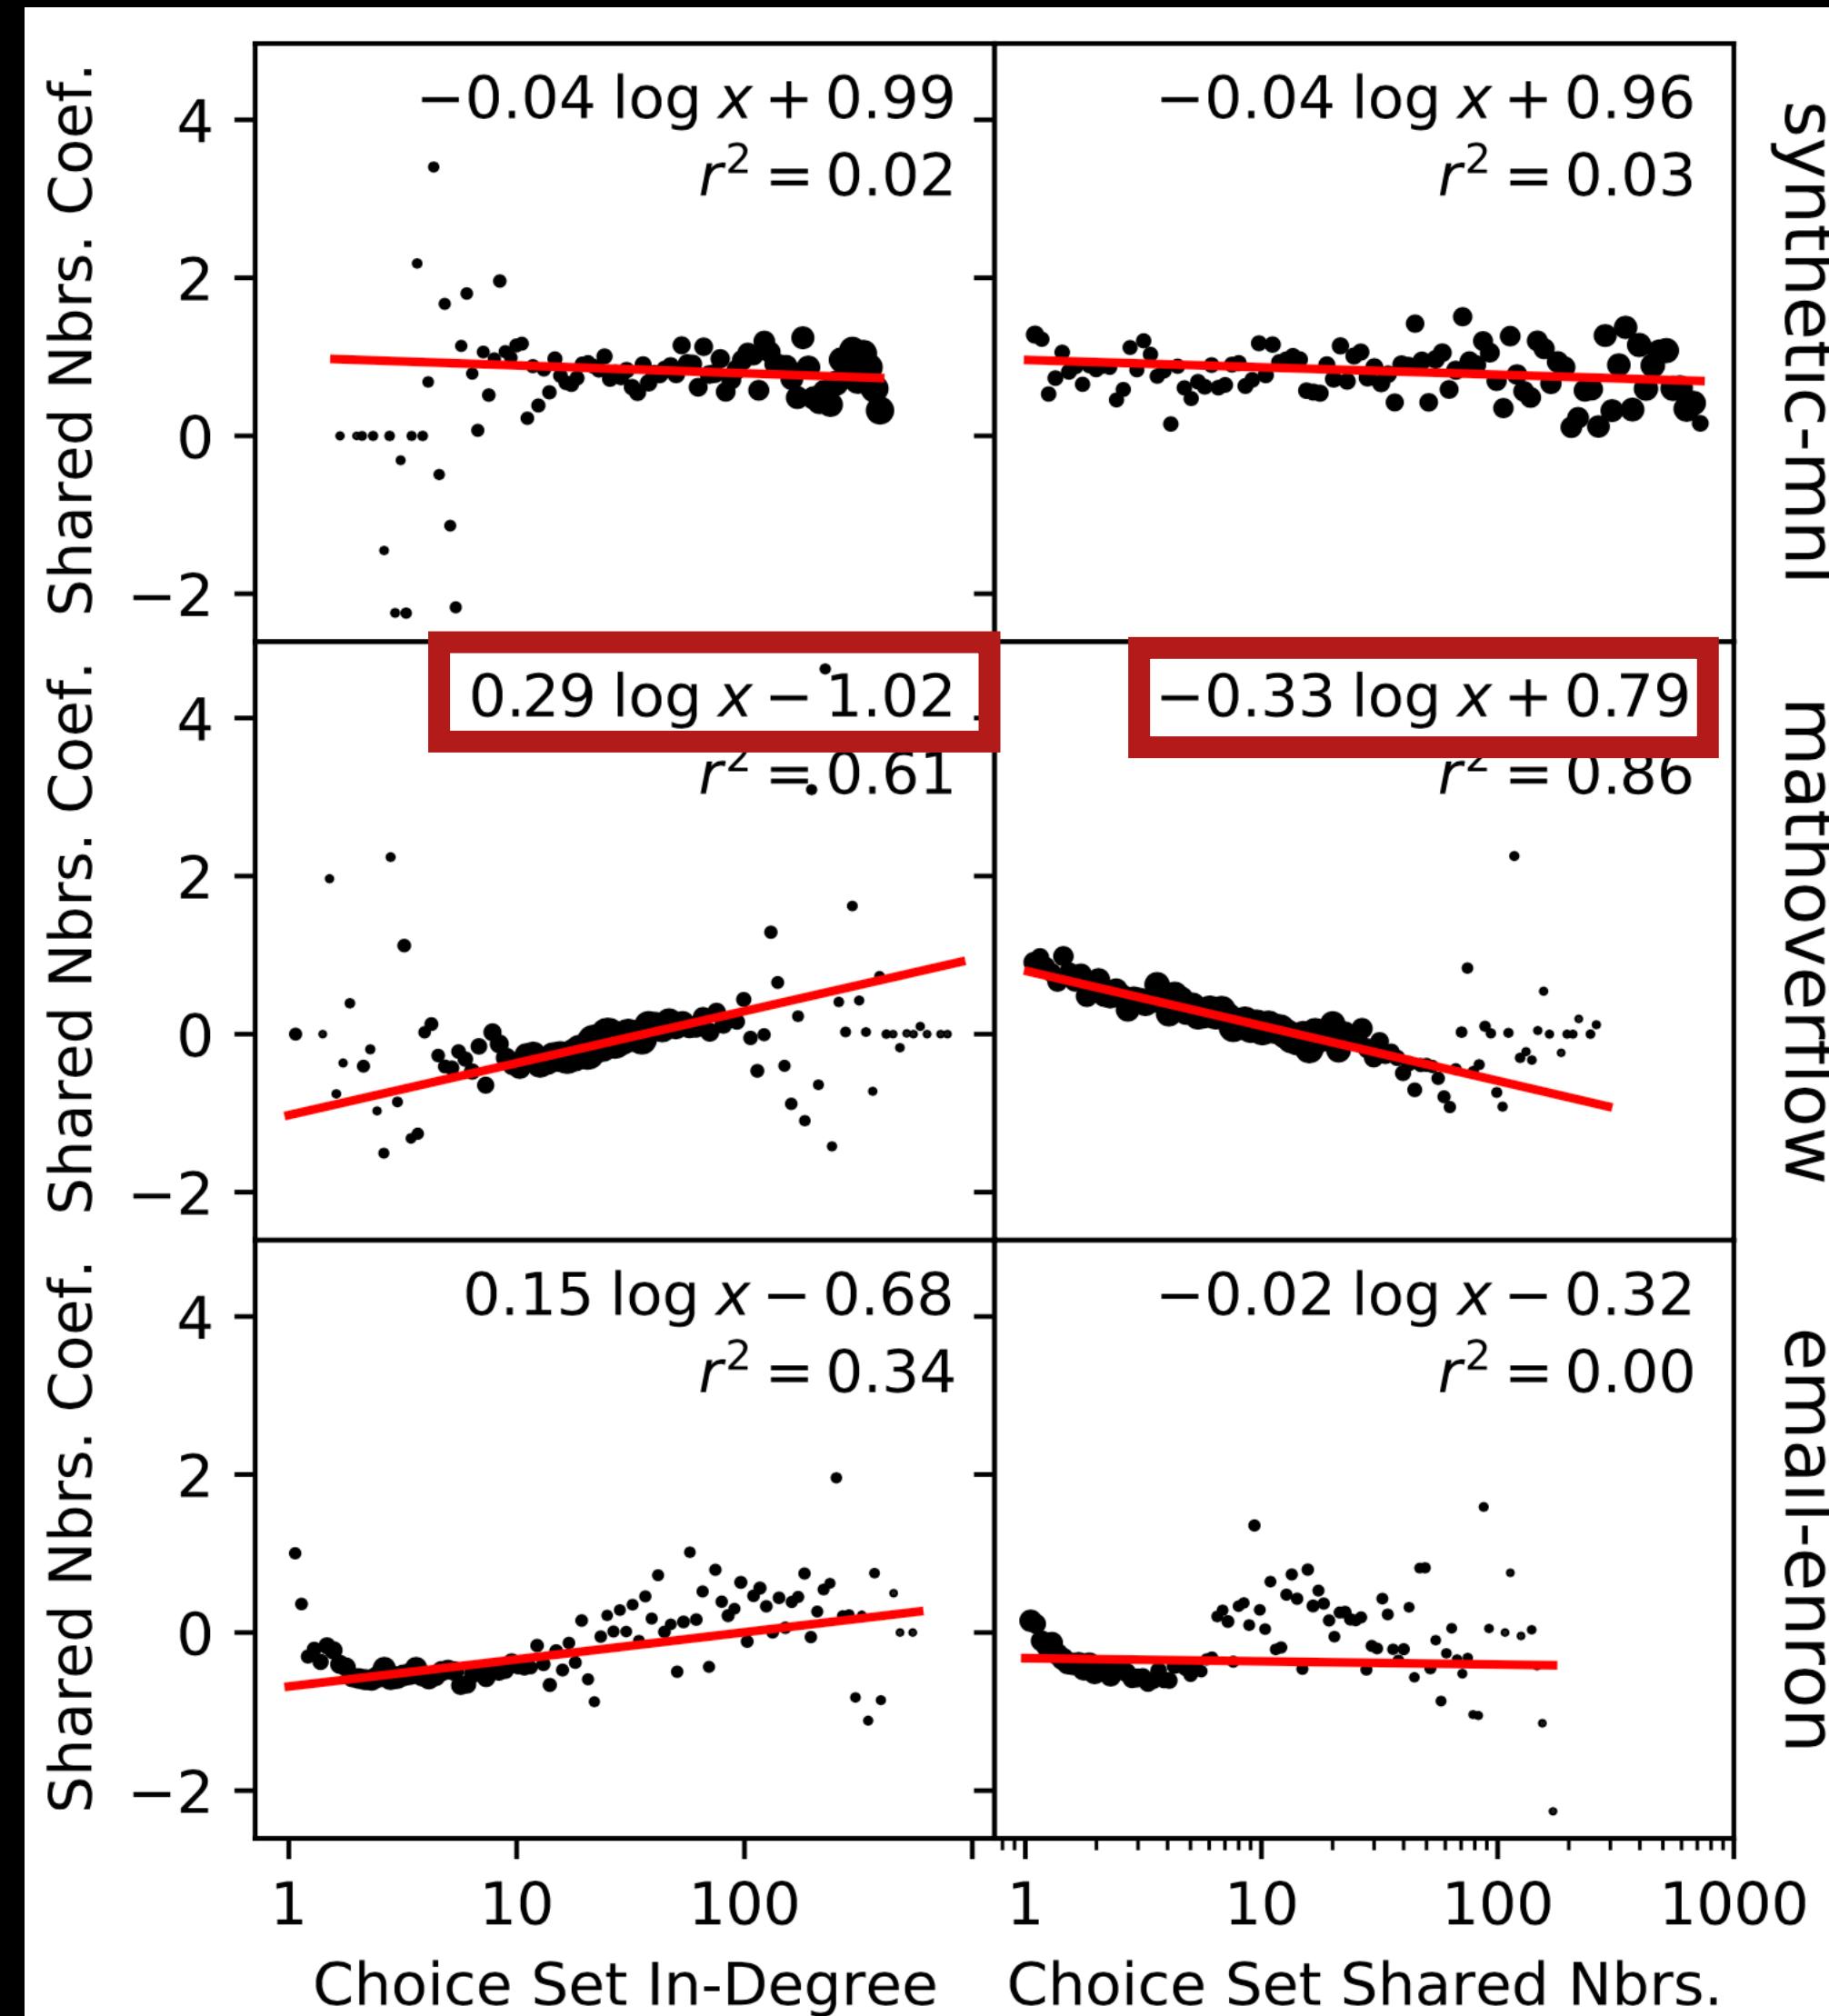
Motivation: the problem of intercepts



what happens with mean vector zero?

synthetic-mnl mathoverflow email-enron

Motivation: the problem of intercepts



what happens with mean vector zero?

$$\Pr(i \mid C) = \frac{\exp([\theta + Ax_C]^T x_i)}{\sum_{j \in C} \exp([\theta + Ax_C]^T x_j)}$$

d intercepts d^2 slopes

One option: *Decomposed Linear Context Logit (DLCL)*

One option: *Decomposed Linear Context Logit (DLCL)*

sub-model for the context effect exerted by each feature

One option: *Decomposed Linear Context Logit (DLCL)*

sub-model for the context effect exerted by each feature

→ $u_{i,C,k}$: utility of i in C in the context of feature k

$$u_{i,C,k} = [B_k + F_k(C)]^T x_i$$

(matrix subscript = column)

One option: *Decomposed Linear Context Logit (DLCL)*

sub-model for the context effect exerted by each feature

→ $u_{i,C,k}$: utility of i in C in the context of feature k

$$u_{i,C,k} = [B_k + F_k(C)]^T x_i$$

(matrix subscript = column)

→ B : matrix of intercepts

One option: *Decomposed Linear Context Logit (DLCL)*

sub-model for the context effect exerted by each feature

→ $u_{i,C,k}$: utility of i in C in the context of feature k

$$u_{i,C,k} = [B_k + F_k(C)]^T x_i$$

(matrix subscript = column)

→ B : matrix of intercepts

→ LCL assumptions yields $F_k(C) = A_k(x_C)_k$

One option: *Decomposed Linear Context Logit (DLCL)*

sub-model for the context effect exerted by each feature

→ $u_{i,C,k}$: utility of i in C in the context of feature k

$$u_{i,C,k} = [B_k + F_k(C)]^T x_i$$

(matrix subscript = column)

→ B : matrix of intercepts

→ LCL assumptions yields $F_k(C) = A_k(x_C)_k$

→ combine sub-models in mixture (w/ proportions π)

One option: *Decomposed Linear Context Logit (DLCL)*

sub-model for the context effect exerted by each feature

→ $u_{i,C,k}$: utility of i in C in the context of feature k

$$u_{i,C,k} = [B_k + F_k(C)]^T x_i$$

(matrix subscript = column)

→ B : matrix of intercepts

→ LCL assumptions yields $F_k(C) = A_k(x_C)_k$

→ combine sub-models in mixture (w/ proportions π)

$$\Pr(i \mid C) = \sum_{k=1}^d \pi_k \frac{\exp([B_k + A_k(x_C)_k]^T x_i)}{\sum_{j \in C} \exp([B_k + A_k(x_C)_k]^T x_j)}$$

4. Identifiability and estimation

LCL identifiability, fully characterized

LCL identifiability, fully characterized

model is *identifiable* from dataset \mathcal{D} if no two parameter values result in the same probability distribution

LCL identifiability, fully characterized

model is *identifiable* from dataset \mathcal{D} if no two parameter values result in the same probability distribution

→ important for inference and interpretation

LCL identifiability, fully characterized

model is *identifiable* from dataset \mathcal{D} if no two parameter values result in the same probability distribution

→ important for inference and interpretation

Theorem 1. A d -feature linear context logit is identifiable from a dataset \mathcal{D} if and only if

$$\text{span} \left\{ \begin{bmatrix} x_C \\ 1 \end{bmatrix} \otimes (x_i - x_C) \mid C \in \mathcal{C}_{\mathcal{D}}, i \in C \right\} = \mathbb{R}^{d^2+d}. \quad (6)$$

($\mathcal{C}_{\mathcal{D}}$: unique choice sets in \mathcal{D} , \otimes : Kronecker product)

LCL identifiability, fully characterized

model is *identifiable* from dataset \mathcal{D} if no two parameter values result in the same probability distribution

→ important for inference and interpretation

Theorem 1. A d -feature linear context logit is identifiable from a dataset \mathcal{D} if and only if

$$\text{span} \left\{ \begin{bmatrix} x_C \\ 1 \end{bmatrix} \otimes (x_i - x_C) \mid C \in \mathcal{C}_{\mathcal{D}}, i \in C \right\} = \mathbb{R}^{d^2+d}. \quad (6)$$

($\mathcal{C}_{\mathcal{D}}$: unique choice sets in \mathcal{D} , \otimes : Kronecker product)

intuition: need varied choice sets containing varied items

LCL identifiability, more intuition

LCL identifiability, more intuition

PROPOSITION 2. *No d -feature linear context logit is identifiable from a dataset \mathcal{D} if it does not include a set of $d + 1$ choice sets with affinely independent mean feature vectors.*

LCL identifiability, more intuition

necessary:

PROPOSITION 2. No d -feature linear context logit is identifiable from a dataset \mathcal{D} if it does not include a set of $d + 1$ choice sets with affinely independent mean feature vectors.

sufficient:

PROPOSITION 3. If a dataset contains $d + 1$ distinct choice sets C_0, \dots, C_d such that

- i. *the set of mean feature vectors $\{x_{C_0}, \dots, x_{C_d}\}$ is affinely independent (the necessary condition from Proposition 2) and*
- ii. *in each choice set C_i , there is some set of $d + 1$ items with affinely independent features,*

then we can uniquely identify a d -feature LCL.

DLCL identifiability.... for another day

Mixture models are rough

see:

Chierichetti, F., Kumar, R., & Tomkins, A. (2018). Learning a mixture of two multinomial logits. In *International Conference on Machine Learning* (pp. 961-969).

Zhao, Z., Piech, P., & Xia, L. (2016). Learning mixtures of Plackett-Luce models. In *International Conference on Machine Learning* (pp. 2906-2914).

Grün, B., & Leisch, F. (2008). Identifiability of finite mixtures of multinomial logit models with varying and fixed effects. *Journal of Classification*, 25(2), 225-247.

LCL and DLCL Estimation

LCL and DLCL Estimation

LCL

negative log-likelihood convex, so use gradient descent

LCL and DLCL Estimation

LCL

negative log-likelihood convex, so use gradient descent

$$-\ell(\theta, A; \mathcal{D}) = \sum_{(i,C) \in \mathcal{D}} -\left(\theta + Ax_C\right)^T x_i + \log \sum_{j \in C} \exp\left(\left[\theta + Ax_C\right]^T x_j\right)$$

LCL and DLCL Estimation

LCL

negative log-likelihood convex, so use gradient descent

$$-\ell(\theta, A; \mathcal{D}) = \sum_{(i,C) \in \mathcal{D}} -\left(\theta + Ax_C\right)^T x_i + \log \sum_{j \in C} \exp\left(\left[\theta + Ax_C\right]^T x_j\right)$$

DLCL

log-likelihood not convex...
but, *expectation-maximization (EM)*
algorithm performs well

Table 9: Adam vs. EM algorithm for DLCL

Dataset	Adam NLL	EM NLL
DISTRICT	3206	3041
DISTRICT-SMART	3303	3144
EXPEDIA	837569	805055
SUSHI	9764	9709
CAR-A	1692	1684
CAR-B	1284	1246
CAR-ALT	7011	6369

5. Results on choice data

Choice datasets

Dataset	Choices	Features	Largest Choice Set
DISTRICT	5376	27	2
DISTRICT-SMART	5376	6	2
SUSHI	5000	6	10
EXPEDIA	276593	5	38
CAR-A	2675	4	2
CAR-B	2206	5	2
CAR-ALT	4654	21	6

Choice datasets

favorite sushi types
(random choice sets in survey)

Dataset	Choices	Features	Largest Choice Set
DISTRICT	5376	27	2
DISTRICT-SMART	5376	6	2
SUSHI	5000	6	10
EXPEDIA	276593	5	38
CAR-A	2675	4	2
CAR-B	2206	5	2
CAR-ALT	4654	21	6

Choice datasets

favorite sushi types
(random choice sets in survey)

hotel bookings
(choice sets = search results)

Dataset	Choices	Features	Largest Choice Set
DISTRICT	5376	27	2
DISTRICT-SMART	5376	6	2
SUSHI	5000	6	10
EXPEDIA	276593	5	38
CAR-A	2675	4	2
CAR-B	2206	5	2
CAR-ALT	4654	21	6

Choice datasets

favorite sushi types
(random choice sets in survey)

hotel bookings
(choice sets = search results)

car purchasing survey
(handcrafted choice sets)

Dataset	Choices	Features	Largest Choice Set
DISTRICT	5376	27	2
DISTRICT-SMART	5376	6	2
SUSHI	5000	6	10
EXPEDIA	276593	5	38
CAR-A	2675	4	2
CAR-B	2206	5	2
CAR-ALT	4654	21	6

LCL improves model fit

whole-dataset negative log-likelihood (lower = better)

	MNL	LCL	Mixed logit	DLCL
DISTRICT	3313	3130	3258	3206
DISTRICT-SMART	3426	3278*	3351	3303 [†]
EXPEDIA	839505	837649*	839055	837569[†]
SUSHI	9821	9773*	9793	9764
CAR-A	1702	1694	1696	1692
CAR-B	1305	1295	1297	1284
CAR-ALT	7393	6733*	7301	7011 [†]

*significant likelihood ratio test vs MNL ($p < 0.001$)

†significant likelihood ratio test vs mixed logit ($p < 0.001$)

LCL can improve out-of-sample prediction performance

mean (std. dev) relative rank (0 = perfect prediction, 1 = always wrong)

	MNL	LCL	Mixed logit	DLCL
DISTRICT	.3680 (.4823)	.3253 (.4685)	.3188 (.4660)	.3225 (.4674)
DISTRICT-SMART	.4006 (.4900)	.3764 (.4845)	.3271 (.4692)	.3448 (.4753)
EXPEDIA	.3859 (.2954)	.3800* (.2945)	.3201 (.2825)	.3195[†] (.2823)
SUSHI	.2727 (.2751)	.2737 (.2781)	.2724 (.2765)	.2732 (.2765)
CAR-A	.3570 (.4791)	.3570 (.4791)	.3570 (.4791)	.3570 (.4791)
CAR-B	.3326 (.4711)	.3213 (.4670)	.3303 (.4703)	.3235 (.4678)
CAR-ALT	.2944 (.2875)	.2661* (.2819)	.2931 (.2966)	.2798 (.2837)

*significant Wilcoxon test vs MNL ($p < 0.001$)

LCL can test individual effects for significance

LCL can test individual effects for significance

constrain all but one entry of A to 0

LCL can test individual effects for significance

constrain all but one entry of A to 0
→ loss still convex

LCL can test individual effects for significance

constrain all but one entry of A to 0

- loss still convex
- test *strength* of single effect

LCL can test individual effects for significance

constrain all but one entry of A to 0

- loss still convex
- test *strength* of single effect
- constrained model nests cond. logit

LCL can test individual effects for significance

constrain all but one entry of A to 0

- loss still convex
- test *strength* of single effect
- constrained model nests cond. logit

\bar{A}_{pq} : learned param in constrained model

LCL can test individual effects for significance

- constrain all but one entry of A to 0
- loss still convex
- test *strength* of single effect
- constrained model nests cond. logit

\bar{A}_{pq} : learned param in constrained model

Table 5: Five largest context effects in SUSHI.

Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>popularity</i> on <i>popularity</i>	-0.28	-0.11	3.0	0.081
<i>availability</i> on <i>is maki</i>	0.24	0.04	0.39	0.53
<i>oiliness</i> on <i>oiliness</i>	-0.20	-0.26	23	1.5×10^{-6}
<i>popularity</i> on <i>availability</i>	0.19	0.09	2.3	0.13
<i>availability</i> on <i>oiliness</i>	-0.18	-0.04	0.74	0.39

LCL can test individual effects for significance

- constrain all but one entry of A to 0
- loss still convex
- test *strength* of single effect
- constrained model nests cond. logit

\bar{A}_{pq} : learned param in constrained model

Table 5: Five largest context effects in SUSHI.

Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>popularity</i> on <i>popularity</i>	-0.28	-0.11	3.0	0.081
<i>availability</i> on <i>is maki</i>	0.24	0.04	0.39	0.53
<i>oiliness</i> on <i>oiliness</i>	-0.20	-0.26	23	1.5×10^{-6}
<i>popularity</i> on <i>availability</i>	0.19	0.09	2.3	0.13
<i>availability</i> on <i>oiliness</i>	-0.18	-0.04	0.74	0.39

Table 6: Five largest context effects in EXPEDIA.

Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>location score</i> on <i>price</i>	-0.47	-0.13	10	0.002
<i>on promotion</i> on <i>price</i>	0.27	0.13	17	3.5×10^{-5}
<i>review score</i> on <i>price</i>	-0.19	-0.13	29	8.6×10^{-8}
<i>star rating</i> on <i>price</i>	0.15	0.20	65	6.0×10^{-16}
<i>price</i> on <i>star rating</i>	0.10	0.03	4.0	0.046

LCL can test individual effects for significance

- constrain all but one entry of A to 0
- loss still convex
- test *strength* of single effect
- constrained model nests cond. logit

\bar{A}_{pq} : learned param in constrained model

Table 5: Five largest context effects in SUSHI.

Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>popularity</i> on <i>popularity</i>	-0.28	-0.11	3.0	0.081
<i>availability</i> on <i>is maki</i>	0.24	0.04	0.39	0.53
<i>oiliness</i> on <i>oiliness</i>	-0.20	-0.26	23	1.5×10^{-6}
<i>popularity</i> on <i>availability</i>	0.19	0.09	2.3	0.13
<i>availability</i> on <i>oiliness</i>	-0.18	-0.04	0.74	0.39

Table 6: Five largest context effects in EXPEDIA.

Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>location score</i> on <i>price</i>	-0.47	-0.13	10	0.002
<i>on promotion</i> on <i>price</i>	0.27	0.13	17	3.5×10^{-5}
<i>review score</i> on <i>price</i>	-0.19	-0.13	29	8.6×10^{-8}
<i>star rating</i> on <i>price</i>	0.15	0.20	65	6.0×10^{-16}
<i>price</i> on <i>star rating</i>	0.10	0.03	4.0	0.046

Table 7: Five largest context effects in CAR-ALT.

Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>truck</i> on <i>truck</i>	1.06	0.83	239	$< 10^{-16}$
<i>van</i> on <i>van</i>	0.94	0.97	309	$< 10^{-16}$
<i>suv</i> on <i>station wagon</i>	0.89	0.98	-0.21	1.0
<i>station wagon</i> on <i>station wagon</i>	0.88	0.93	153	$< 10^{-16}$
<i>sports car</i> on <i>station wagon</i>	0.86	0.96	-0.21	1.0

LCL can test individual effects for significance

constrain all but one entry of A to 0
 → loss still convex
 → test *strength* of single effect
 → constrained model nests cond. logit

\bar{A}_{pq} : learned param in constrained model

are these effects *causal*?

Table 5: Five largest context effects in SUSHI.

Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>popularity</i> on <i>popularity</i>	-0.28	-0.11	3.0	0.081
<i>availability</i> on <i>is maki</i>	0.24	0.04	0.39	0.53
<i>oiliness</i> on <i>oiliness</i>	-0.20	-0.26	23	1.5×10^{-6}
<i>popularity</i> on <i>availability</i>	0.19	0.09	2.3	0.13
<i>availability</i> on <i>oiliness</i>	-0.18	-0.04	0.74	0.39

Table 6: Five largest context effects in EXPEDIA.

Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>location score</i> on <i>price</i>	-0.47	-0.13	10	0.002
<i>on promotion</i> on <i>price</i>	0.27	0.13	17	3.5×10^{-5}
<i>review score</i> on <i>price</i>	-0.19	-0.13	29	8.6×10^{-8}
<i>star rating</i> on <i>price</i>	0.15	0.20	65	6.0×10^{-16}
<i>price</i> on <i>star rating</i>	0.10	0.03	4.0	0.046

Table 7: Five largest context effects in CAR-ALT.

Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>truck</i> on <i>truck</i>	1.06	0.83	239	$< 10^{-16}$
<i>van</i> on <i>van</i>	0.94	0.97	309	$< 10^{-16}$
<i>suv</i> on <i>station wagon</i>	0.89	0.98	-0.21	1.0
<i>station wagon</i> on <i>station wagon</i>	0.88	0.93	153	$< 10^{-16}$
<i>sports car</i> on <i>station wagon</i>	0.86	0.96	-0.21	1.0

LCL can test individual effects for significance

constrain all but one entry of A to 0
 → loss still convex
 → test *strength* of single effect
 → constrained model nests cond. logit

\bar{A}_{pq} : learned param in constrained model

are these effects *causal*?

→ *choice set assignment* (stay tuned)

Table 5: Five largest context effects in SUSHI.

Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>popularity</i> on <i>popularity</i>	-0.28	-0.11	3.0	0.081
<i>availability</i> on <i>is maki</i>	0.24	0.04	0.39	0.53
<i>oiliness</i> on <i>oiliness</i>	-0.20	-0.26	23	1.5×10^{-6}
<i>popularity</i> on <i>availability</i>	0.19	0.09	2.3	0.13
<i>availability</i> on <i>oiliness</i>	-0.18	-0.04	0.74	0.39

Table 6: Five largest context effects in EXPEDIA.

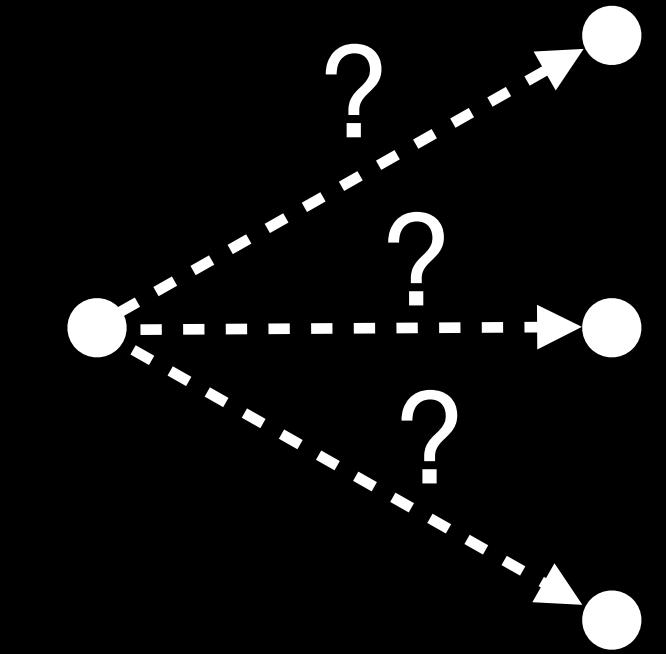
Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>location score</i> on <i>price</i>	-0.47	-0.13	10	0.002
<i>on promotion</i> on <i>price</i>	0.27	0.13	17	3.5×10^{-5}
<i>review score</i> on <i>price</i>	-0.19	-0.13	29	8.6×10^{-8}
<i>star rating</i> on <i>price</i>	0.15	0.20	65	6.0×10^{-16}
<i>price</i> on <i>star rating</i>	0.10	0.03	4.0	0.046

Table 7: Five largest context effects in CAR-ALT.

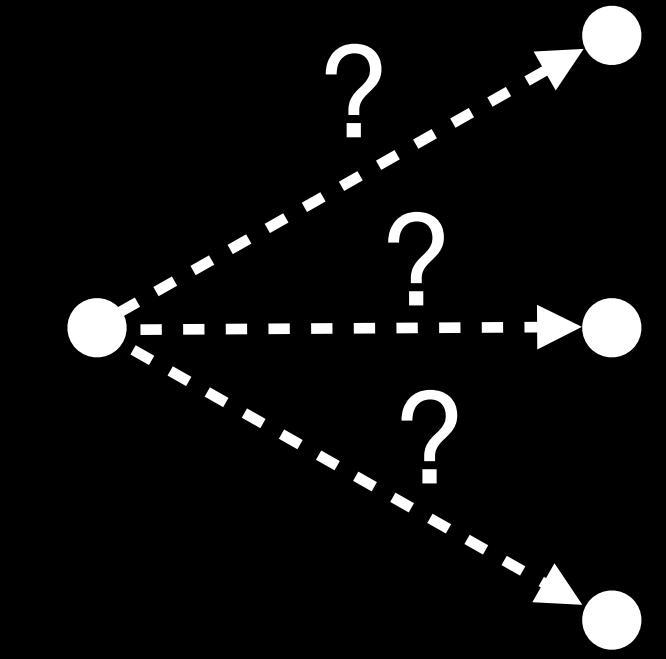
Effect (q on p)	A_{pq}	\bar{A}_{pq}	LRT	p -value
<i>truck</i> on <i>truck</i>	1.06	0.83	239	$< 10^{-16}$
<i>van</i> on <i>van</i>	0.94	0.97	309	$< 10^{-16}$
<i>suv</i> on <i>station wagon</i>	0.89	0.98	-0.21	1.0
<i>station wagon</i> on <i>station wagon</i>	0.88	0.93	153	$< 10^{-16}$
<i>sports car</i> on <i>station wagon</i>	0.86	0.96	-0.21	1.0

6. Social network application

What factors drive edge formation?

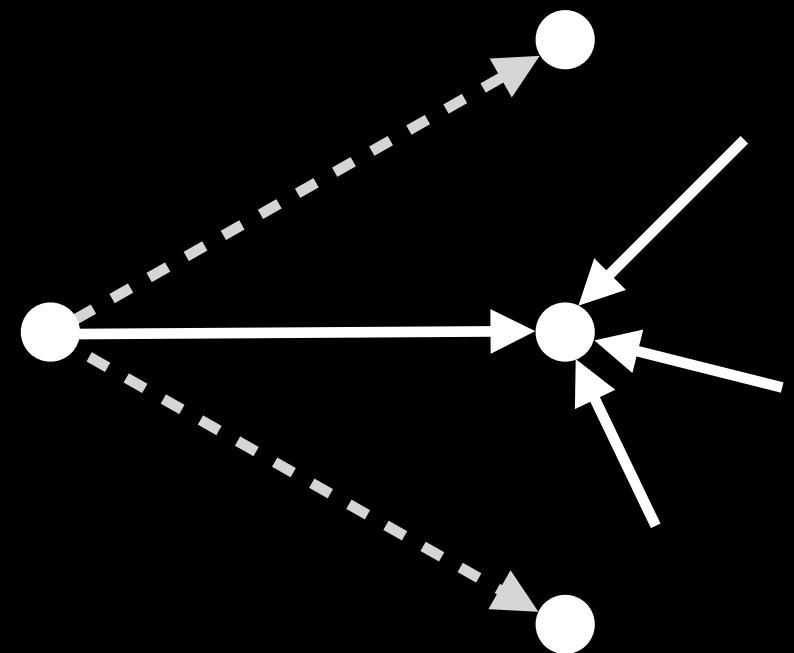


What factors drive edge formation?



Preferential attachment

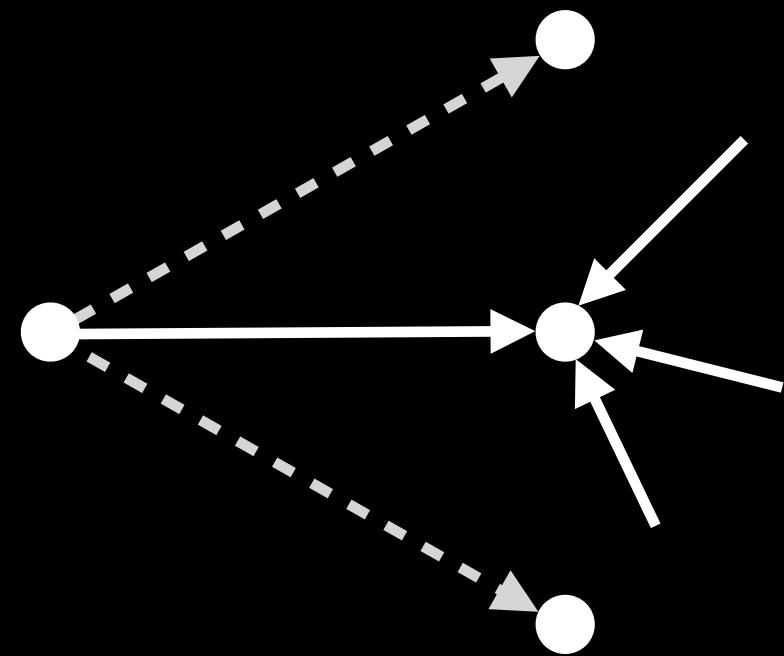
(Barabási & Albert, *Science* 1999)



What factors drive edge formation?

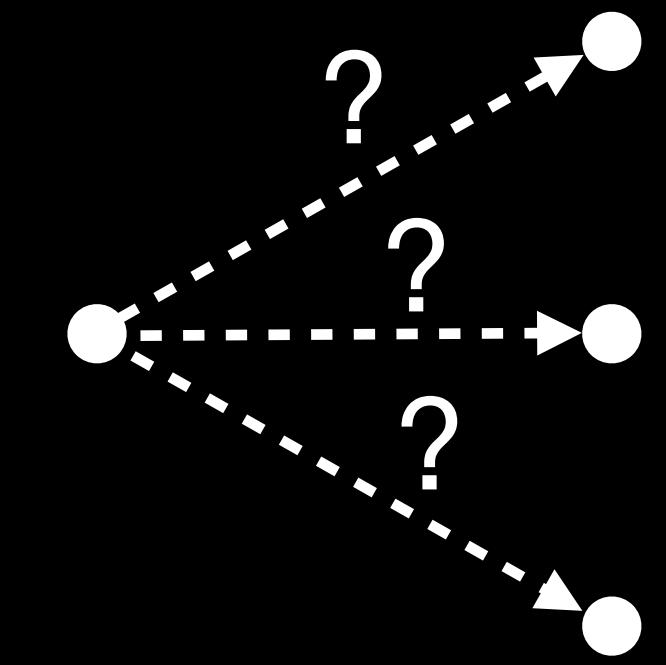
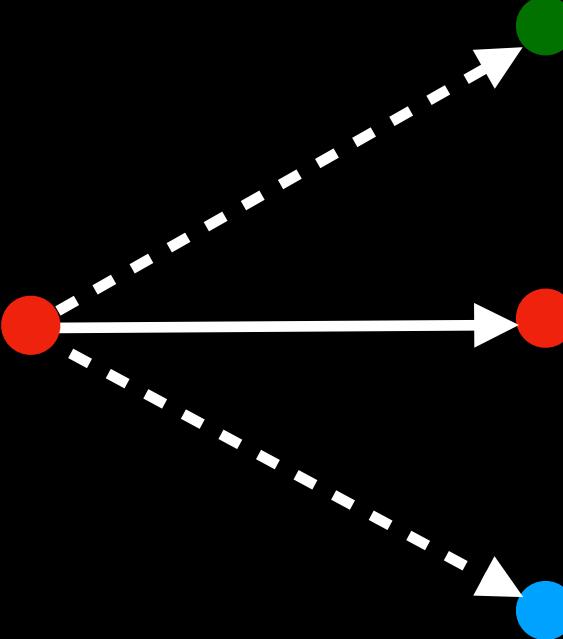
Preferential attachment

(Barabási & Albert, *Science* 1999)



Homophily

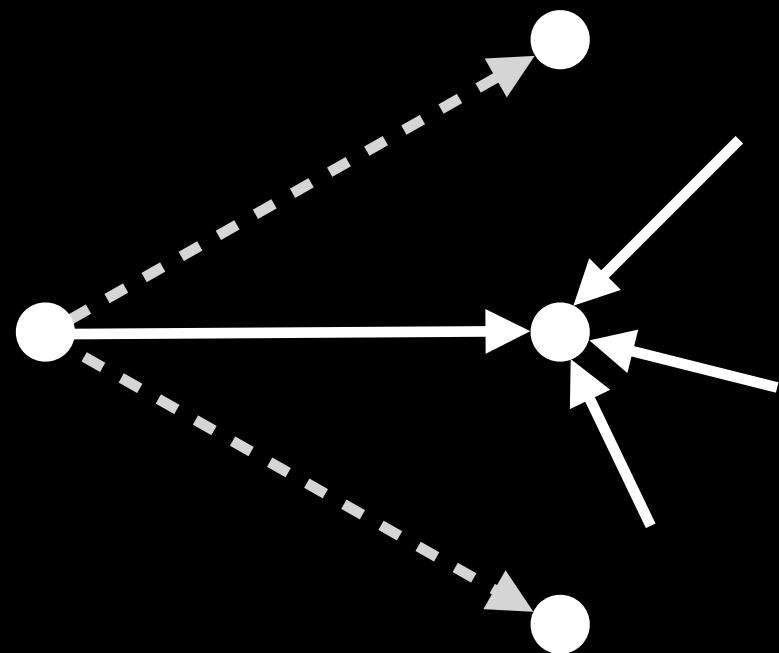
(McPherson et al., *Annual Review of Sociology* 2001)
(Papadopoulos et al., *Nature* 2012)



What factors drive edge formation?

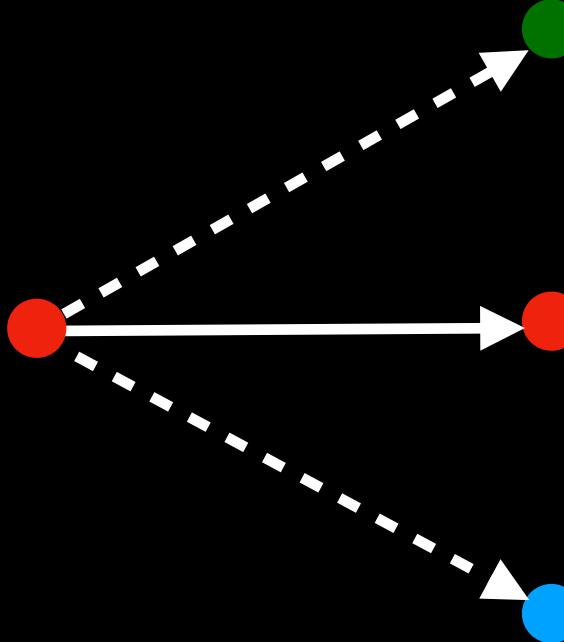
Preferential attachment

(Barabási & Albert, *Science* 1999)



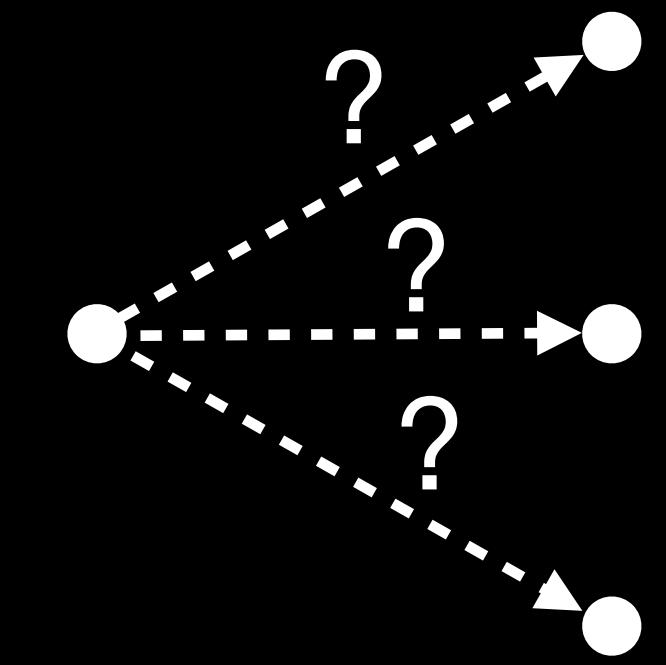
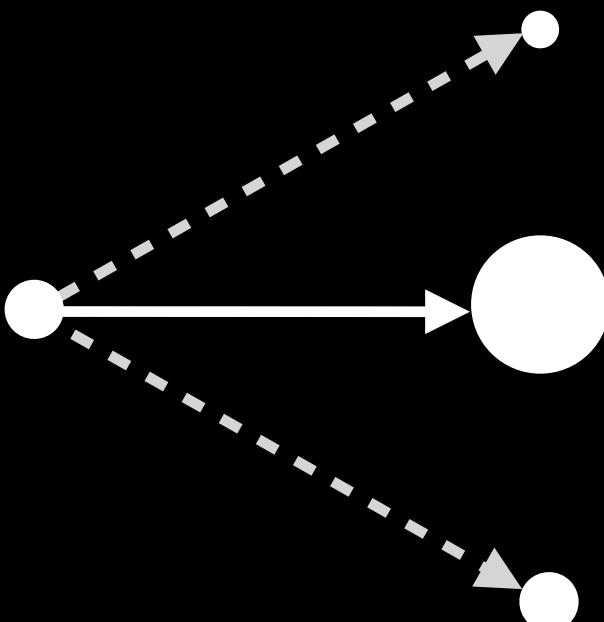
Homophily

(McPherson et al., *Annual Review of Sociology* 2001)
(Papadopoulos et al., *Nature* 2012)



Fitness

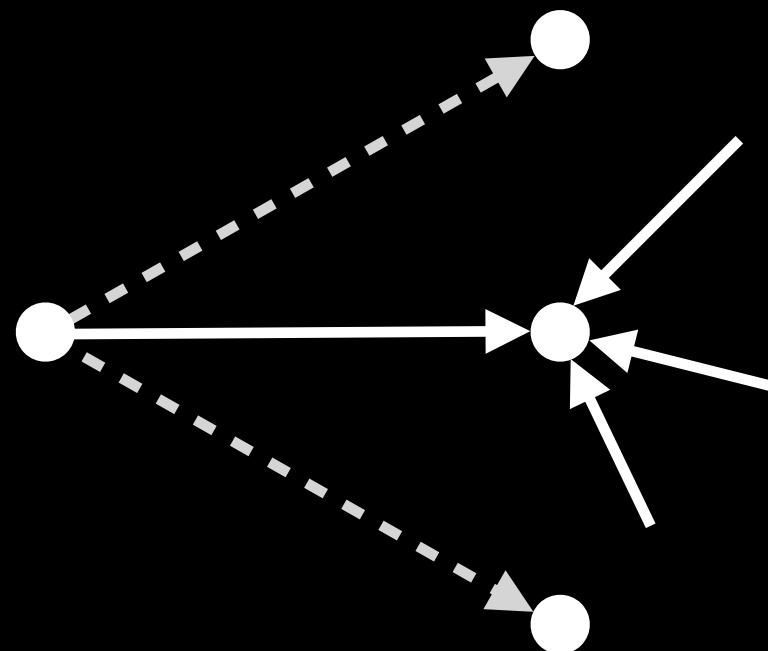
(Bianconi & Barabási, *Europhysics Letters* 2001)
(Caldarelli et al., *Physical Review Letters* 2002)



What factors drive edge formation?

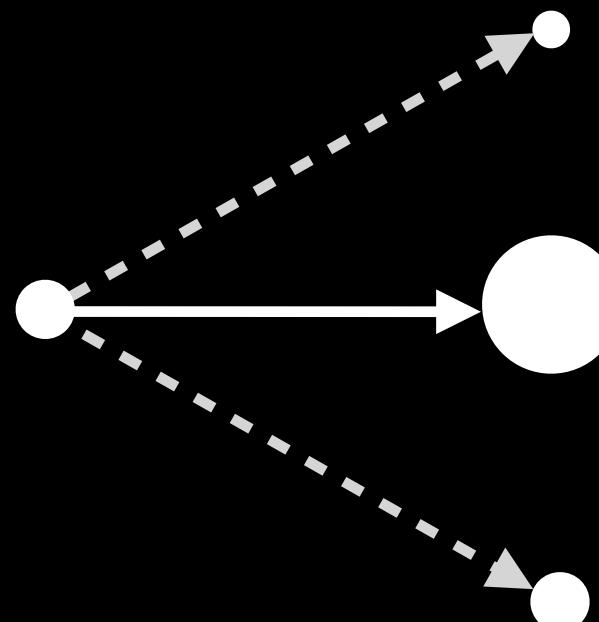
Preferential attachment

(Barabási & Albert, *Science* 1999)



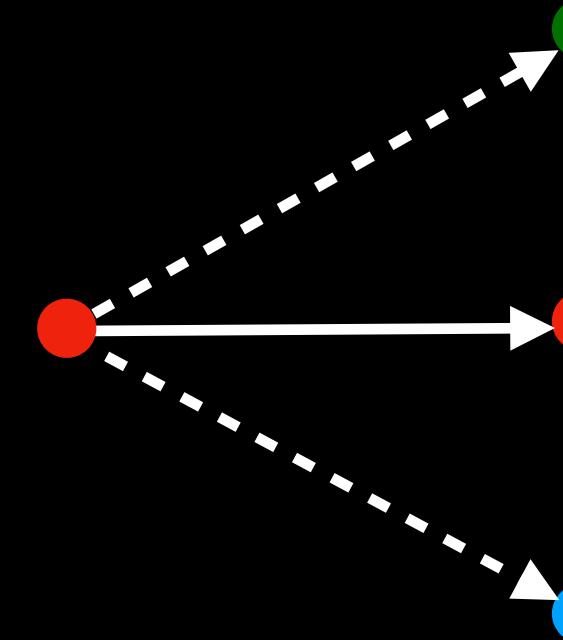
Fitness

(Bianconi & Barabási, *Europhysics Letters* 2001)
(Caldarelli et al., *Physical Review Letters* 2002)



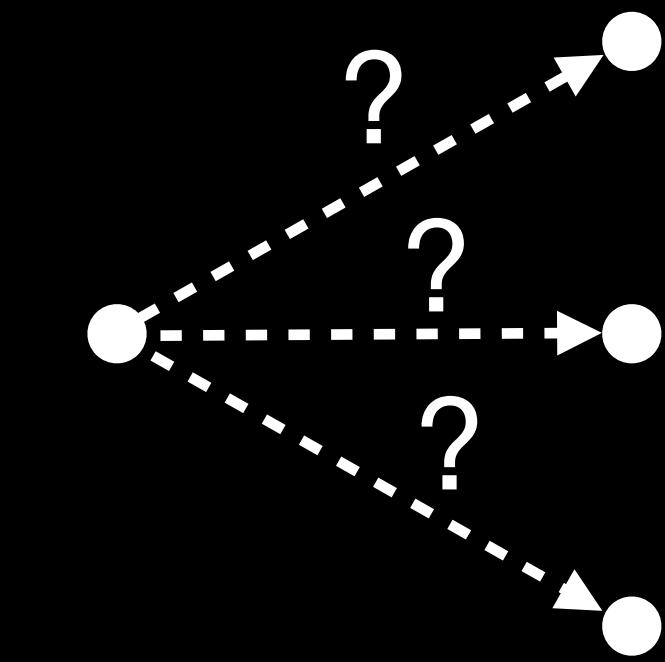
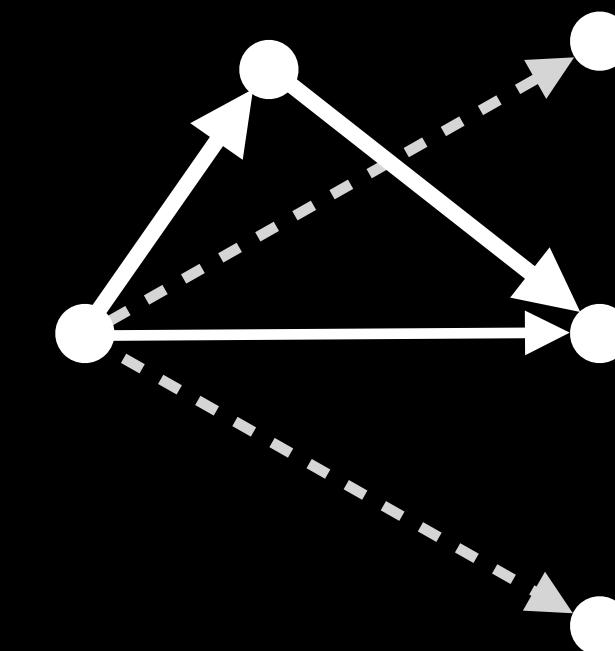
Homophily

(McPherson et al., *Annual Review of Sociology* 2001)
(Papadopoulos et al., *Nature* 2012)



Triadic closure

(Rapoport, *Bulletin of Mathematical Biophysics* 1953)
(Jin et al., *Physical Review E* 2001)



“Choosing to grow a graph”

(Overgoor et al., *SINM '19 & WWW '19*)

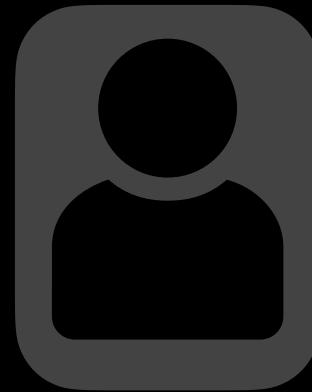
(Gupta & Porter, *arXiv* 2020)

“Choosing to grow a graph”

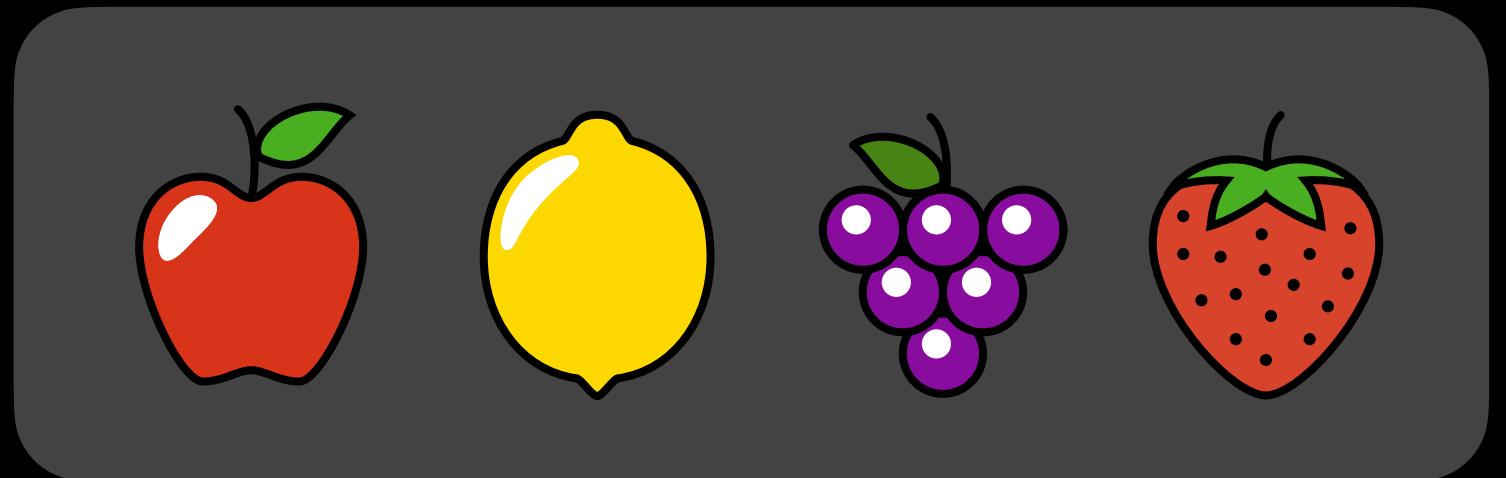
(Overgoor et al., *SINM '19 & WWW '19*)

(Gupta & Porter, *arXiv* 2020)

so far:



chooser



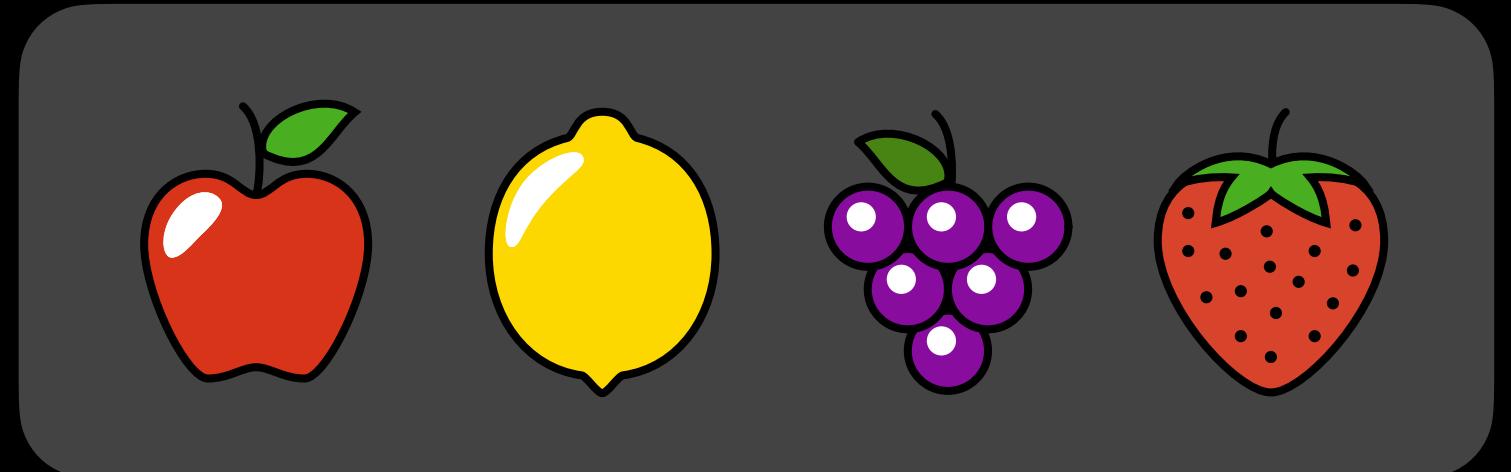
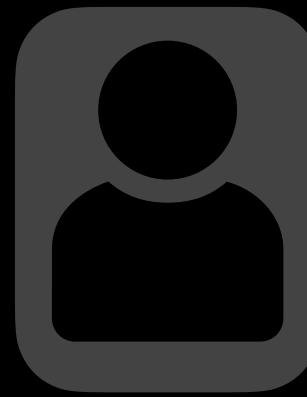
choice set

“Choosing to grow a graph”

(Overgoor et al., *SINM '19 & WWW '19*)

(Gupta & Porter, *arXiv* 2020)

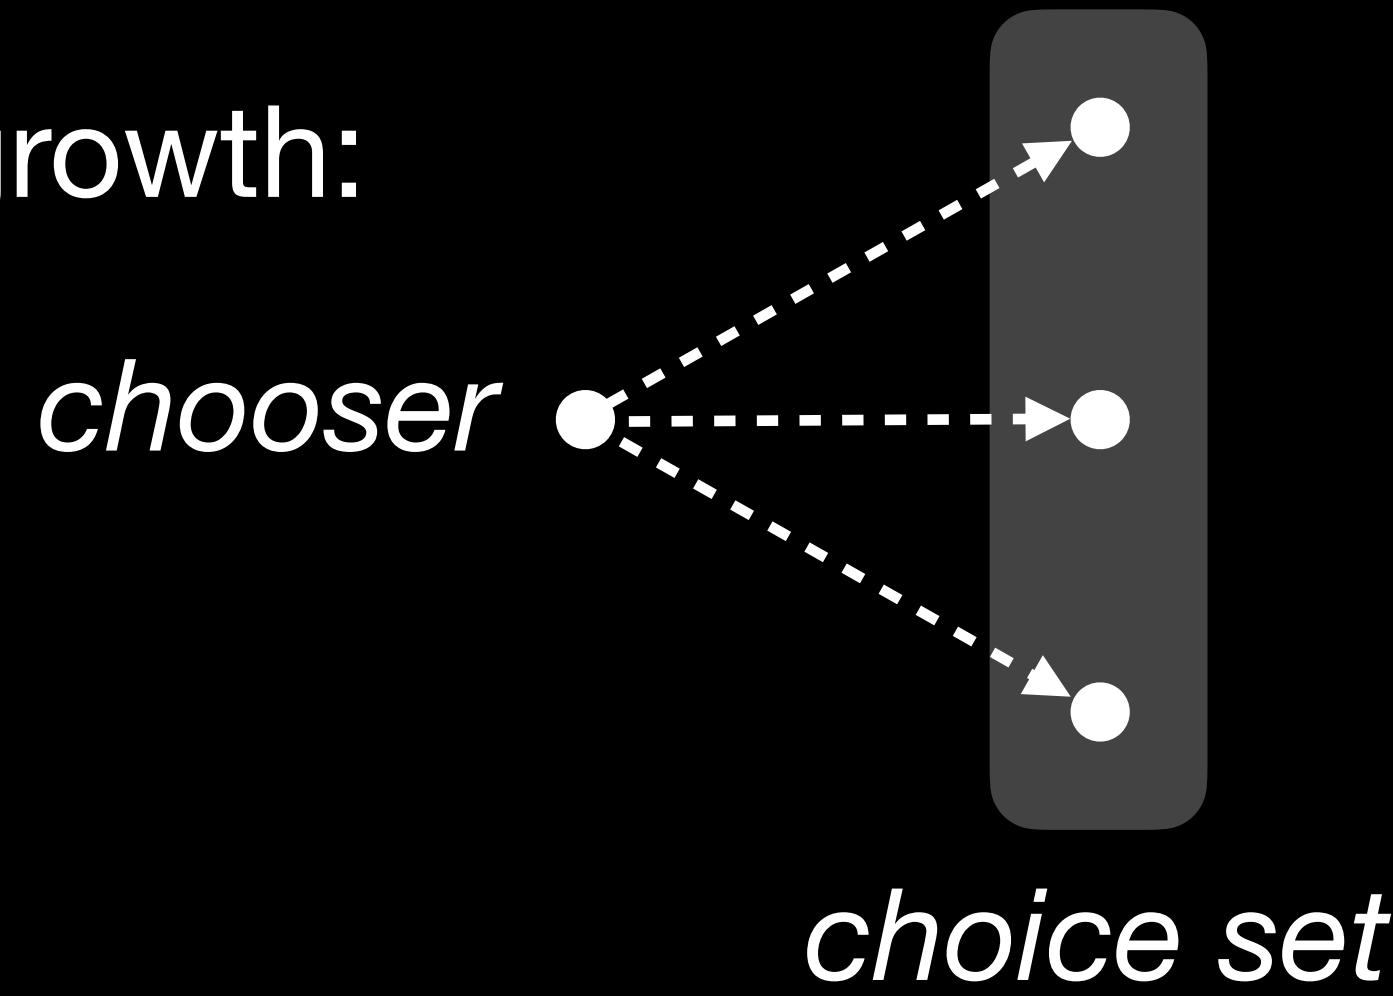
so far:



chooser

choice set

in network growth:



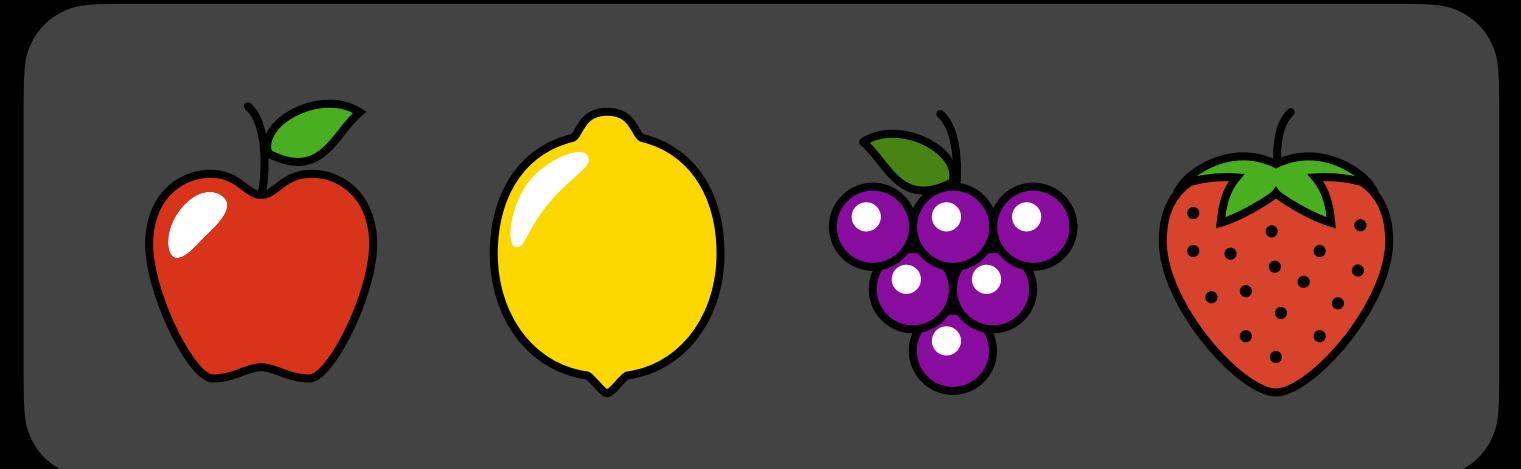
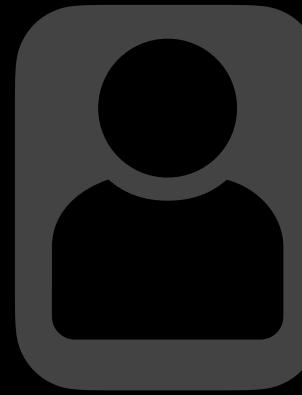
choice set

“Choosing to grow a graph”

(Overgoor et al., S/INM ’19 & WWW ’19)

(Gupta & Porter, arXiv 2020)

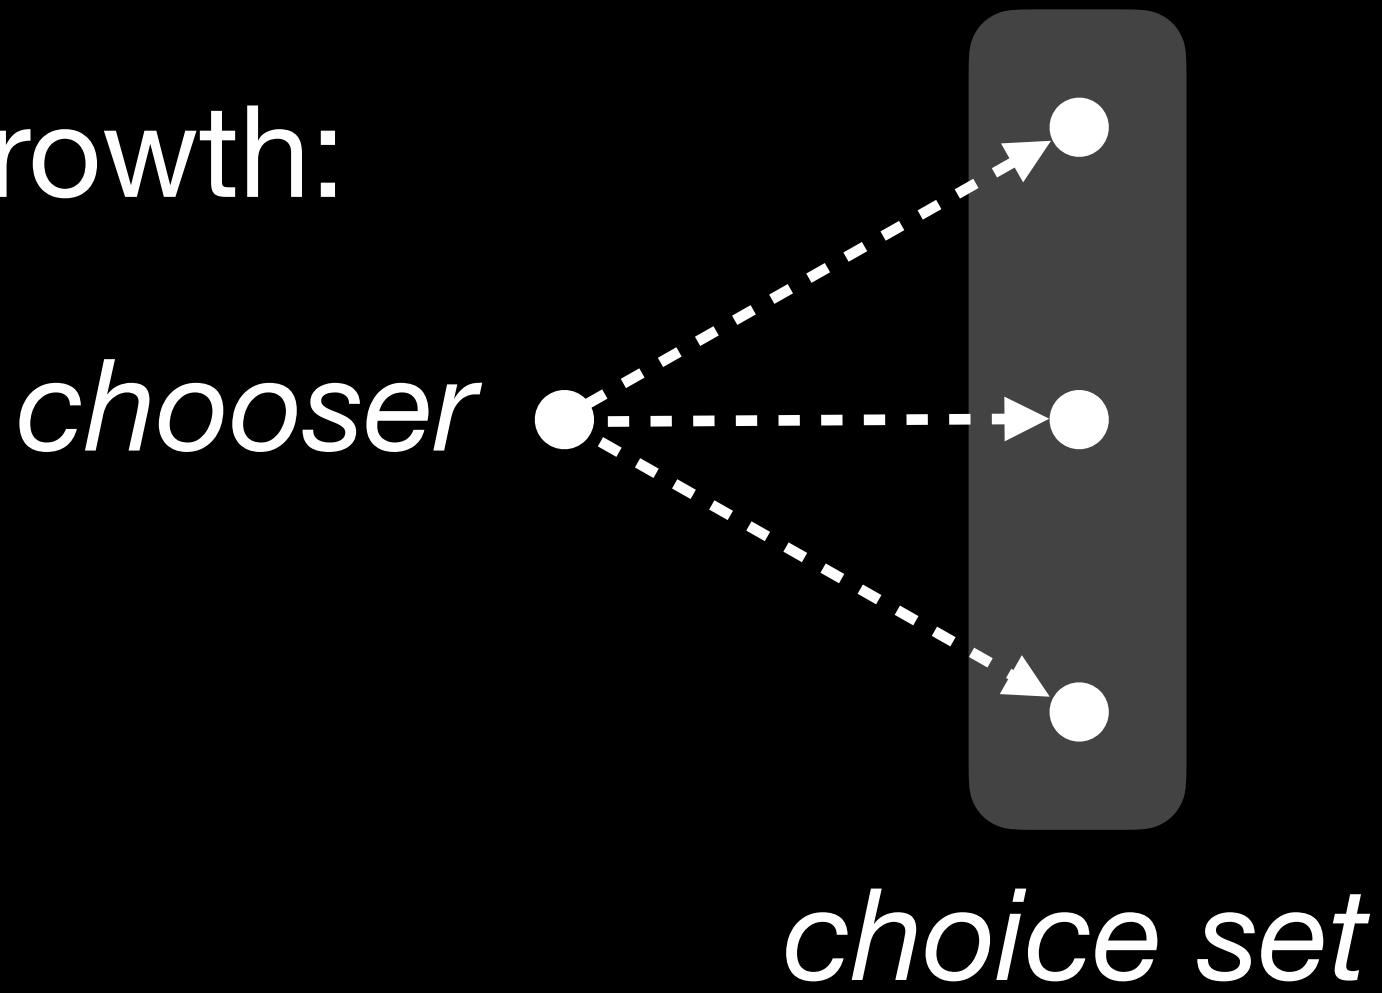
so far:



chooser

choice set

in network growth:



choice set

Key usage

Timestamped edges

→ meaningful choice sets

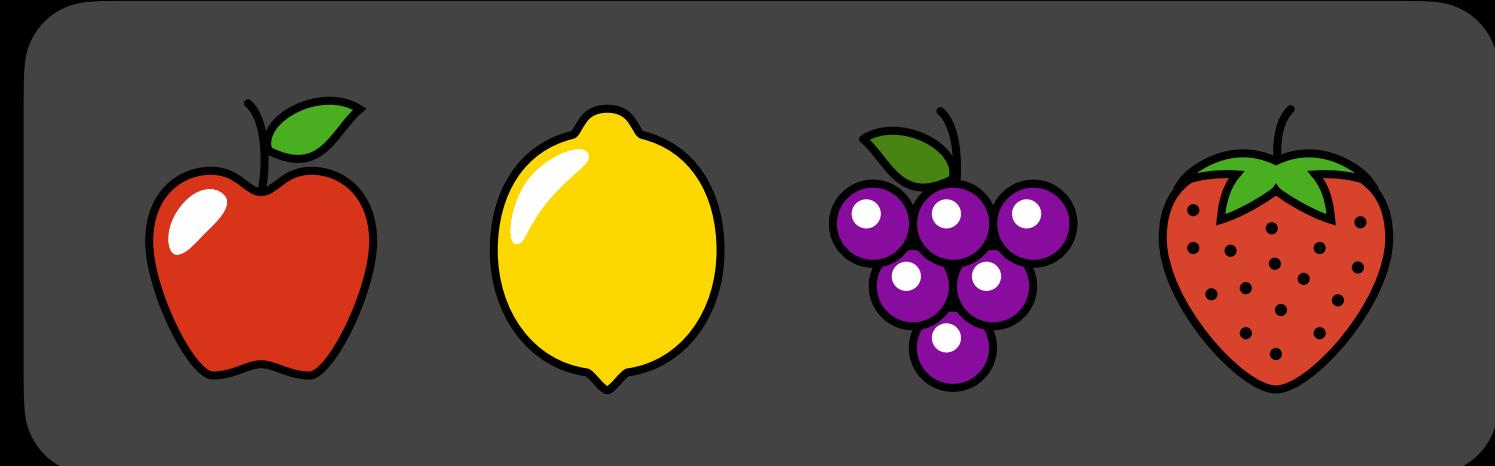
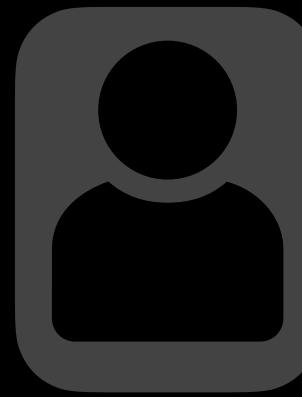
Infer relative importance of edge formation mechanisms from data

“Choosing to grow a graph”

(Overgoor et al., S/INM ’19 & WWW ’19)

(Gupta & Porter, arXiv 2020)

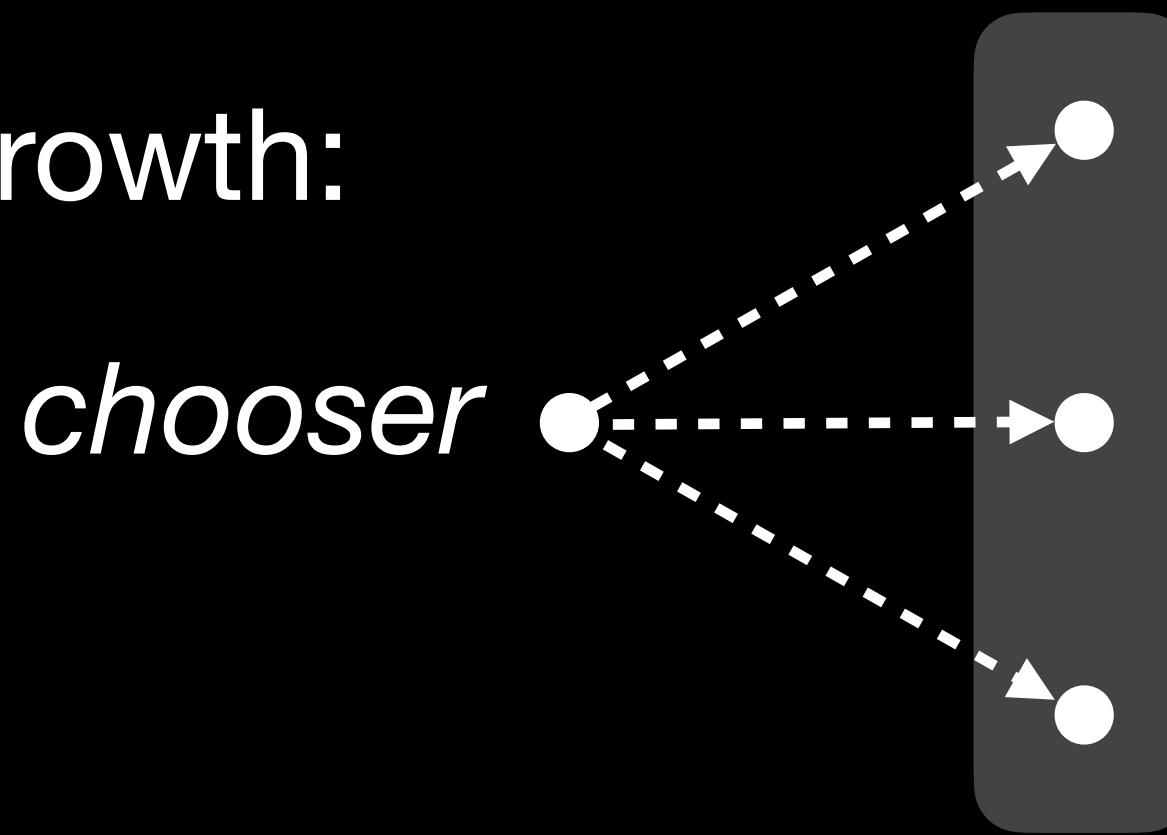
so far:



chooser

choice set

in network growth:



choice set

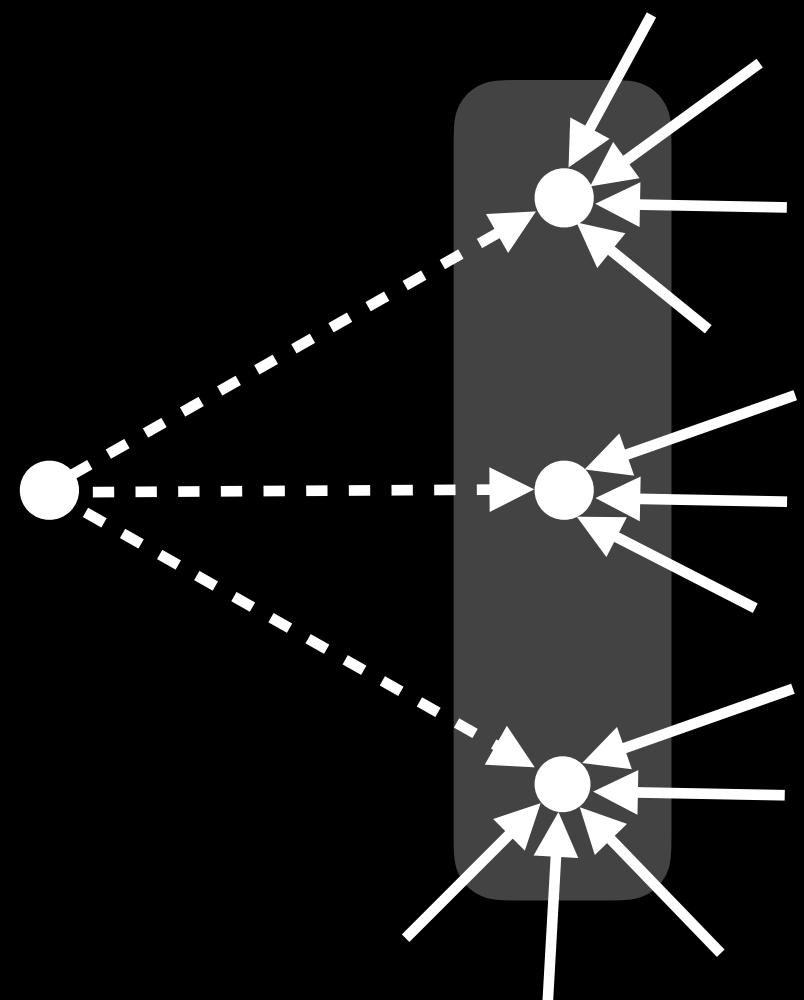
Key usage

Timestamped edges

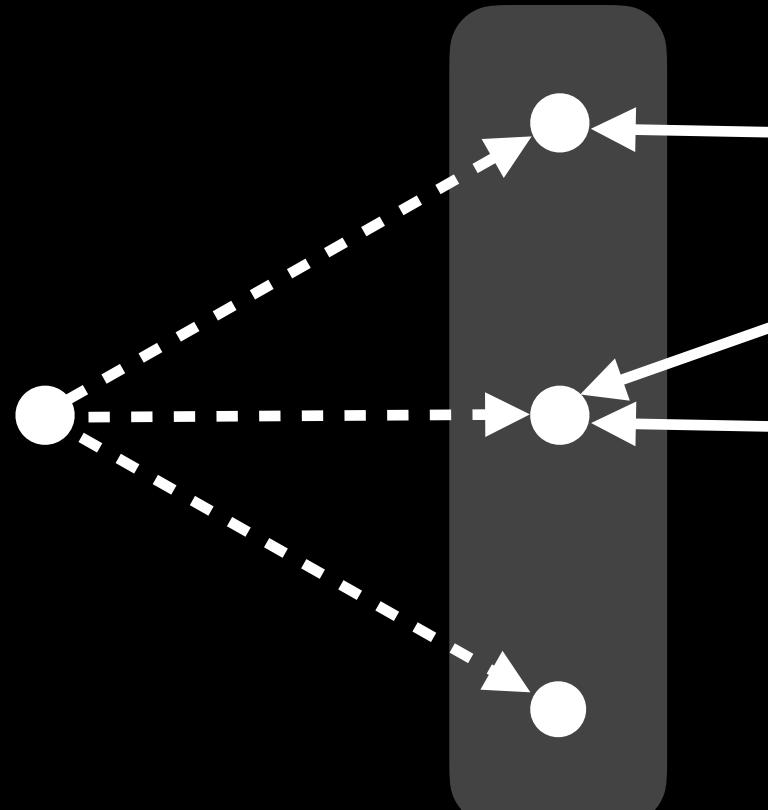
→ meaningful choice sets

Infer relative importance of edge formation mechanisms from data

feature context effects:



vs.



Choosing to close triangles

Triadic closure offers small choice sets

- tractable inference
- varied choice sets

Choosing to close triangles

Triadic closure offers small choice sets

- tractable inference
- varied choice sets

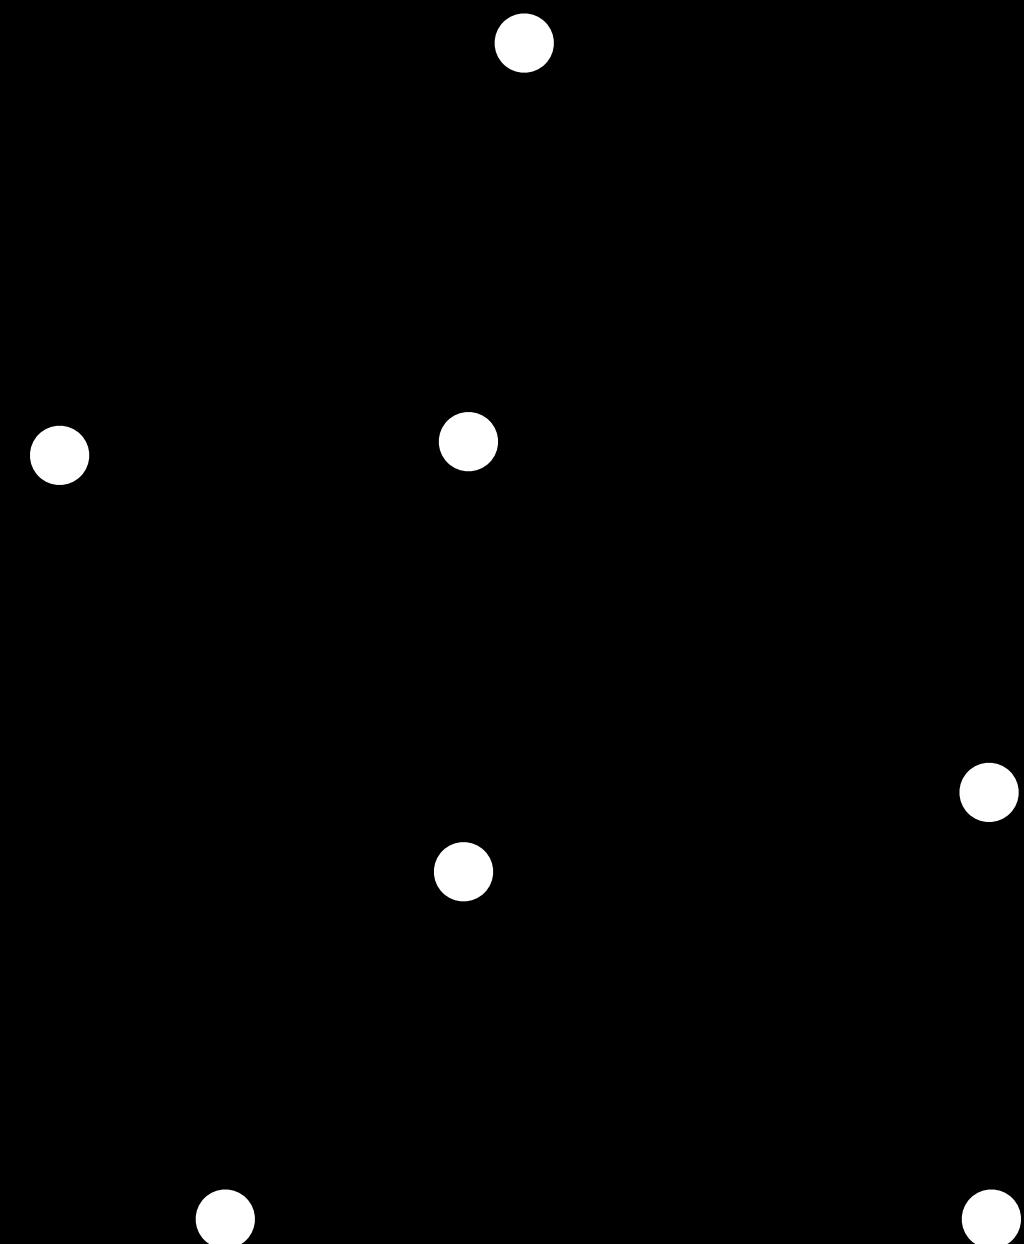
Our data

Timestamped edges
(including repeats)

Choosing to close triangles

Triadic closure offers small choice sets

- tractable inference
- varied choice sets



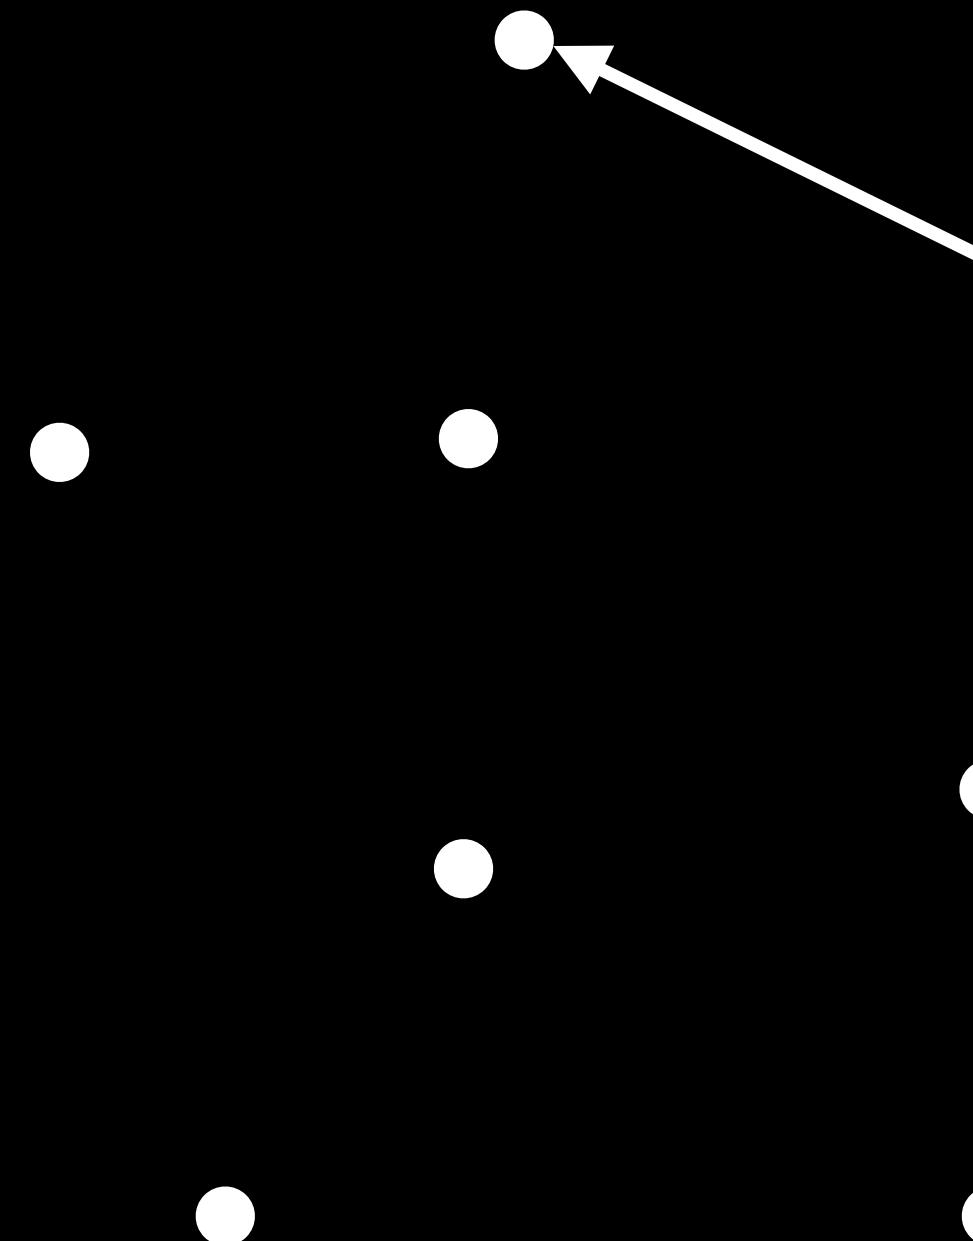
Our data

Timestamped edges
(including repeats)

Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

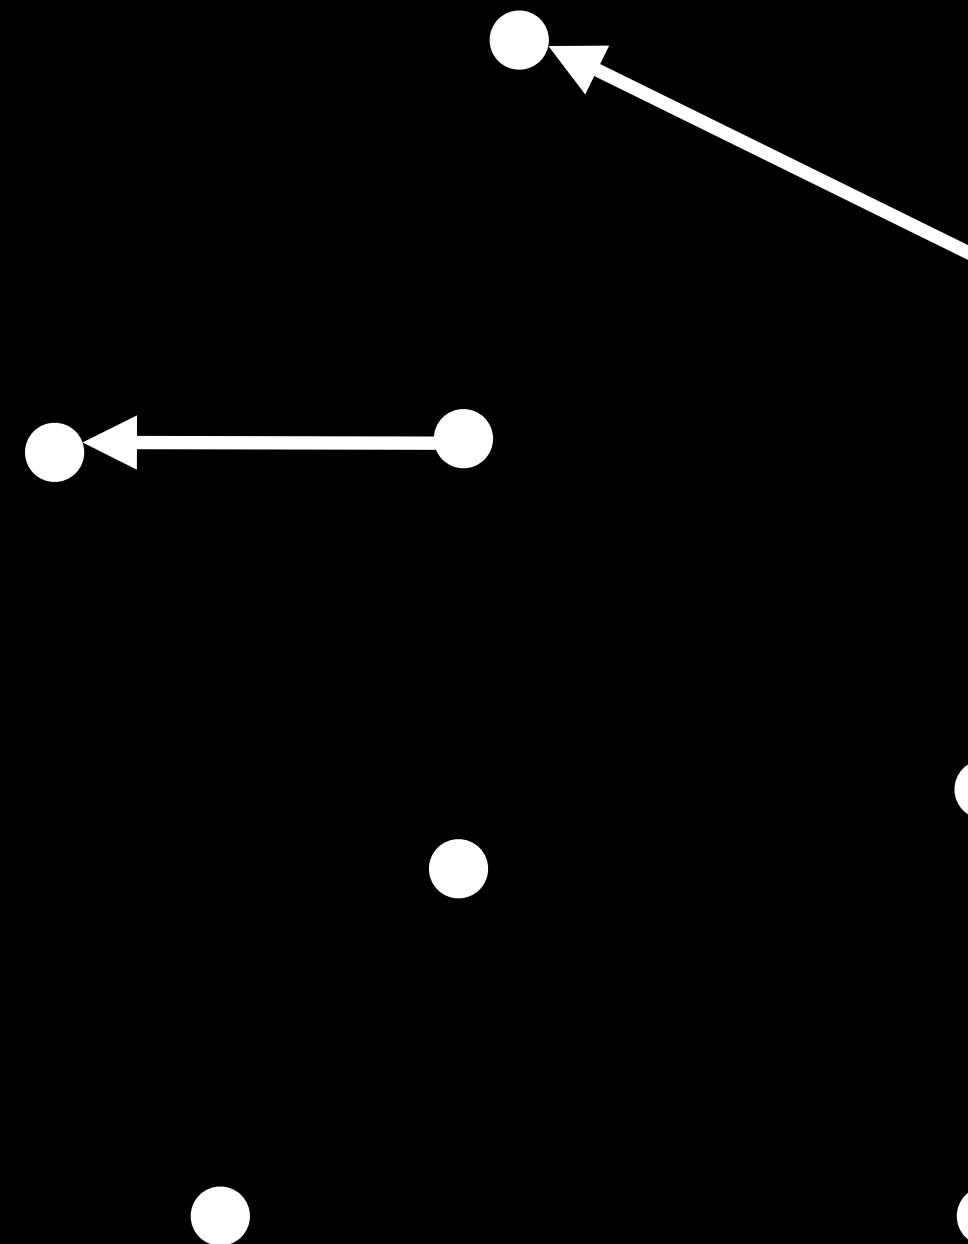
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

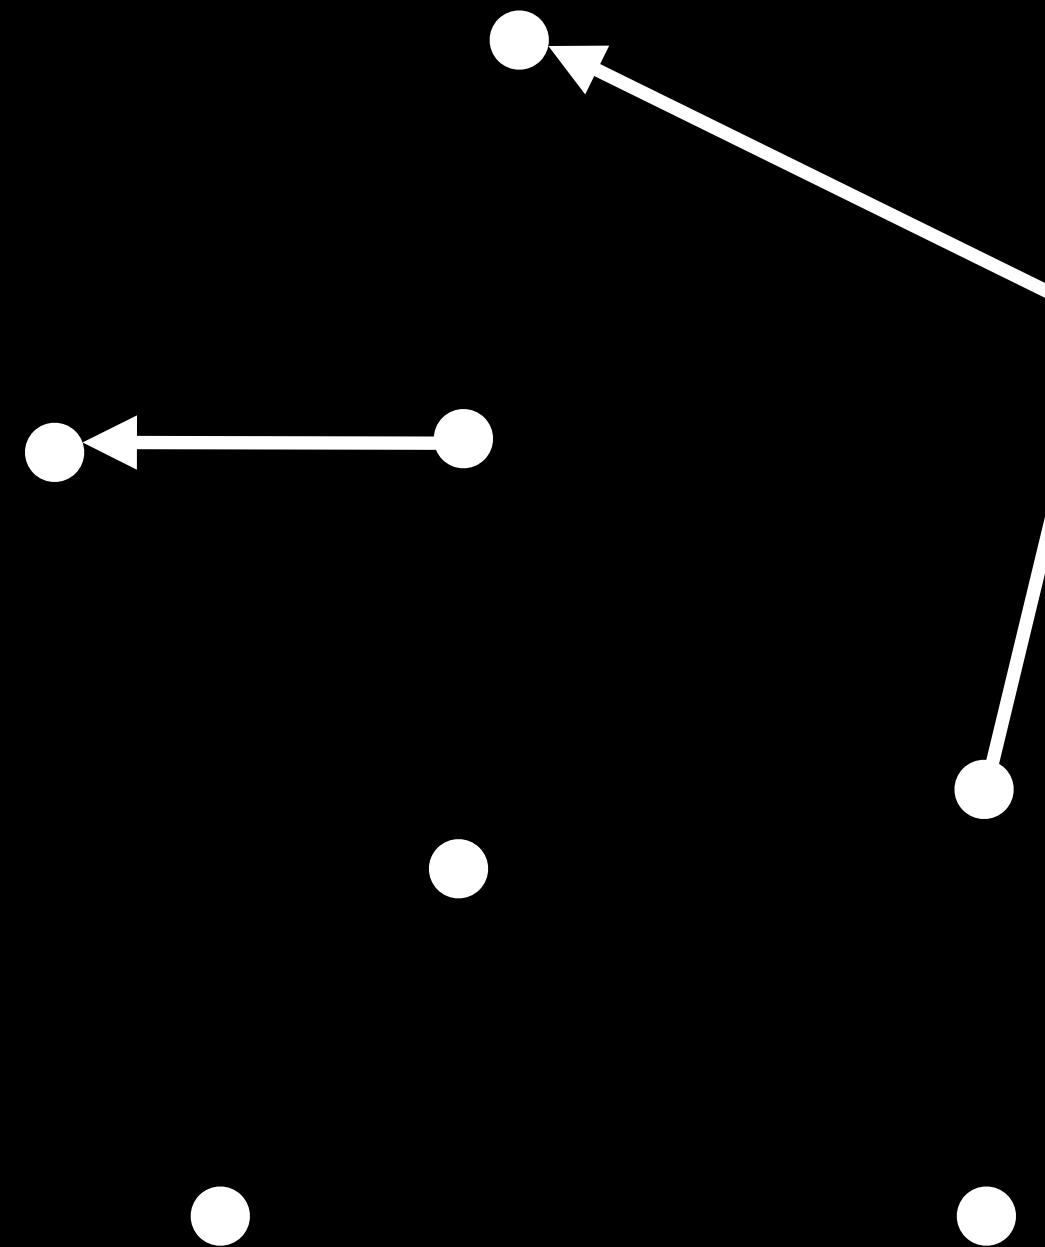
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

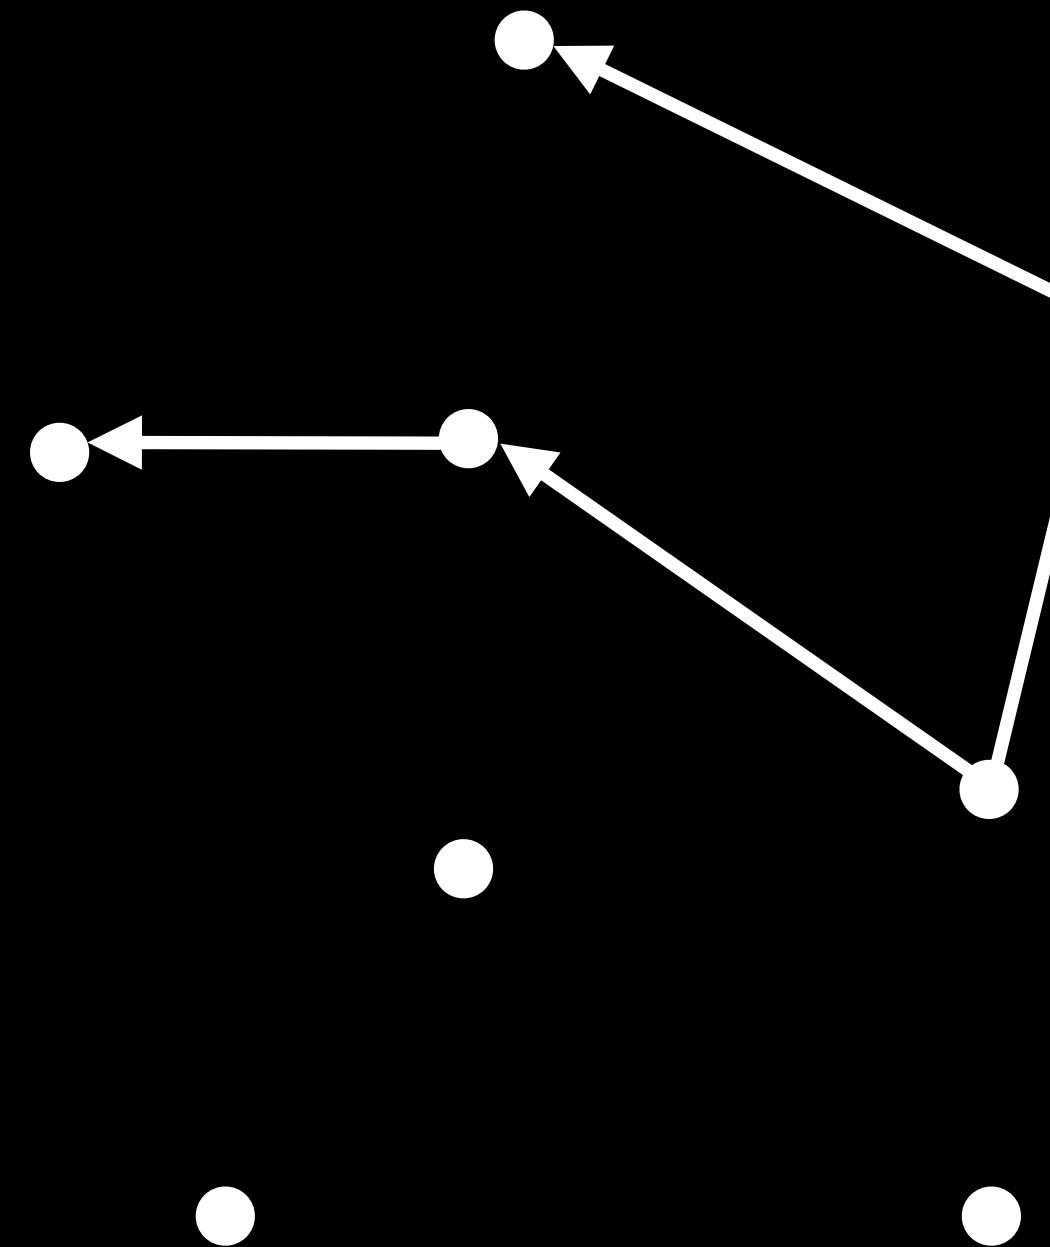
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

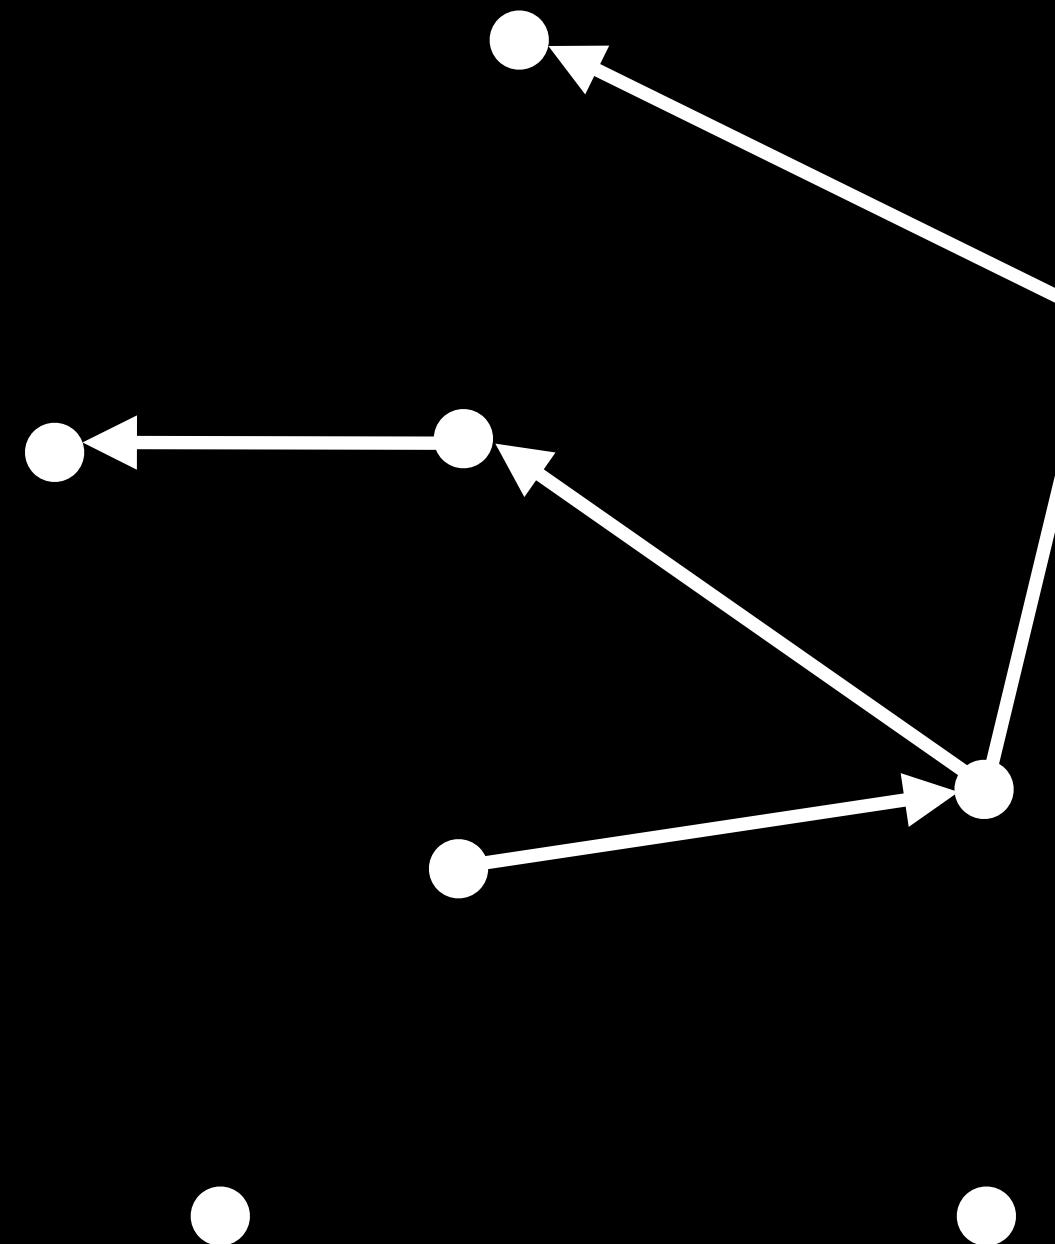
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

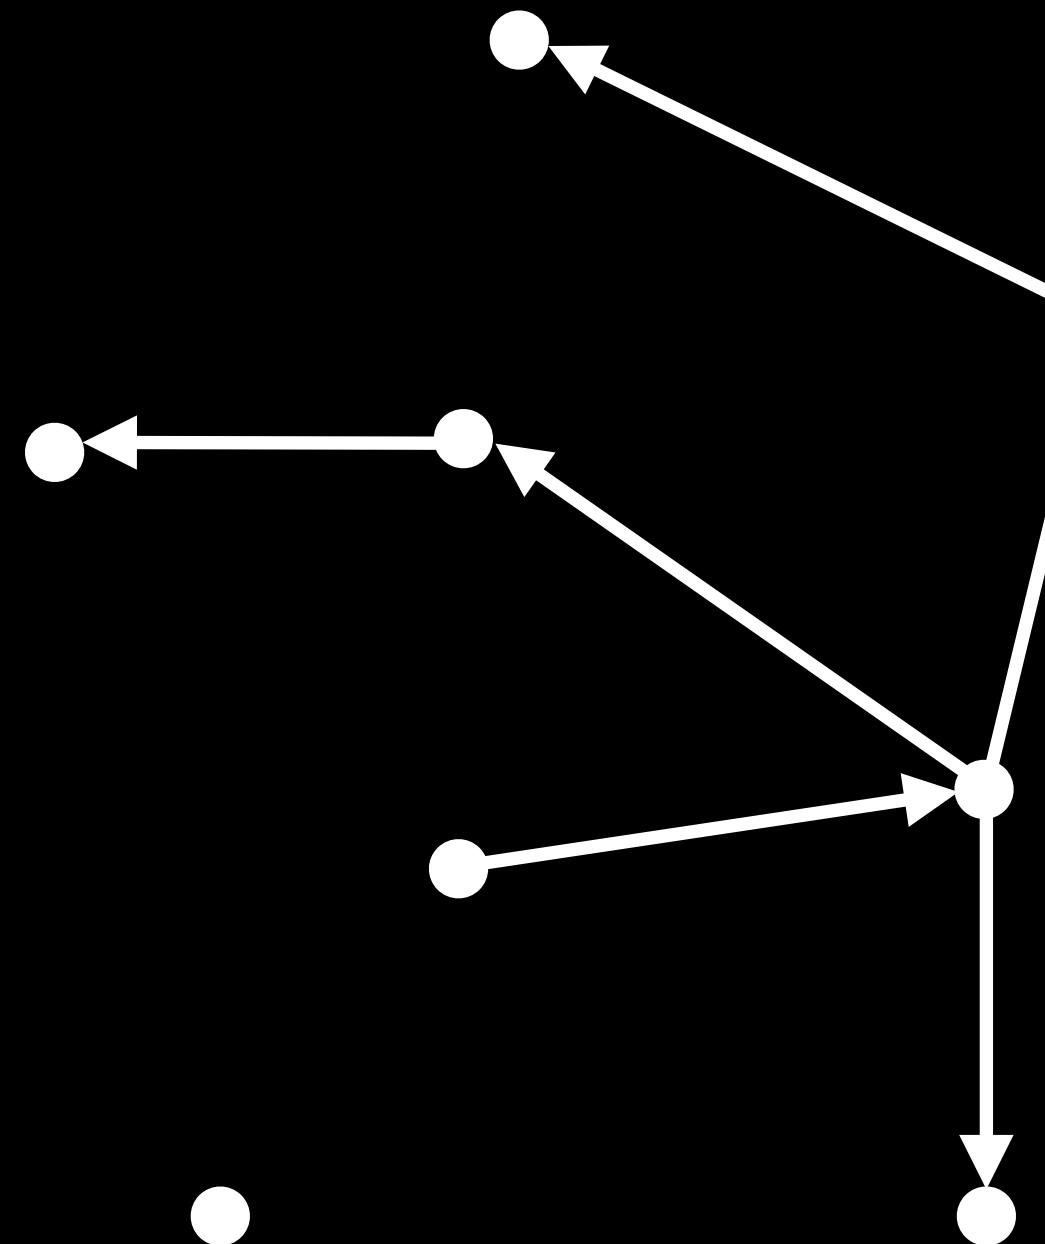
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

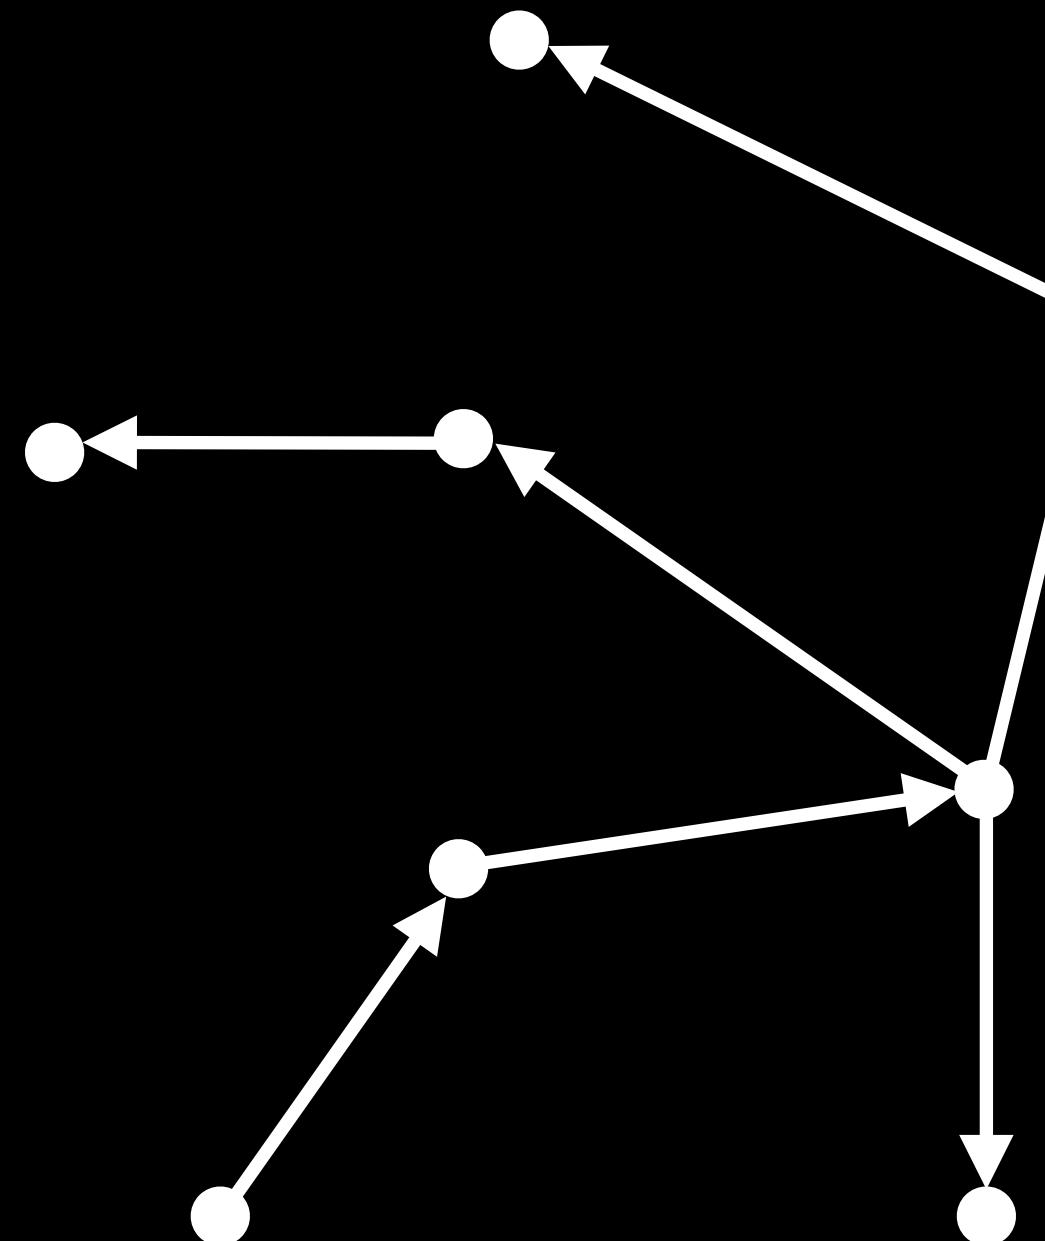
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

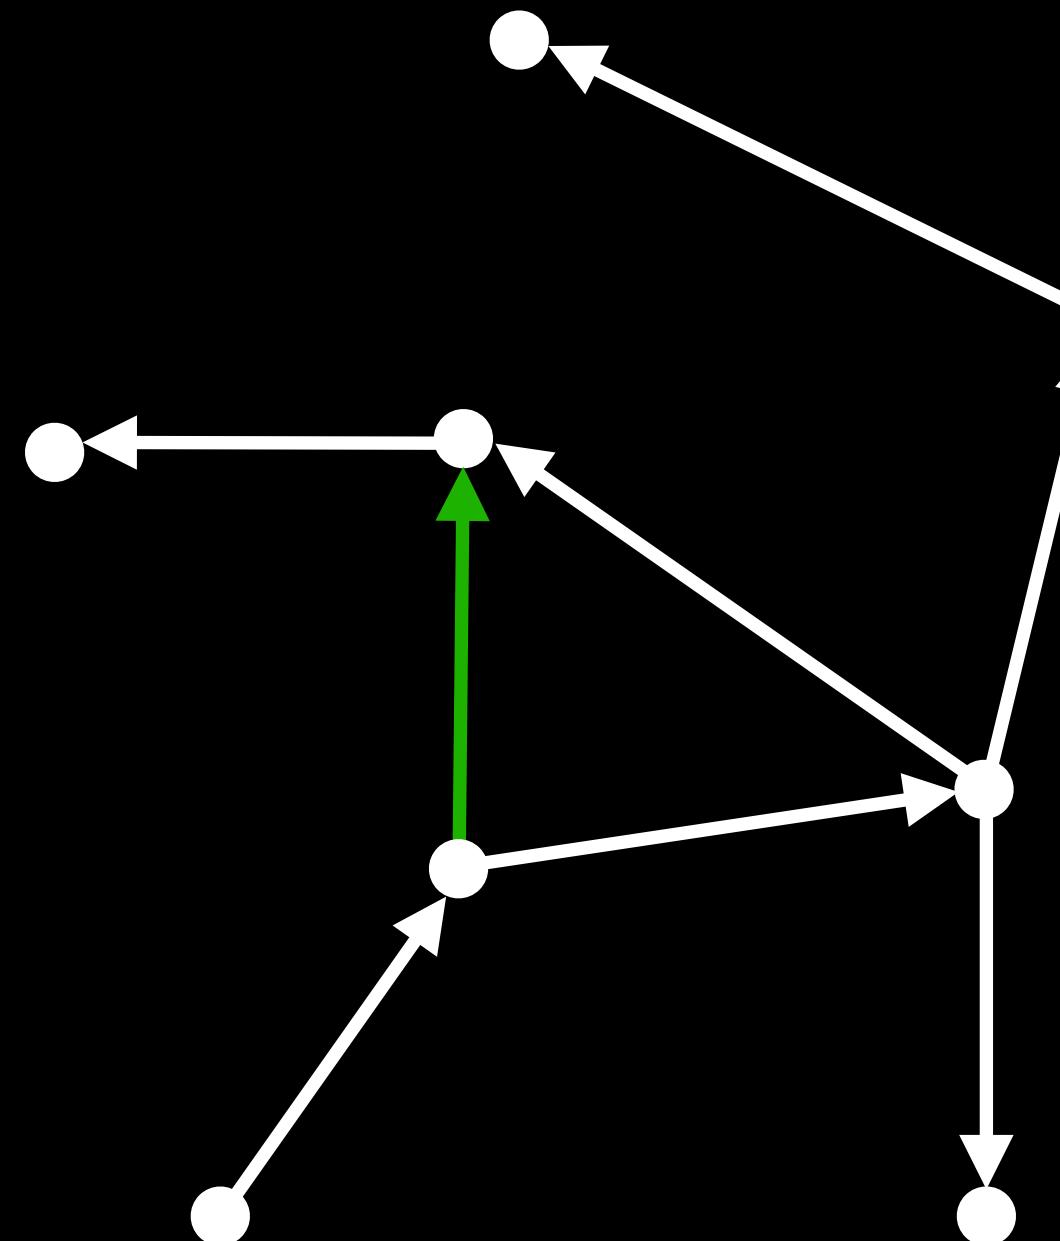
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

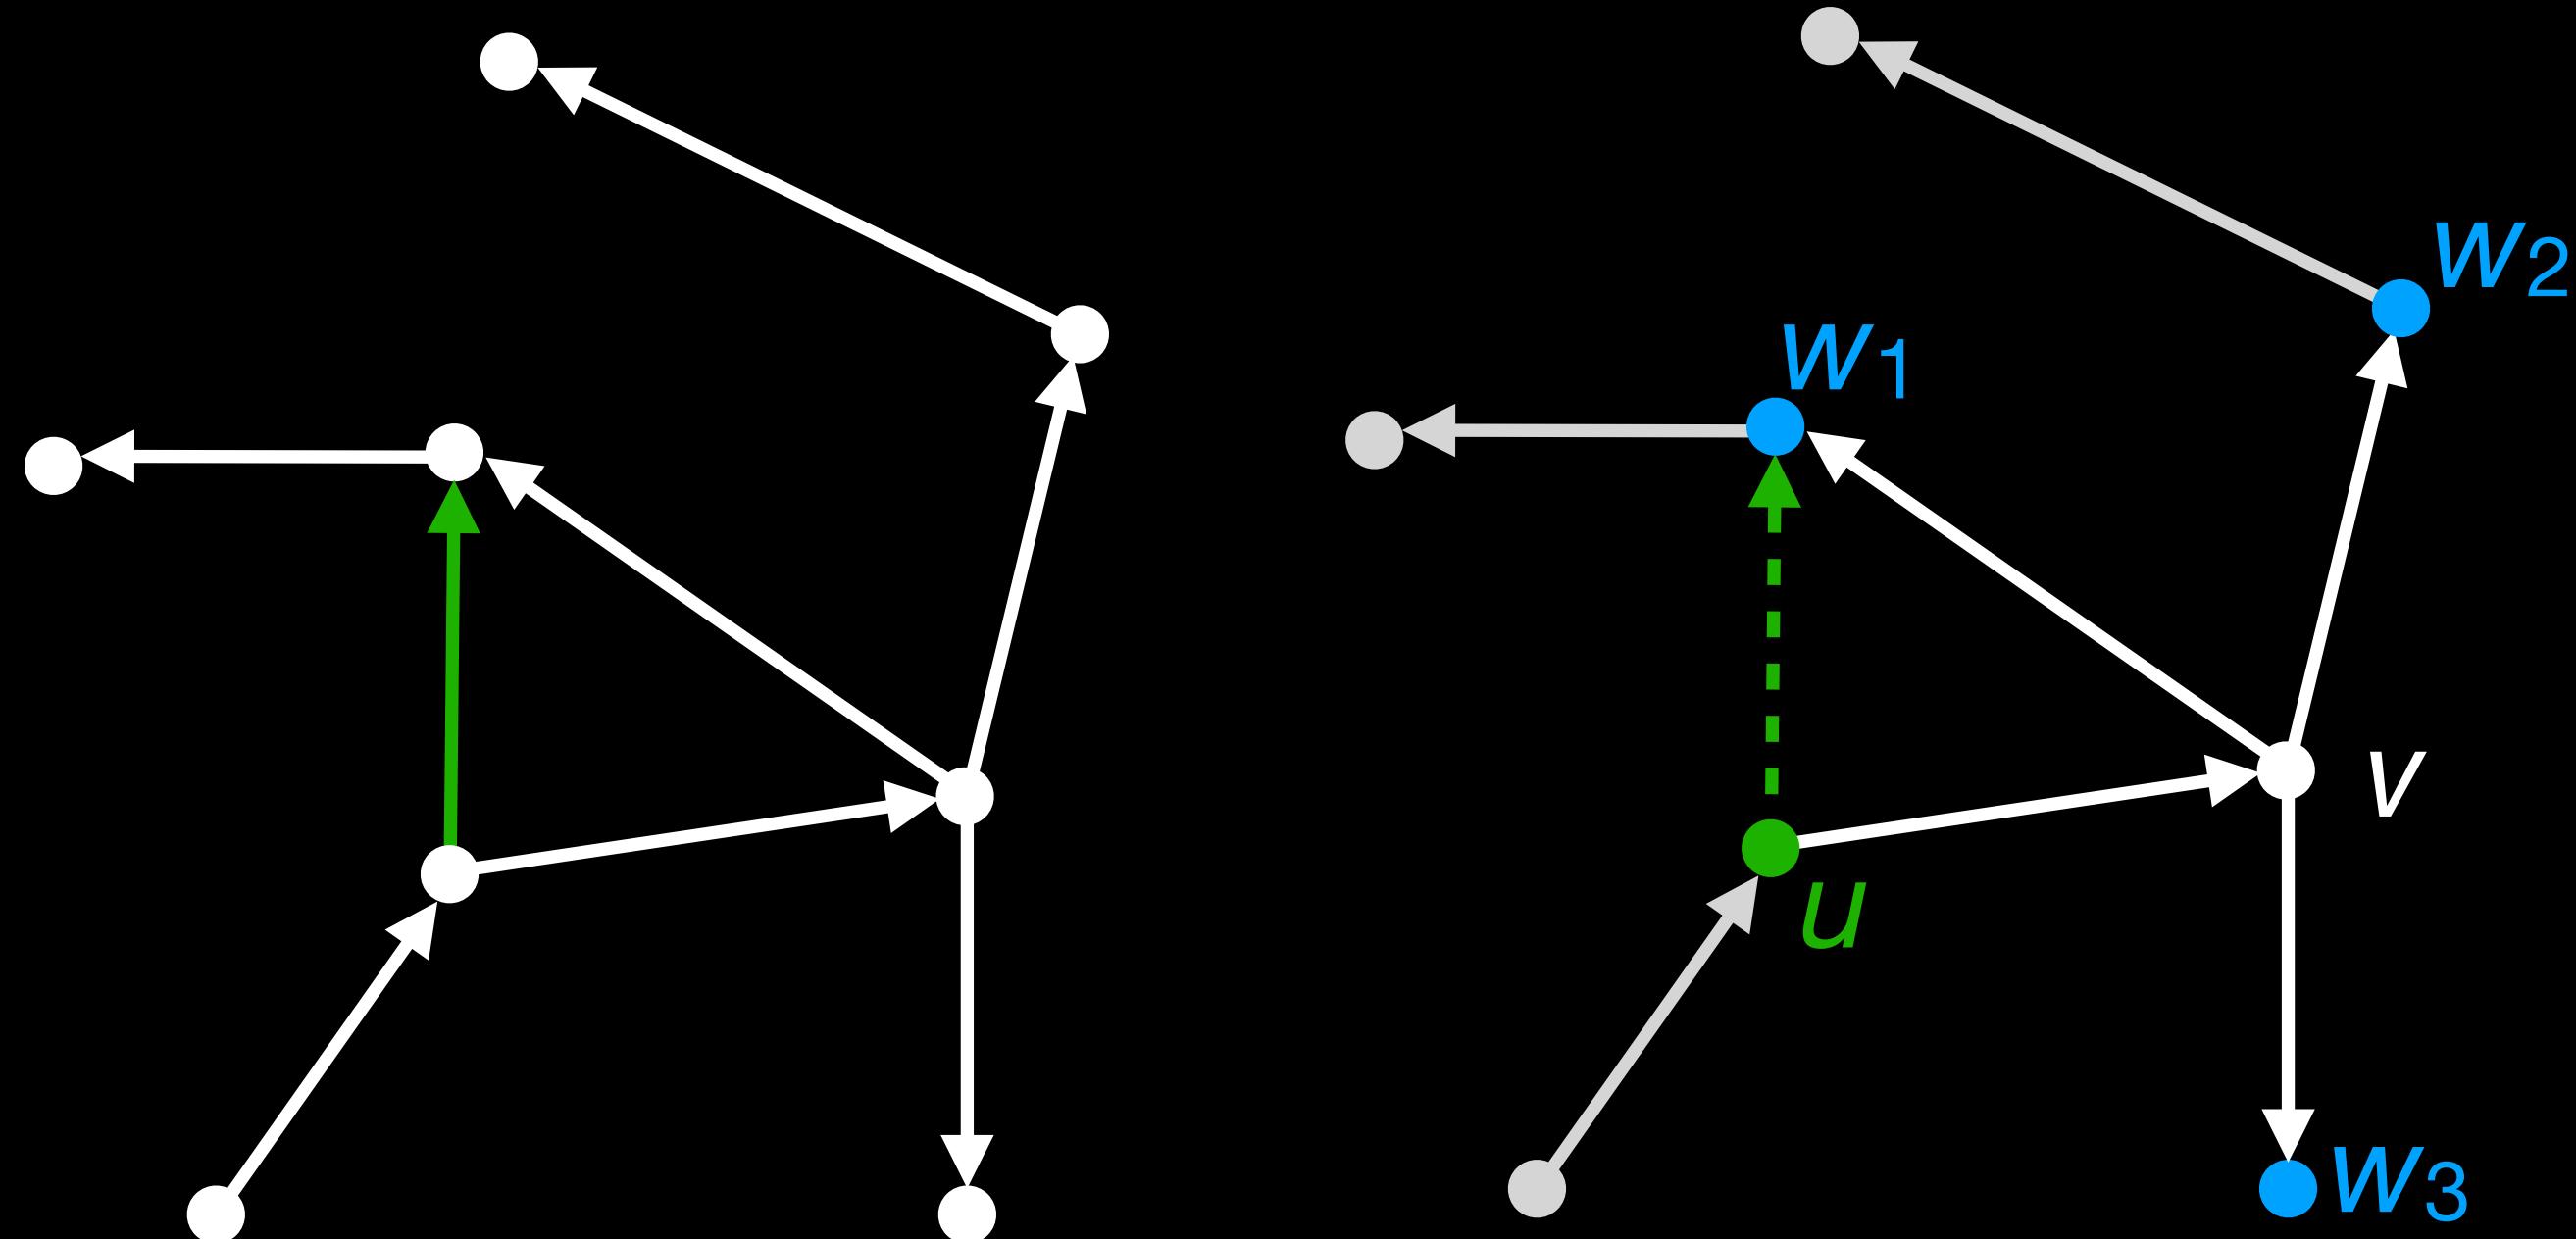
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

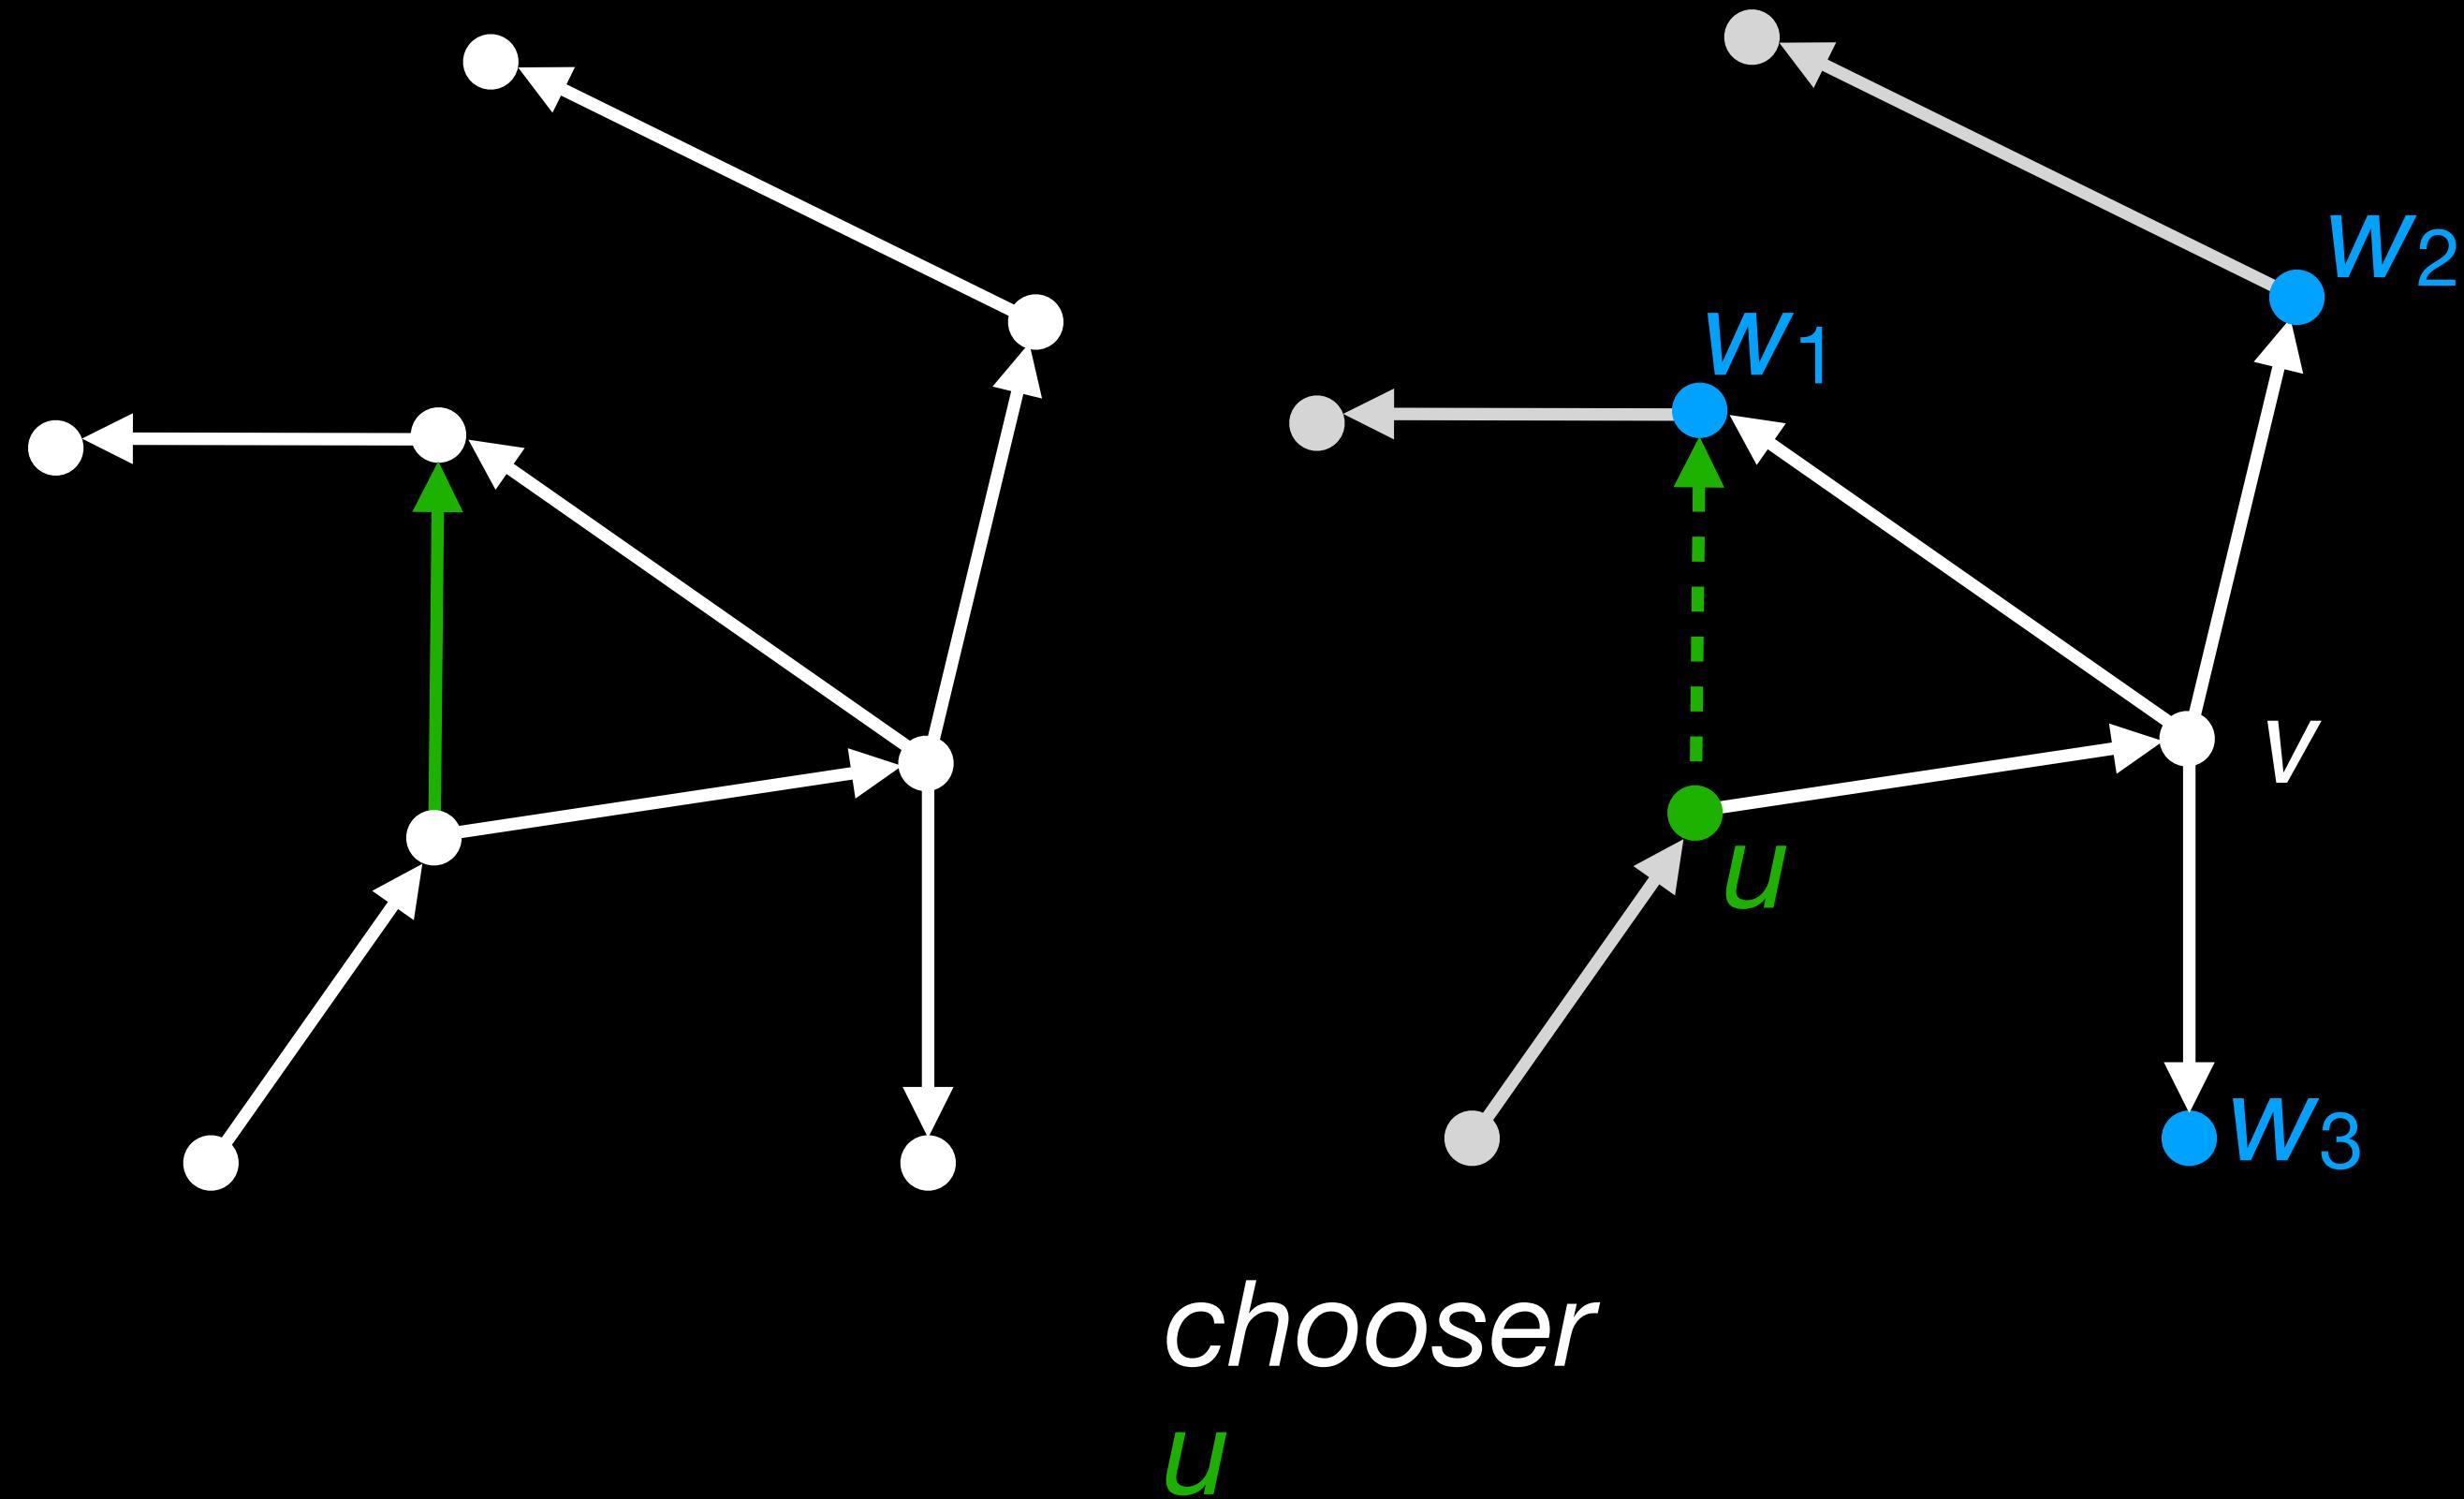
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

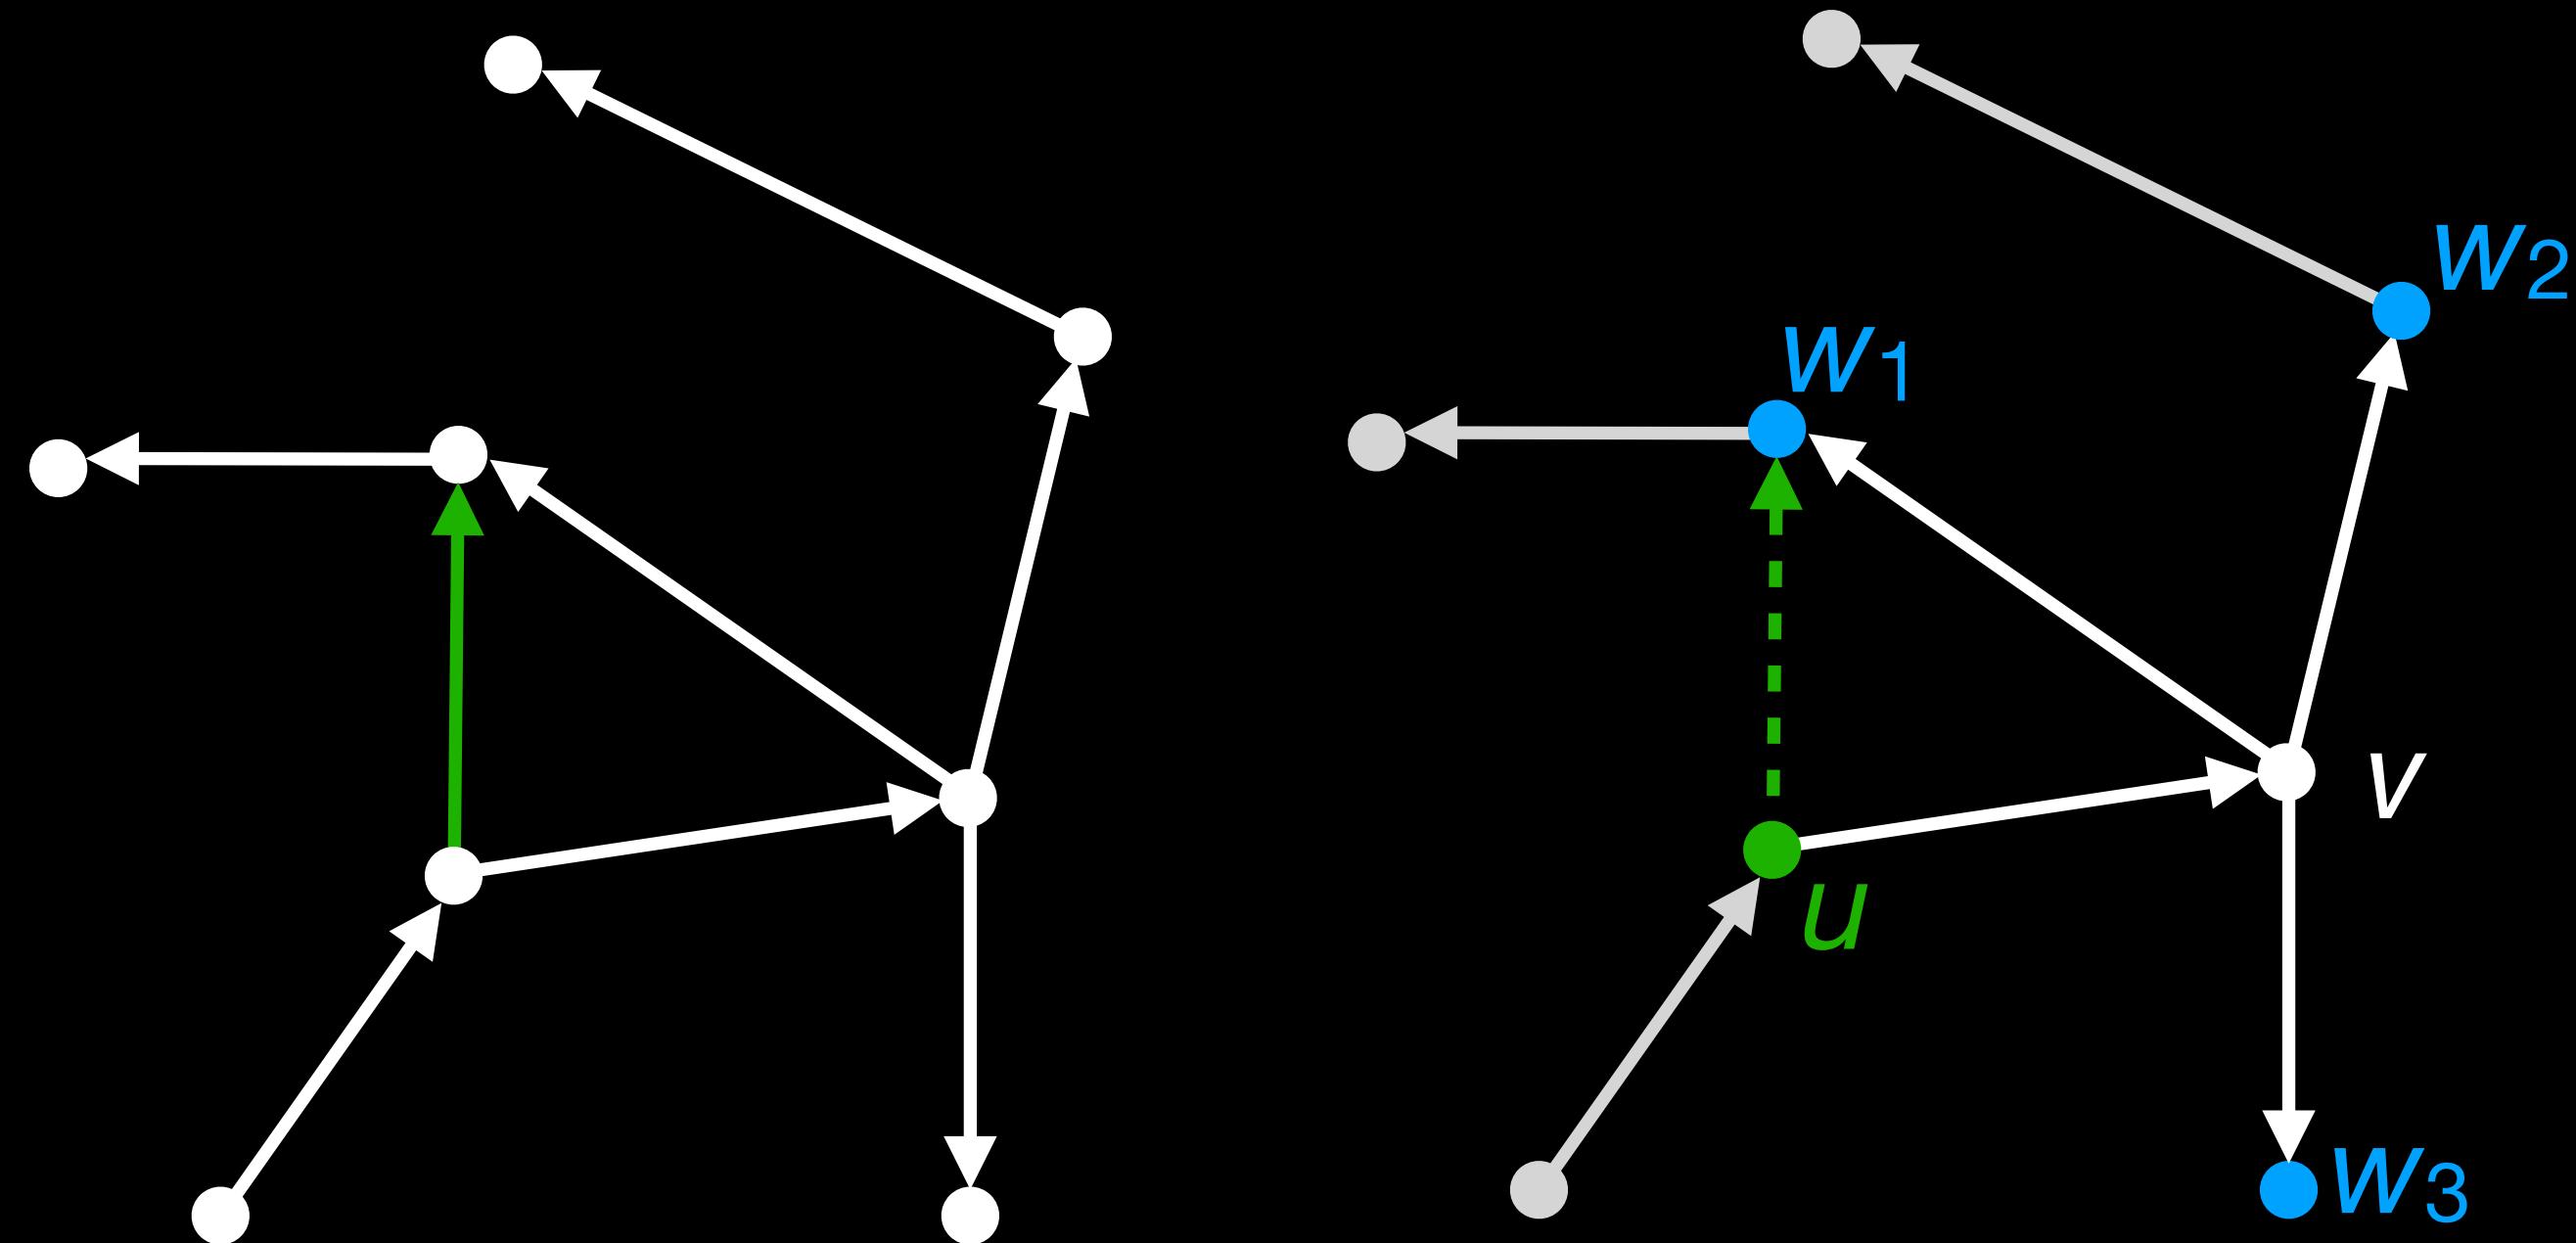
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

Our data
Timestamped edges
(including repeats)

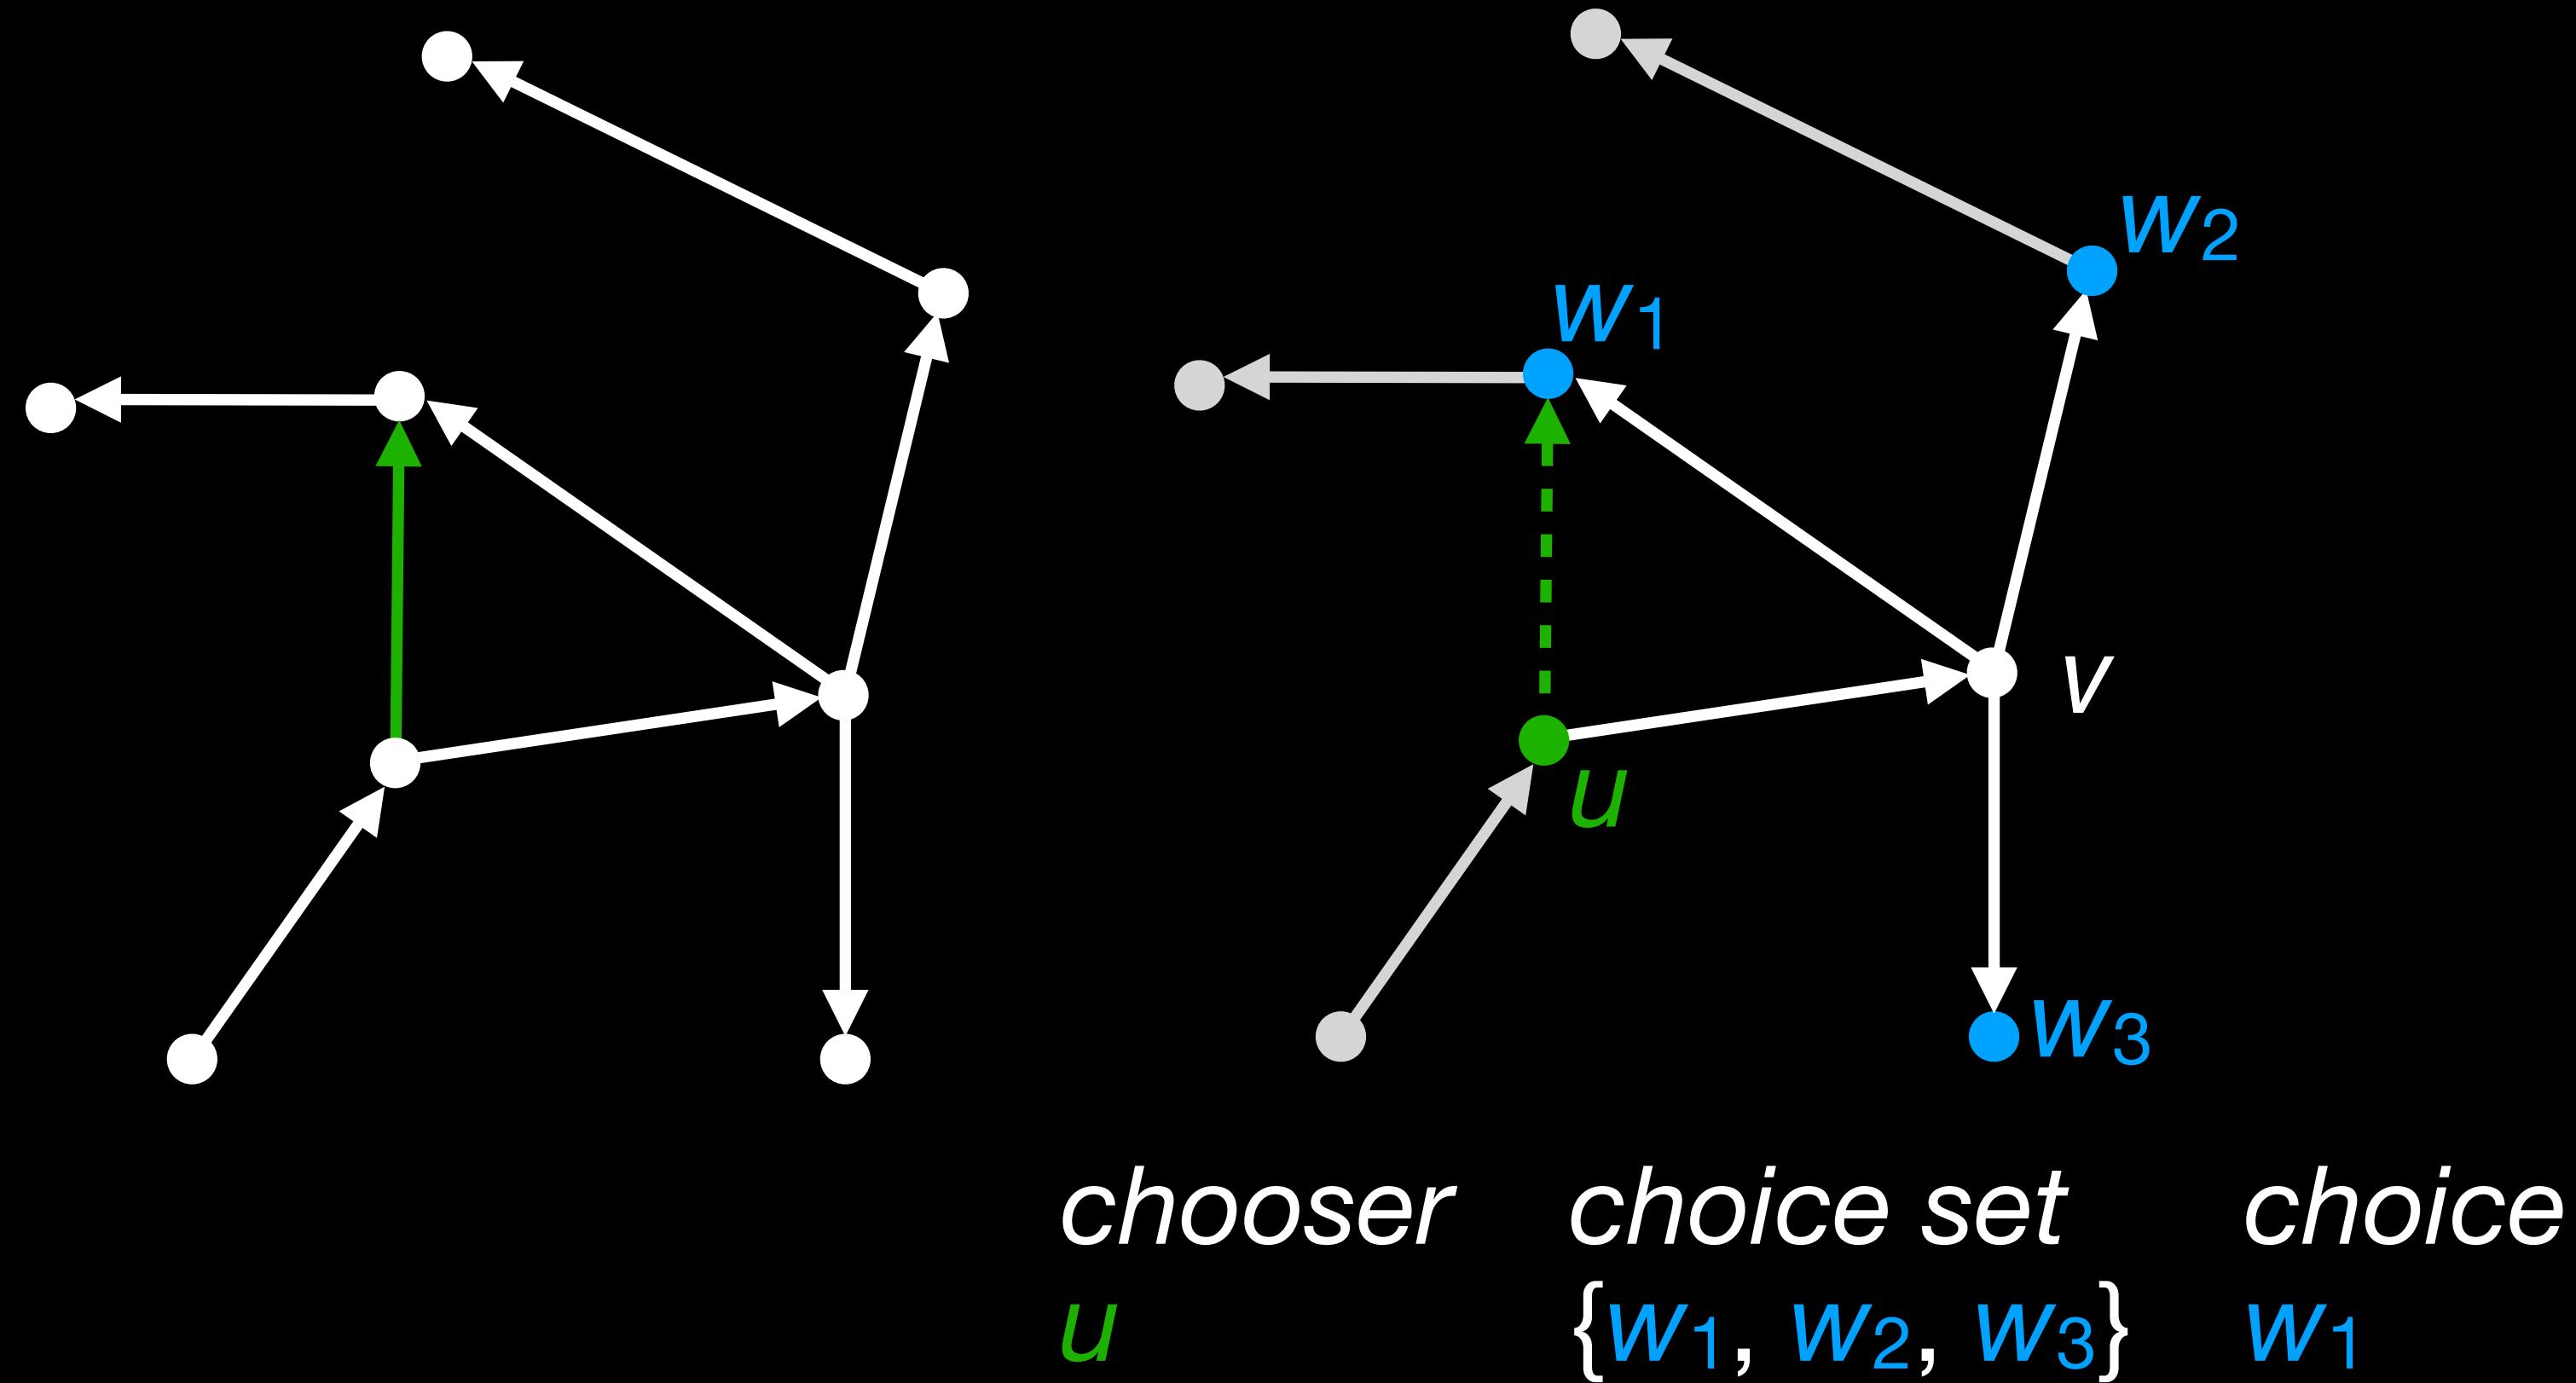


chooser *choice set*
 u $\{w_1, w_2, w_3\}$

Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

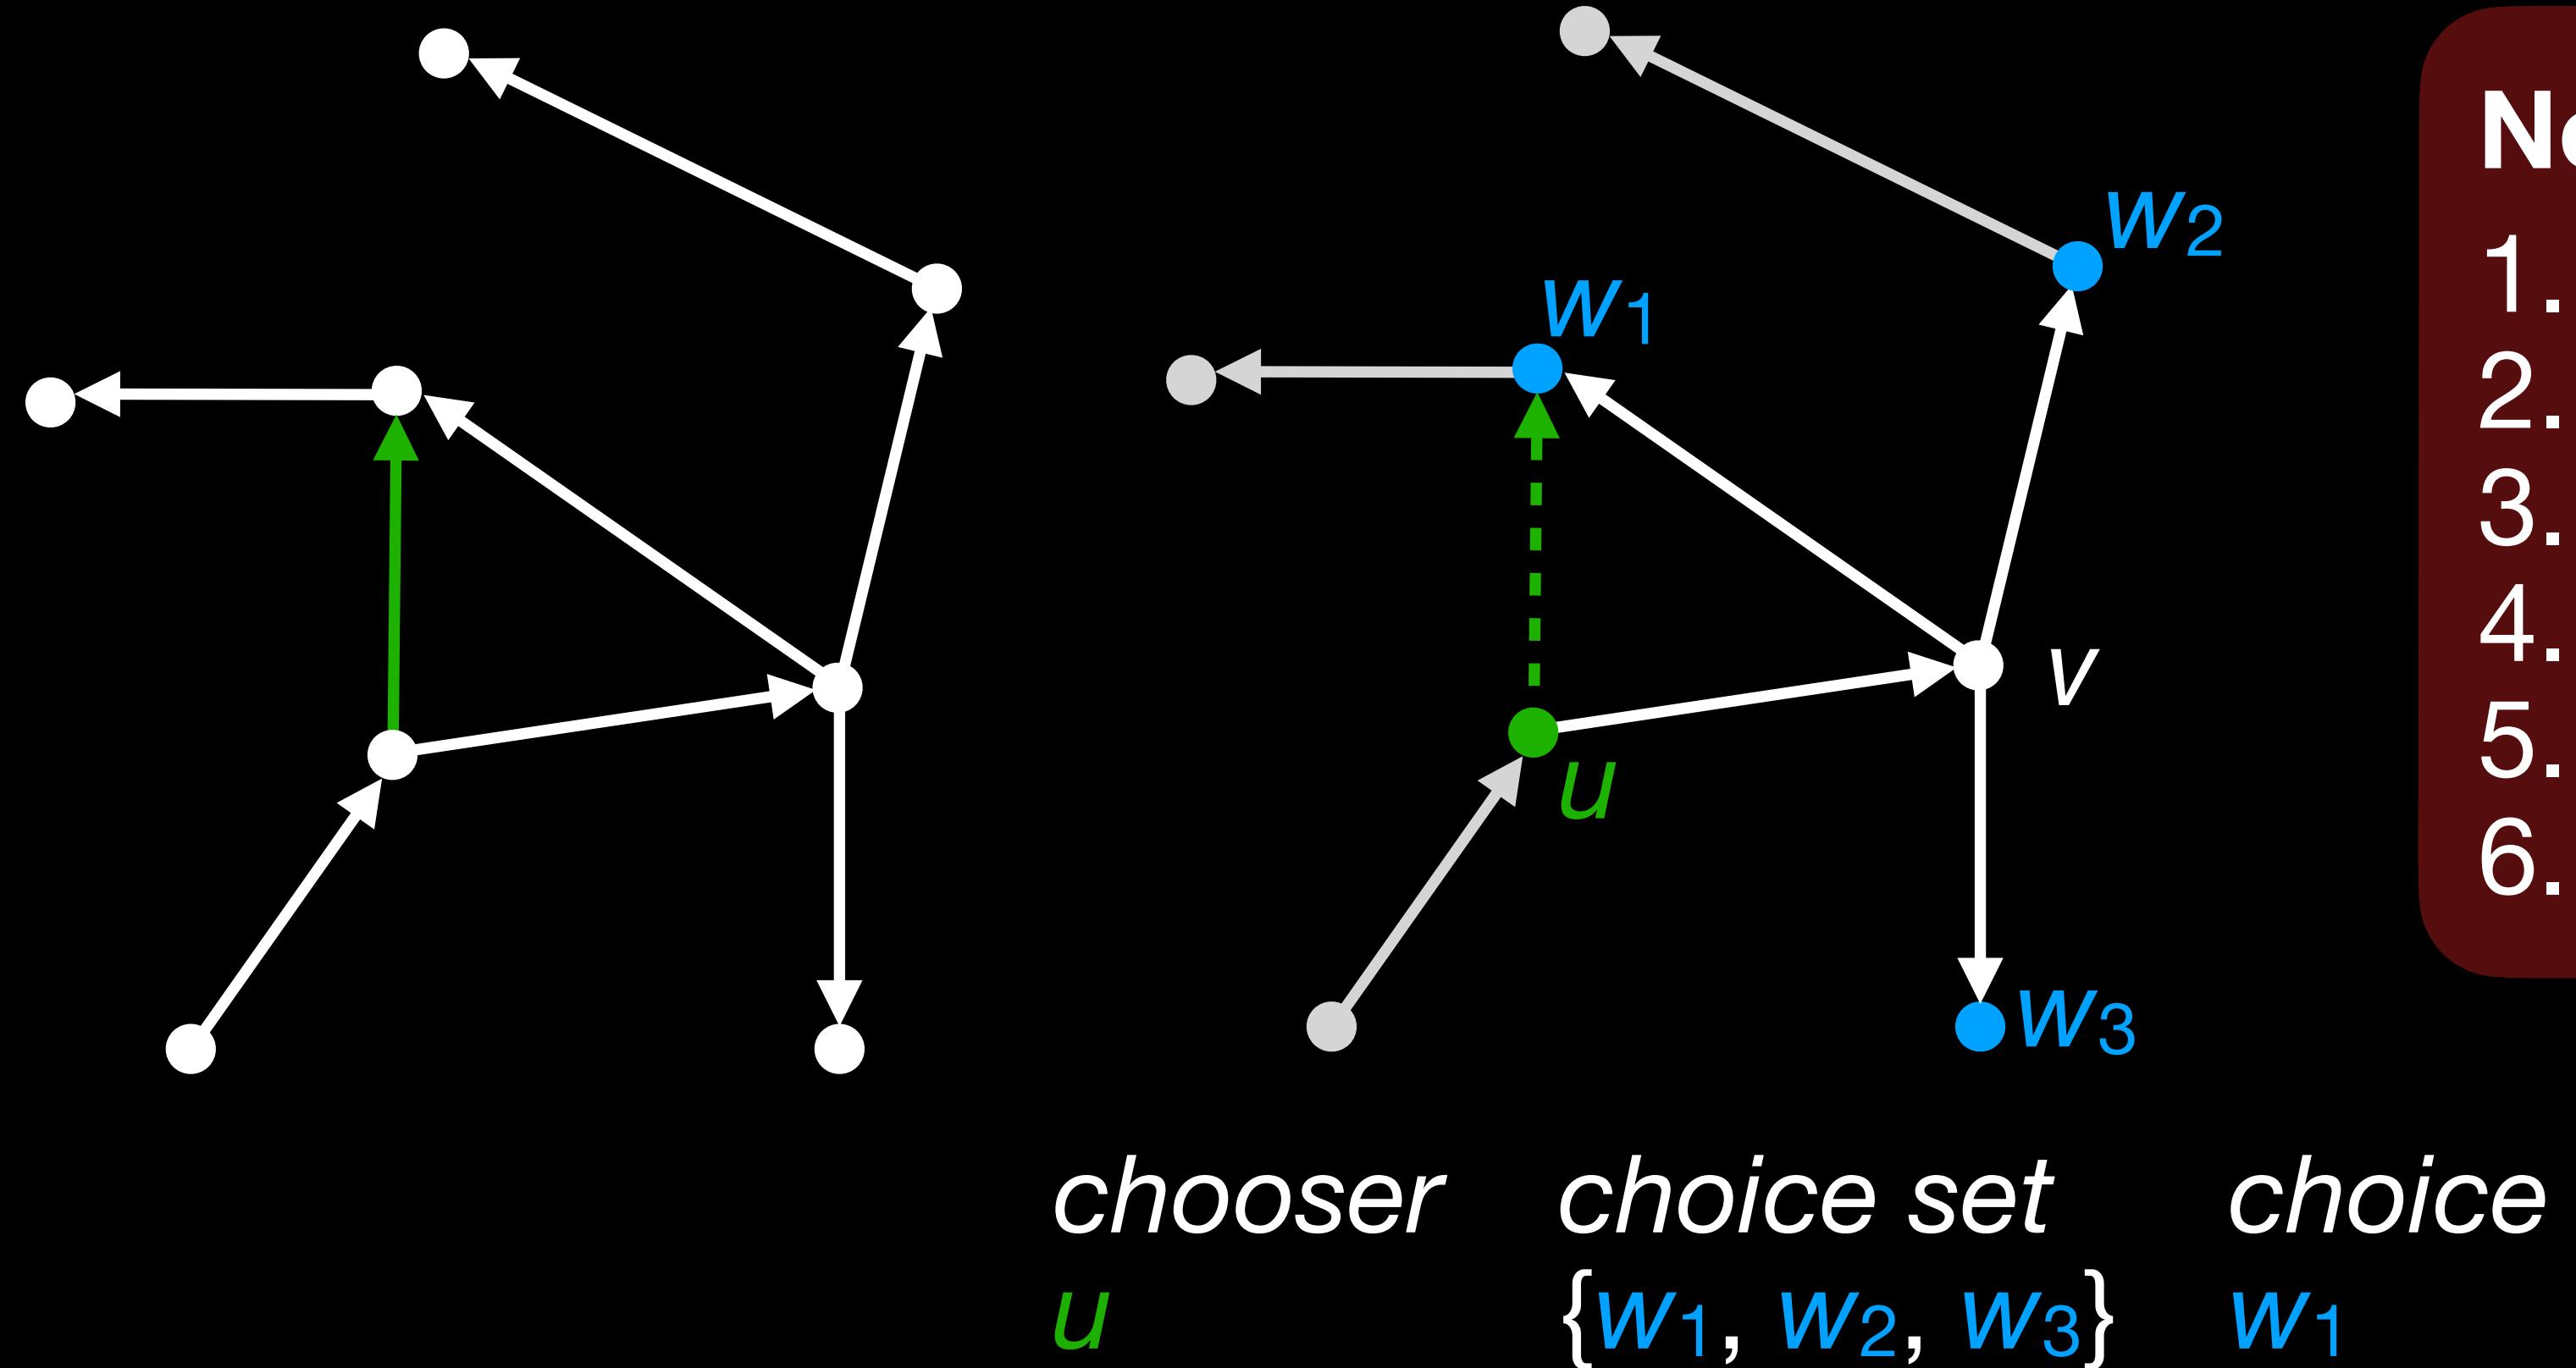
Our data
Timestamped edges
(including repeats)



Choosing to close triangles

Triadic closure offers small choice sets
→ tractable inference
→ varied choice sets

Our data
Timestamped edges
(including repeats)



- Node features**
1. in-degree of w
 2. # shared neighbors of u, w
 3. weight of edge $w \rightarrow u$
 4. time since last edge into w
 5. time since last edge out of w
 6. time since last $w \rightarrow u$ edge

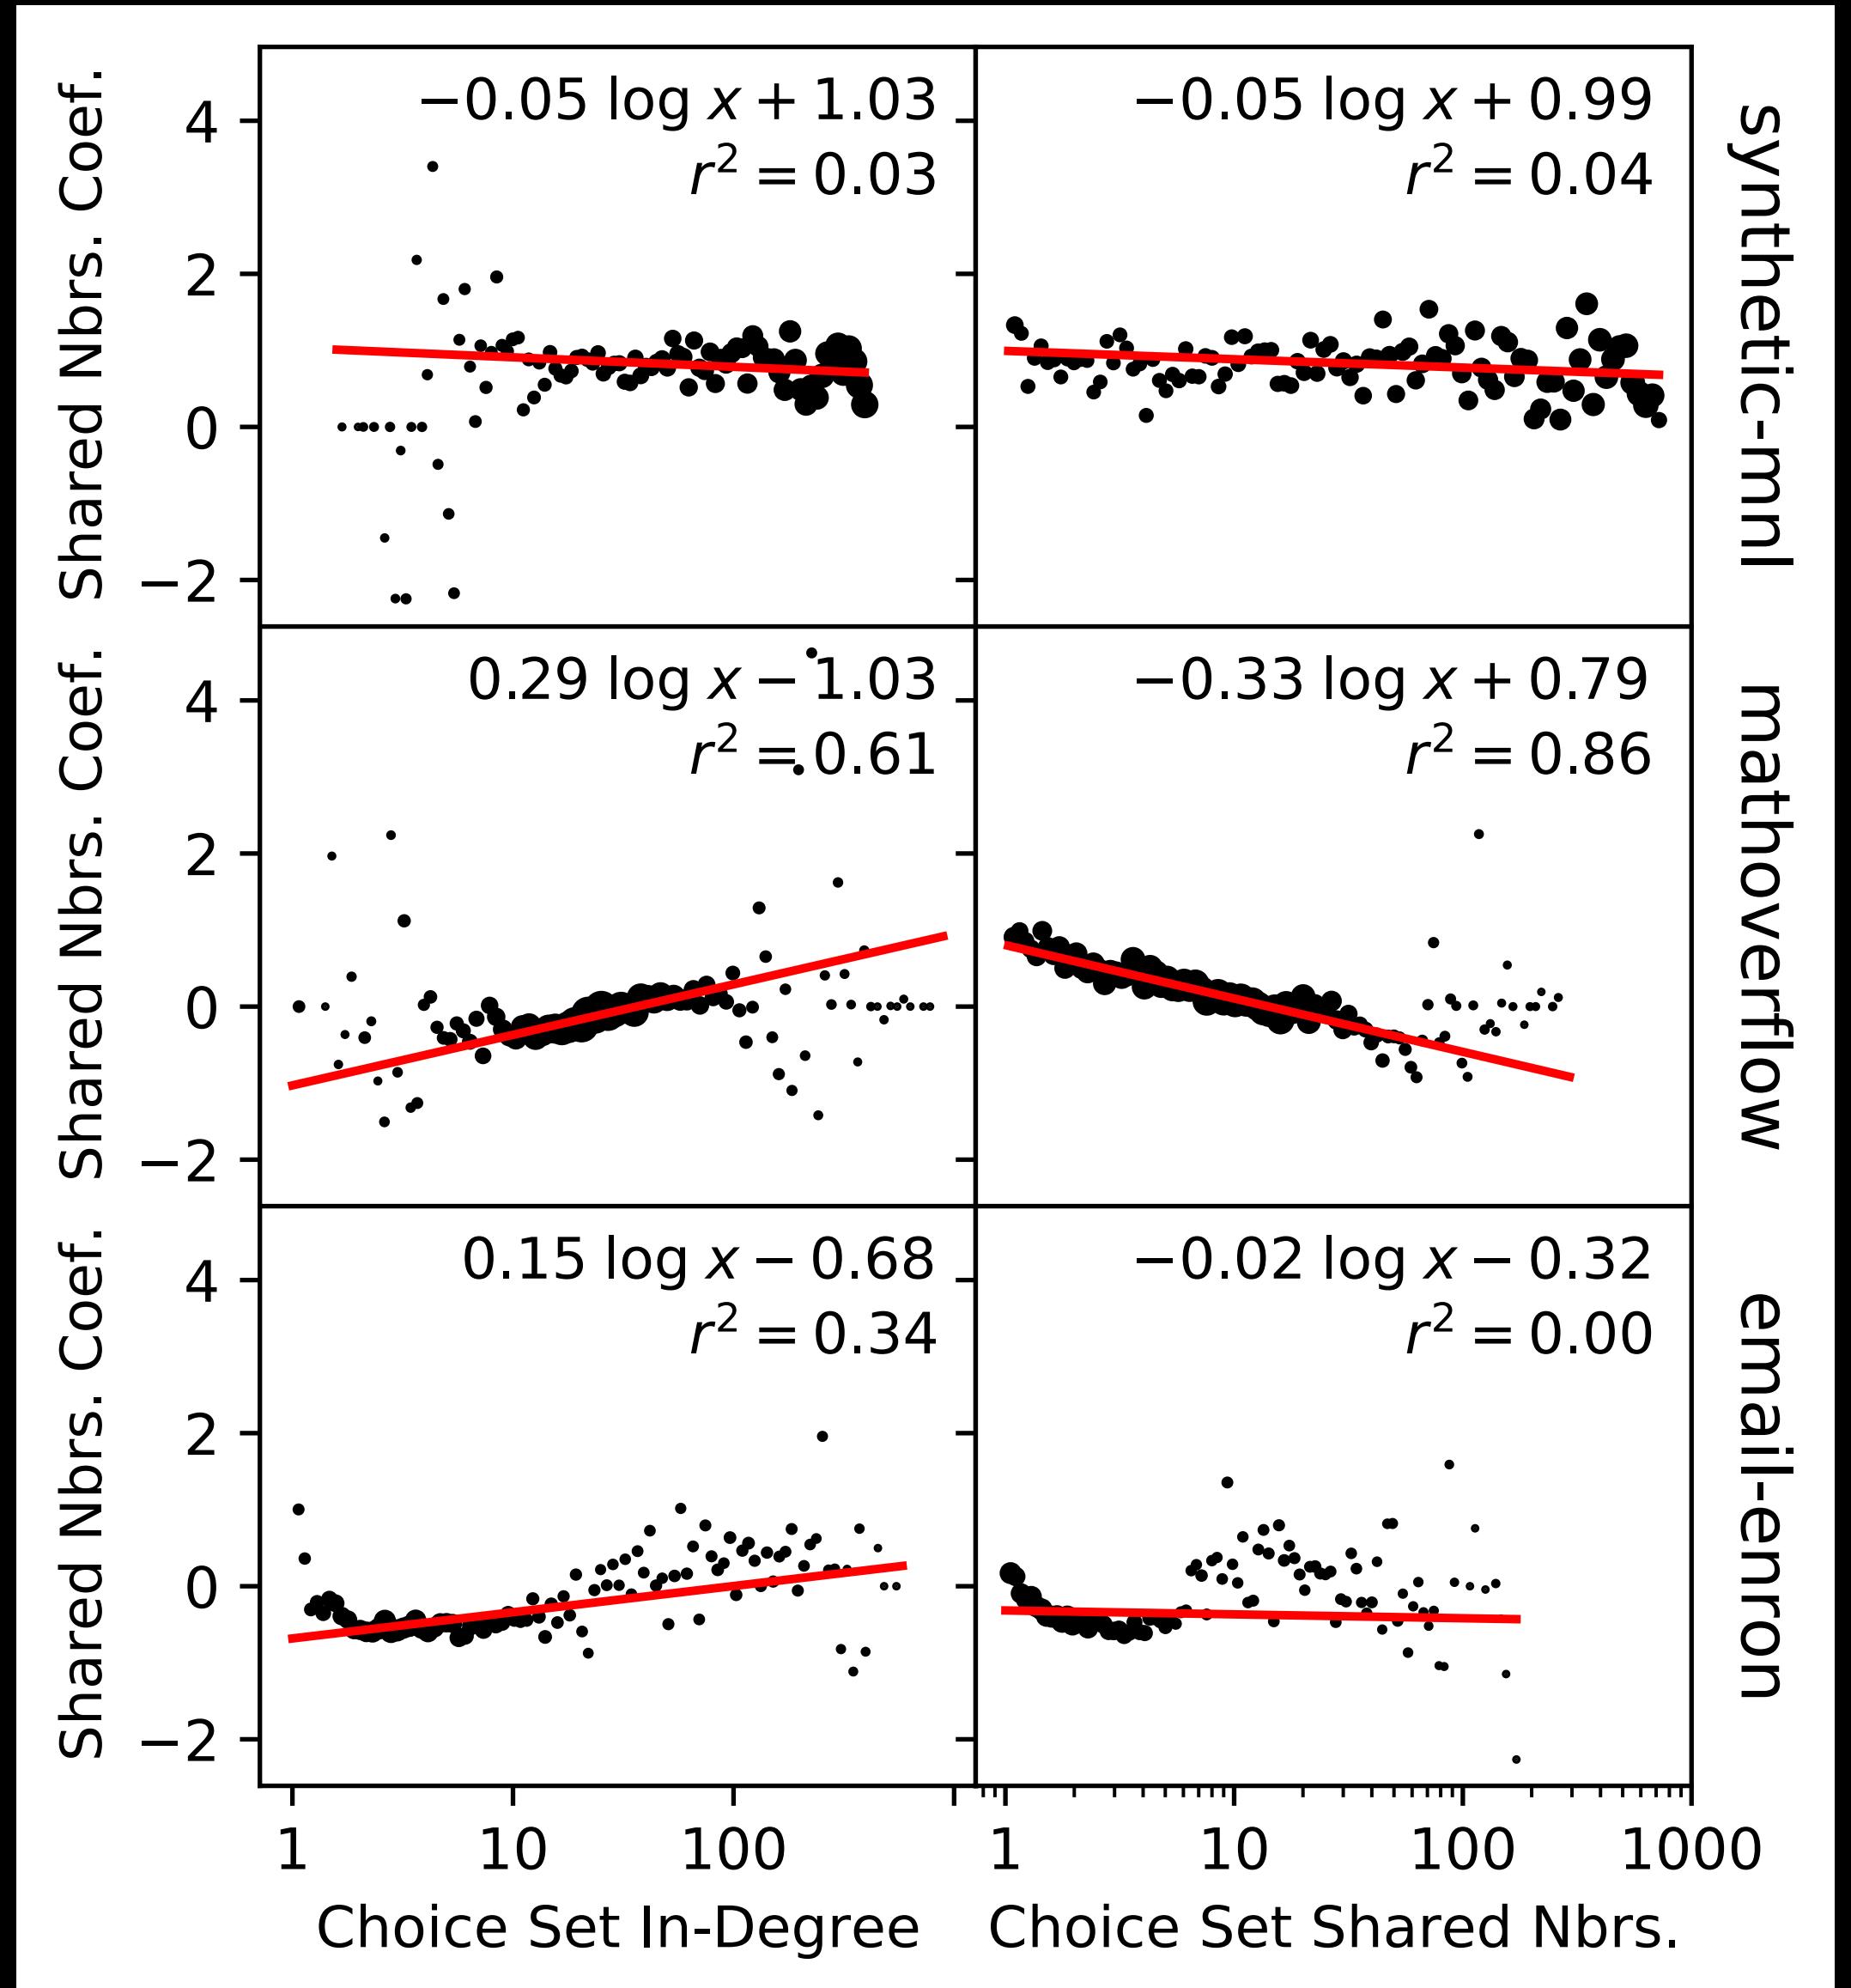
Context matters in triadic closure

Context matters in triadic closure

Datasets

- email-enron
- email-eu
- email-w3c
- wiki-talk
- reddit-hyperlink
- bitcoin-alpha
- bitcoin-otc
- mathoverflow
- college-msg
- facebook-wall
- sms-a
- sms-b
- sms-c

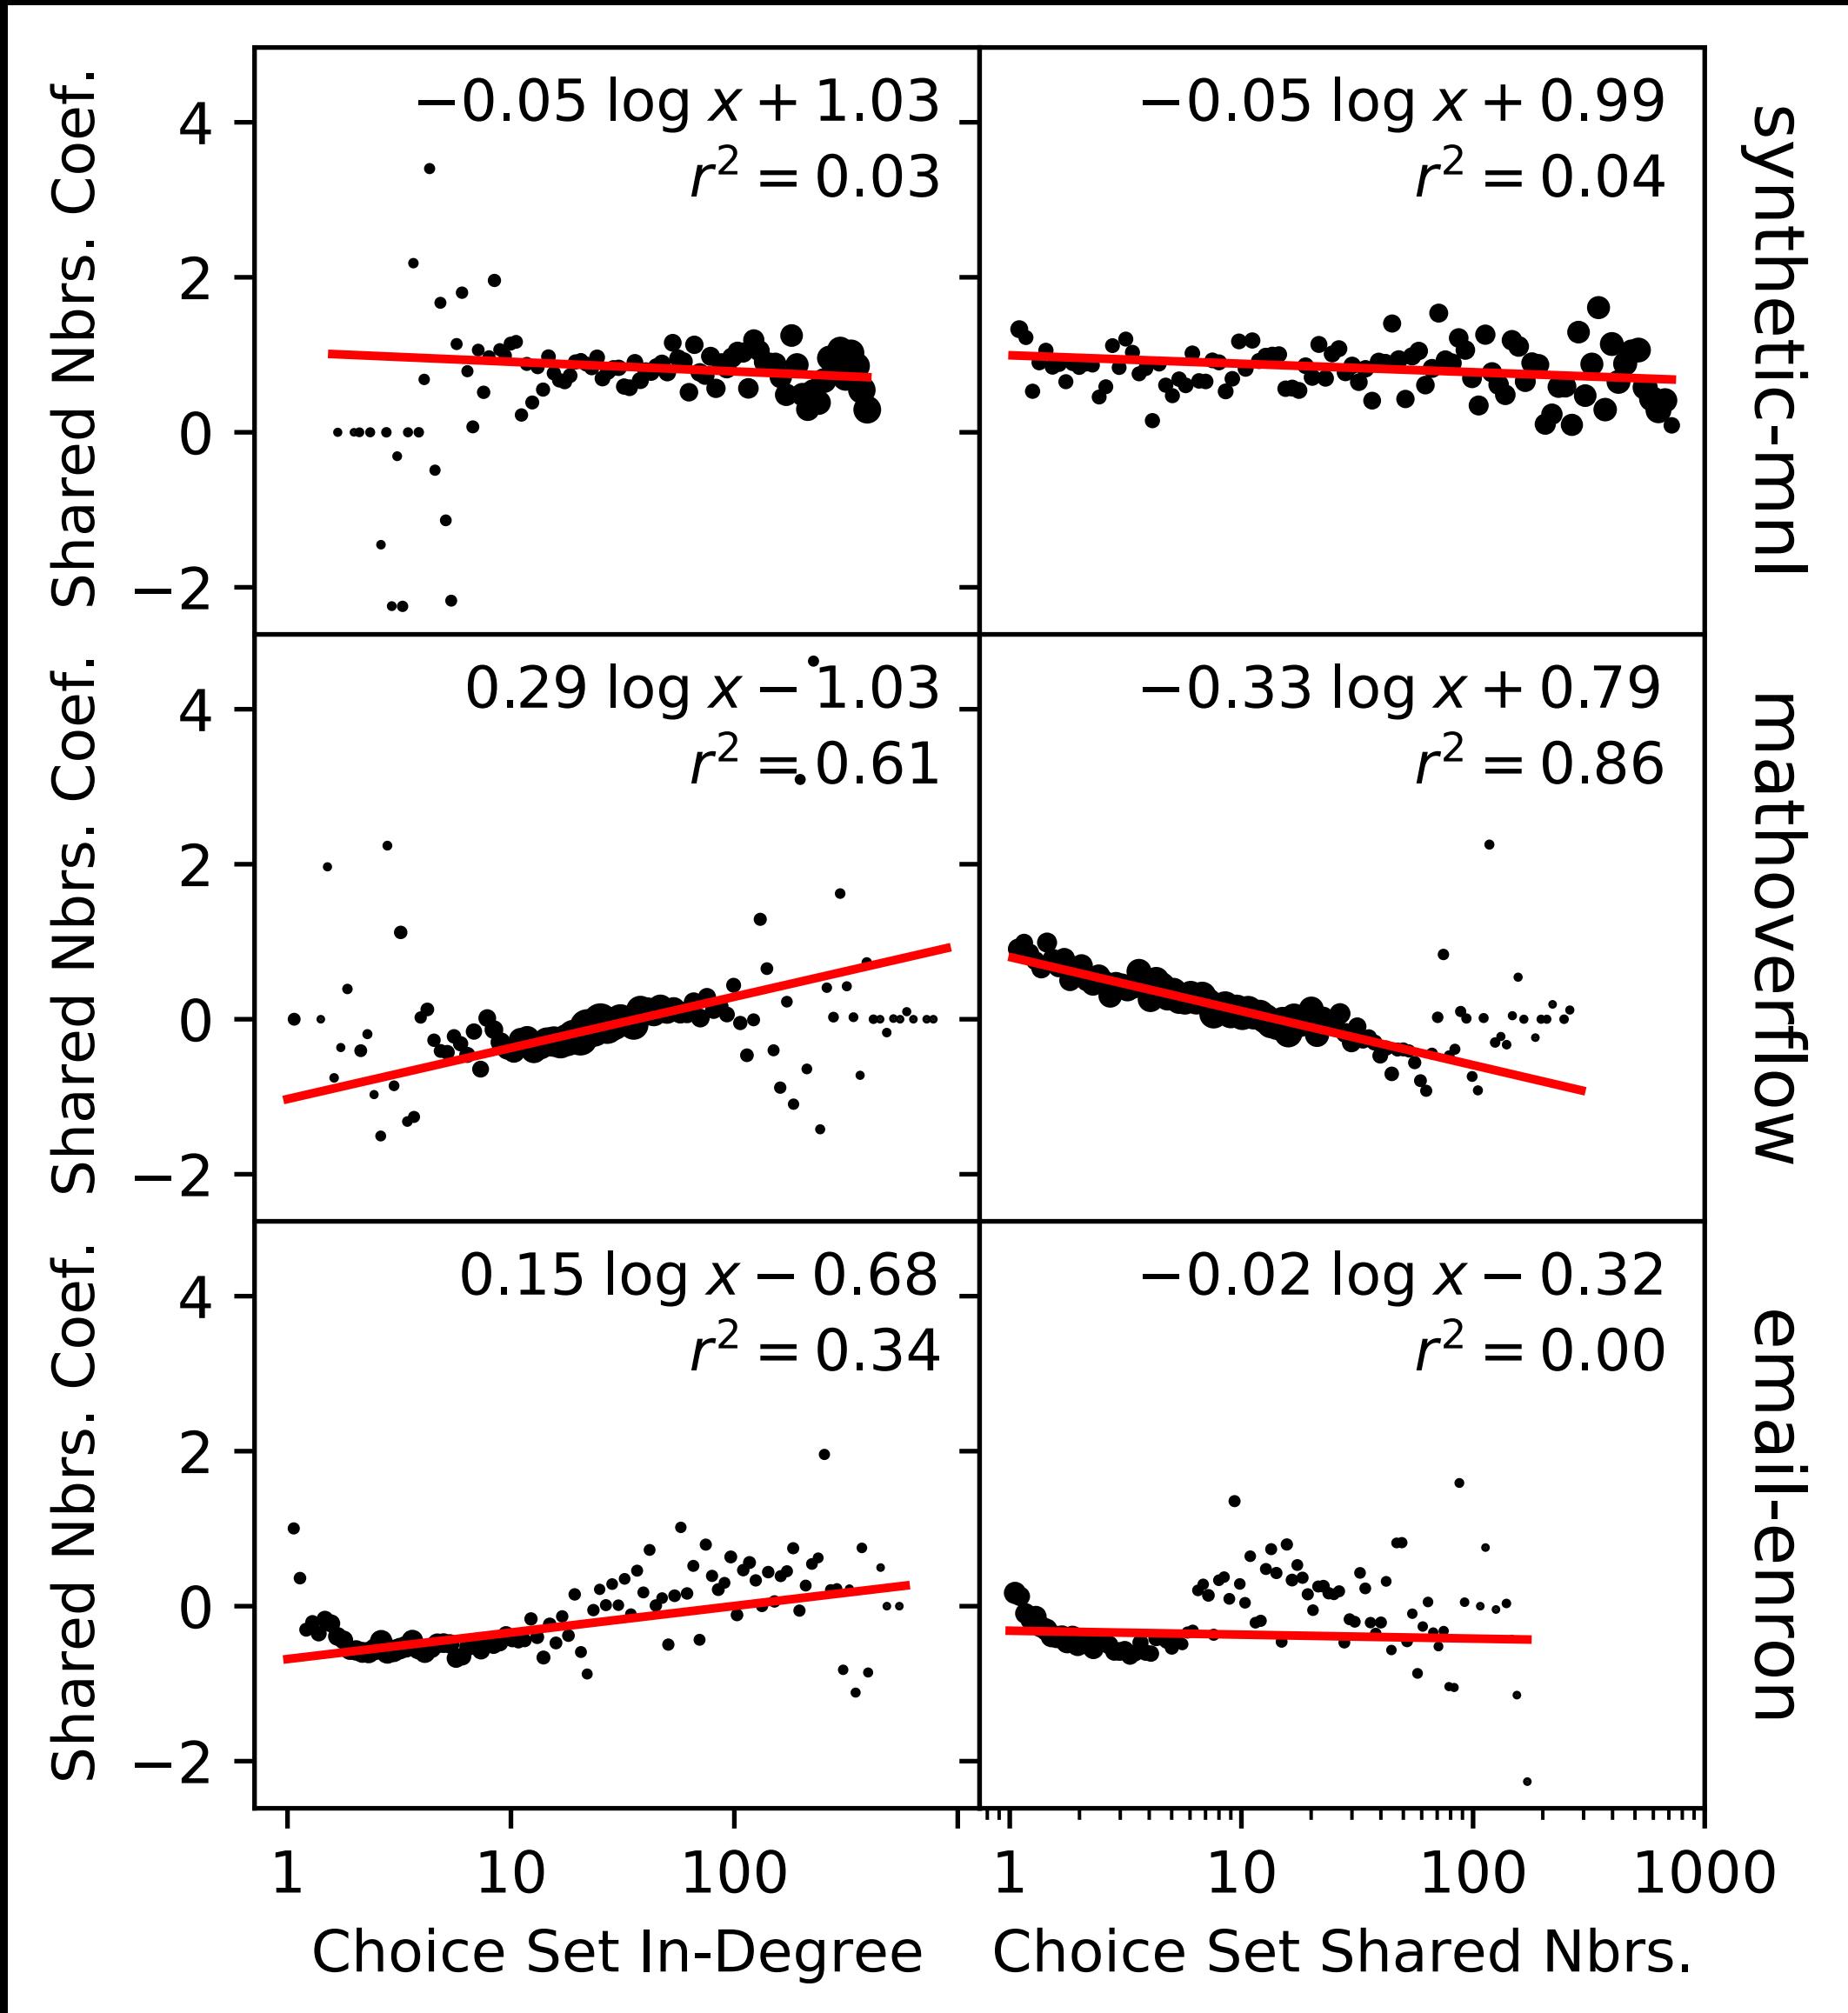
Context matters in triadic closure



synthetic-mnl mathoverflow email-enron

- Datasets**
- email-enron
 - email-eu
 - email-w3c
 - wiki-talk
 - reddit-hyperlink
 - bitcoin-alpha
 - bitcoin-otc
 - mathoverflow
 - college-msg
 - facebook-wall
 - sms-a
 - sms-b
 - sms-c

Context matters in triadic closure

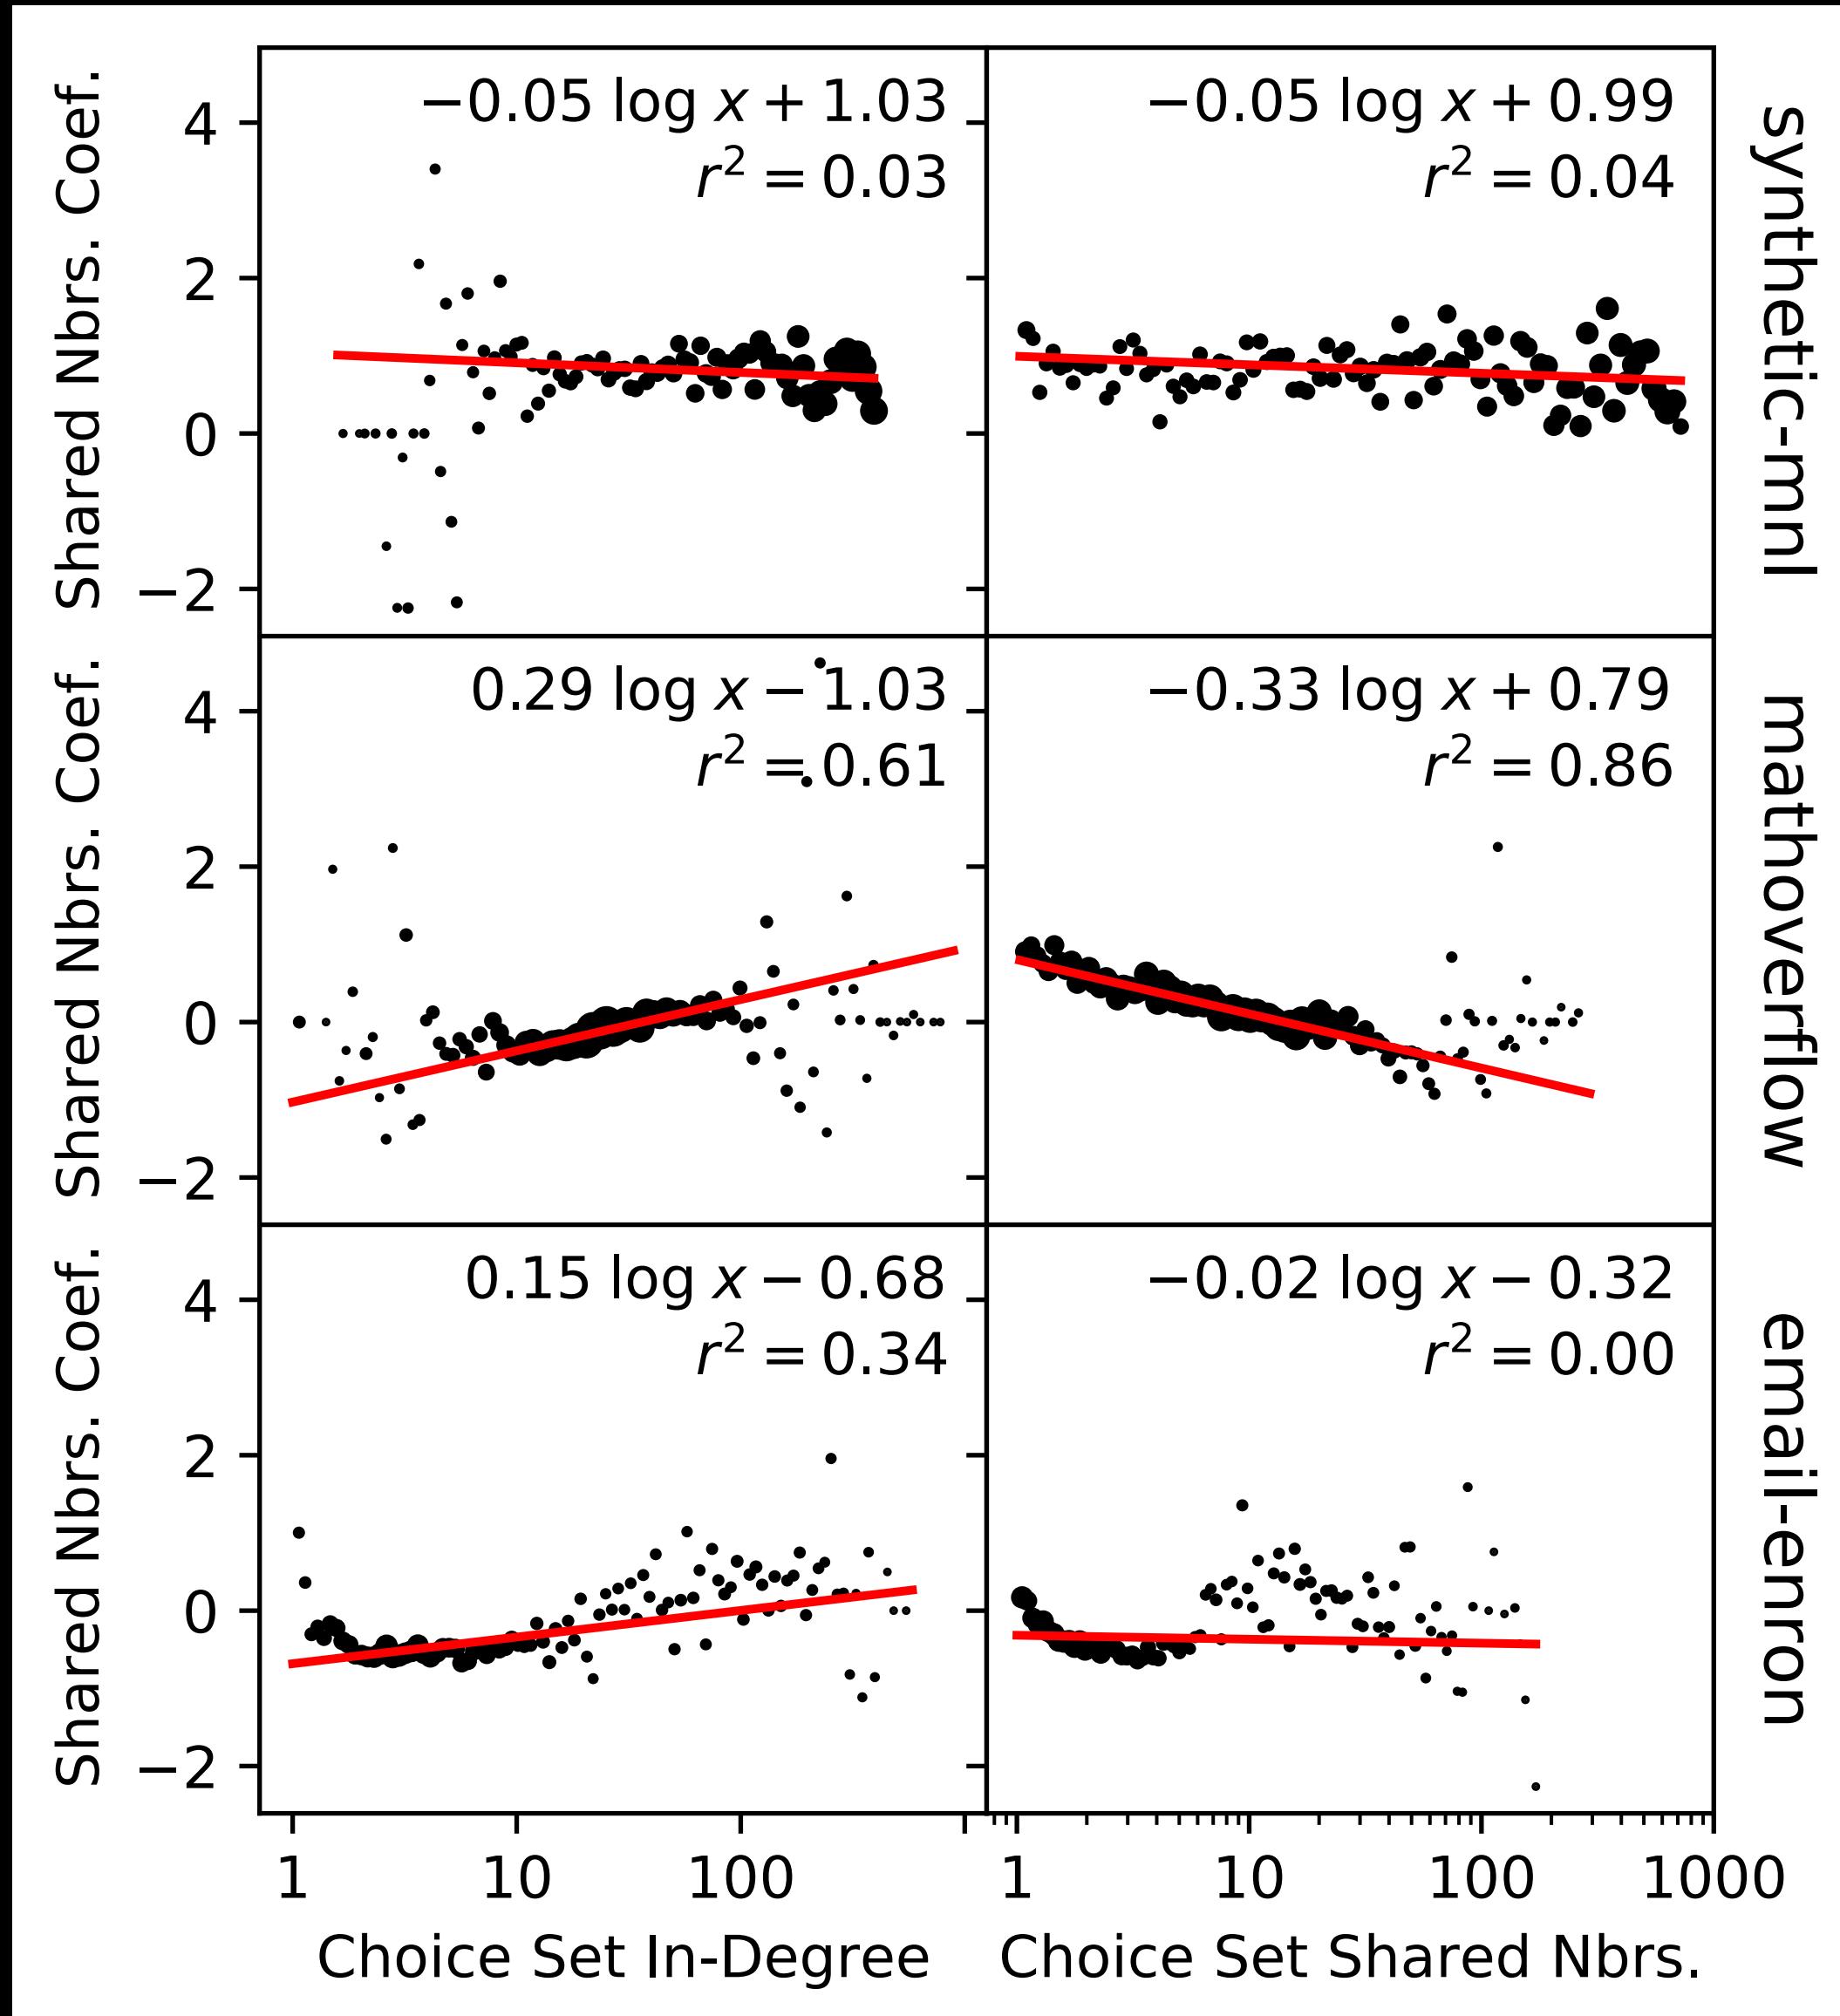


Synthetic data,
no context effects

Datasets

- email-enron
- email-eu
- email-w3c
- wiki-talk
- reddit-hyperlink
- bitcoin-alpha
- bitcoin-otc
- mathoverflow
- college-msg
- facebook-wall
- sms-a
- sms-b
- sms-c

Context matters in triadic closure

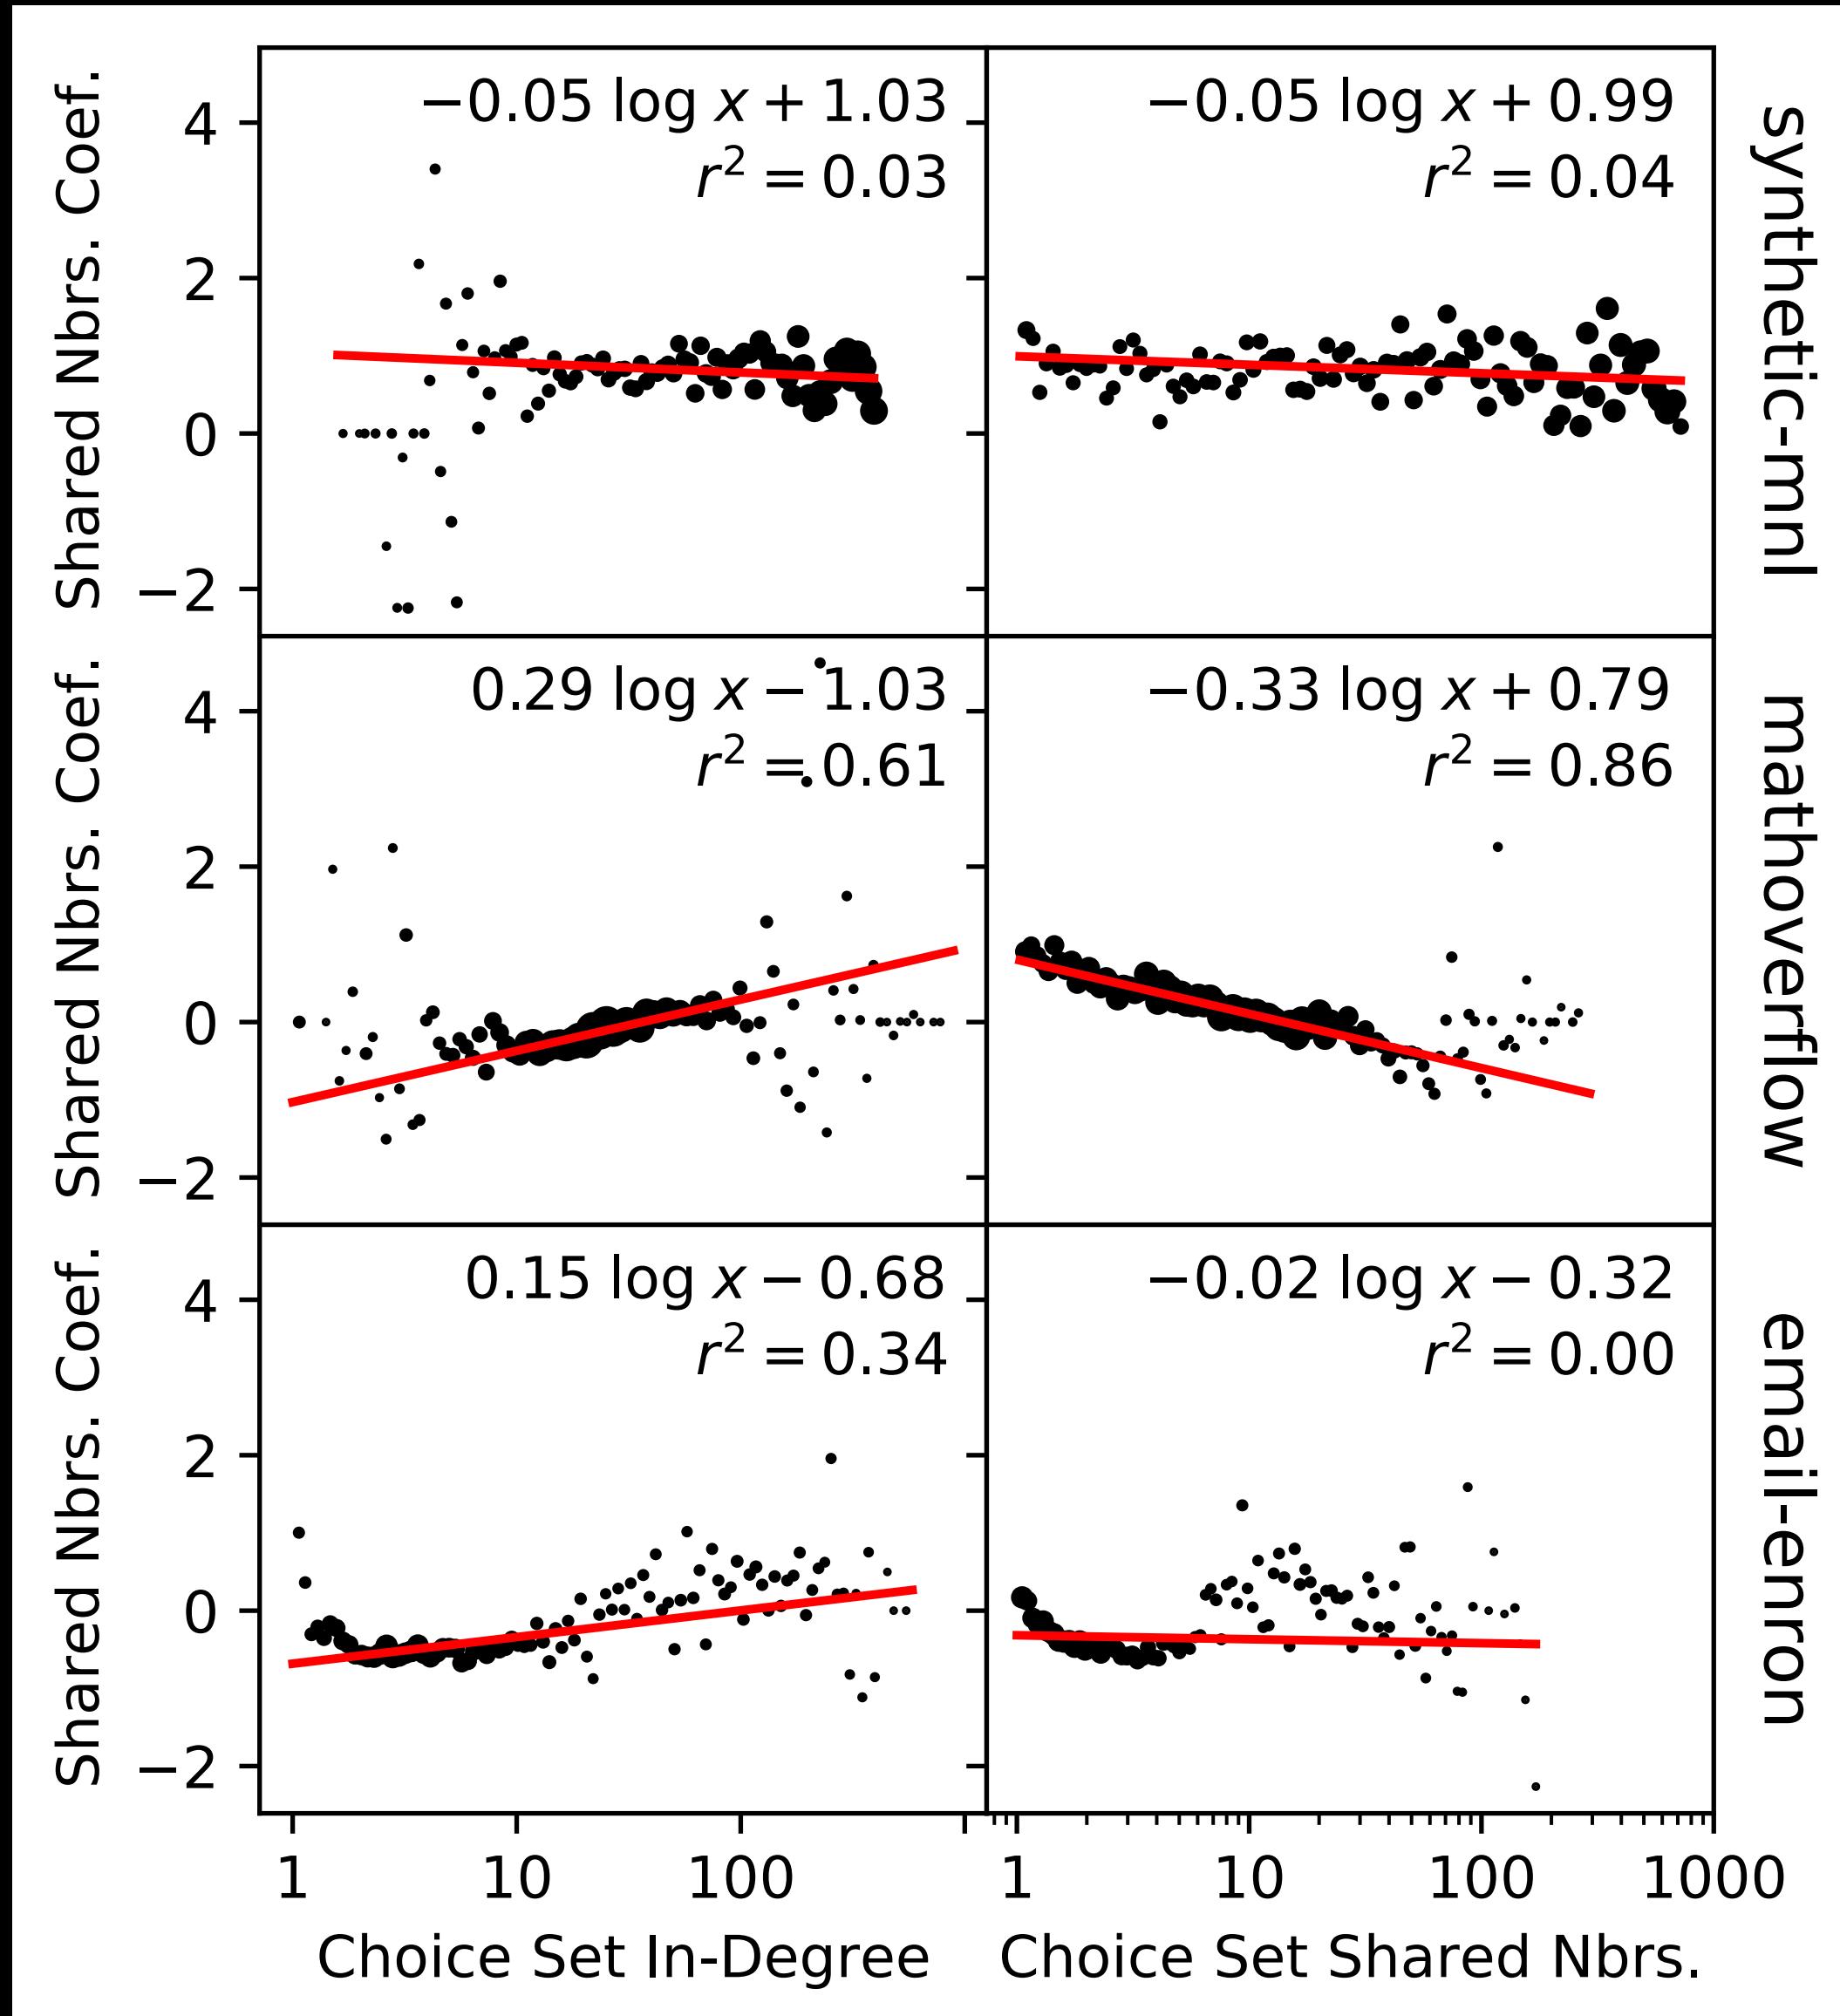


Synthetic data,
no context effects

Commenting network,
linear context effects

- Datasets**
- email-enron
 - email-eu
 - email-w3c
 - wiki-talk
 - reddit-hyperlink
 - bitcoin-alpha
 - bitcoin-otc
 - mathoverflow
 - college-msg
 - facebook-wall
 - sms-a
 - sms-b
 - sms-c

Context matters in triadic closure



Synthetic data,
no context effects

Commenting network,
linear context effects

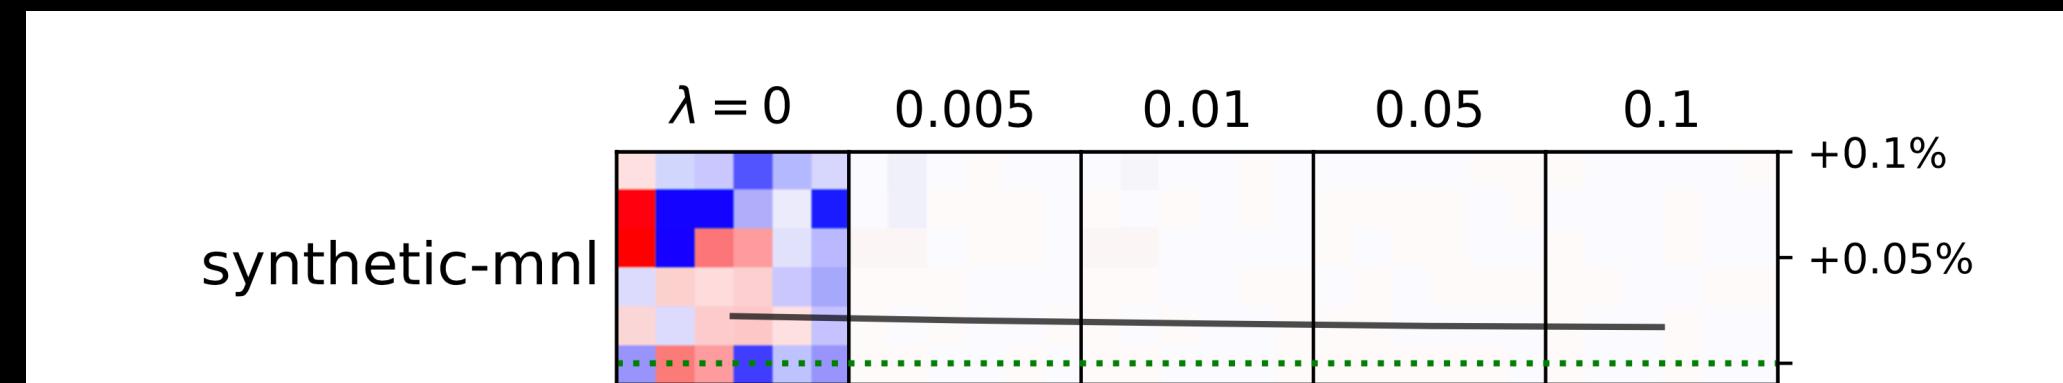
Email network,
nonlinear context effects?

Datasets

- email-enron
- email-eu
- email-w3c
- wiki-talk
- reddit-hyperlink
- bitcoin-alpha
- bitcoin-otc
- mathoverflow
- college-msg
- facebook-wall
- sms-a
- sms-b
- sms-c

LCL reveals interpretable feature context effects

LCL reveals interpretable feature context effects

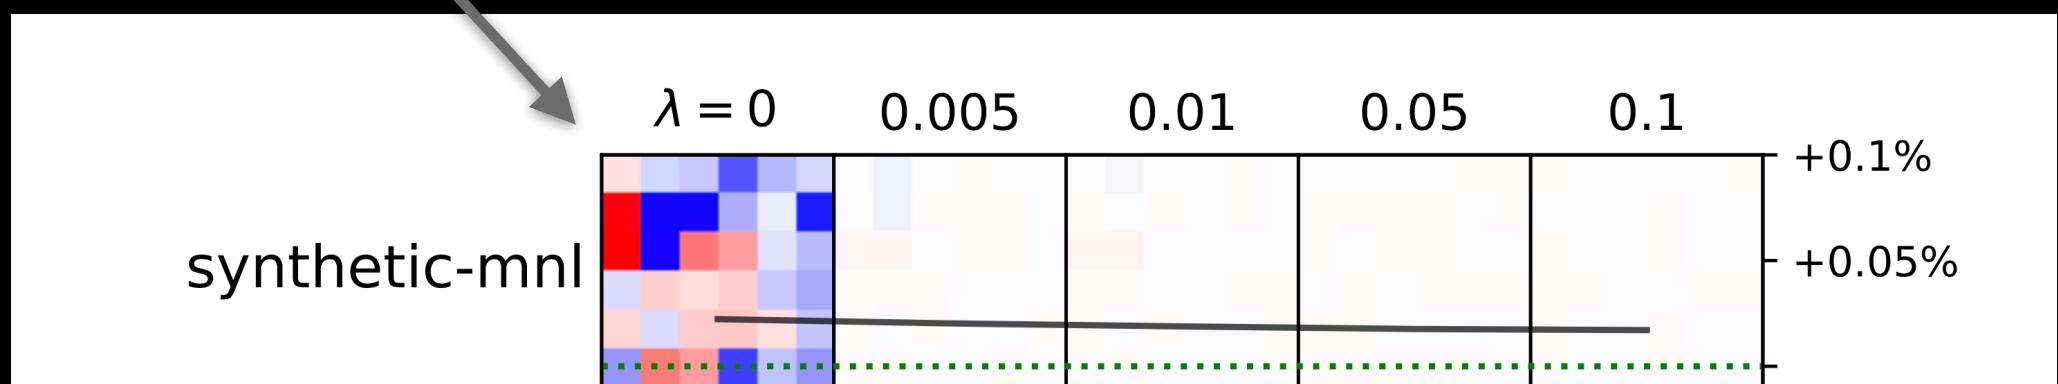


LCL reveals interpretable feature context effects

context effect matrix A

red: +, blue: -, white: 0

(column acts on row)



Node features
(left-right, top-bottom)

1. in-degree
2. shared neighbors
3. reciprocal weight
4. send recency
5. receive recency
6. reciprocal recency

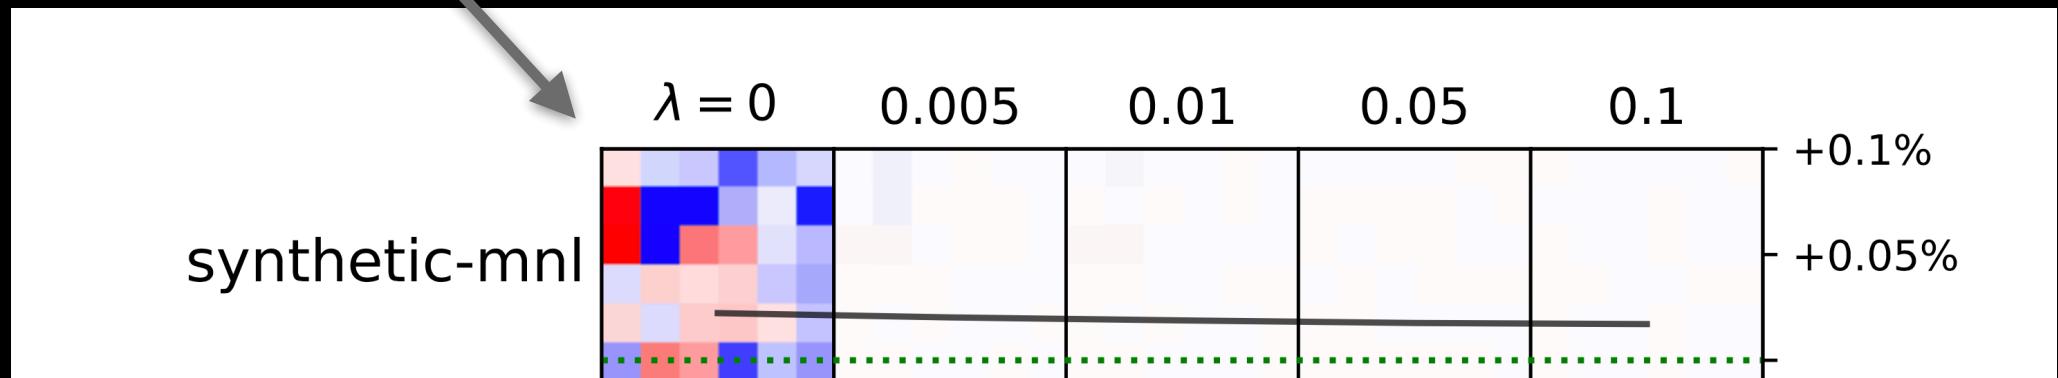
LCL reveals interpretable feature context effects

context effect matrix A

red: +, blue: -, white: 0

(column acts on row)

increasing L_1 regularization on A



Node features
(left-right, top-bottom)

1. in-degree
2. shared neighbors
3. reciprocal weight
4. send recency
5. receive recency
6. reciprocal recency

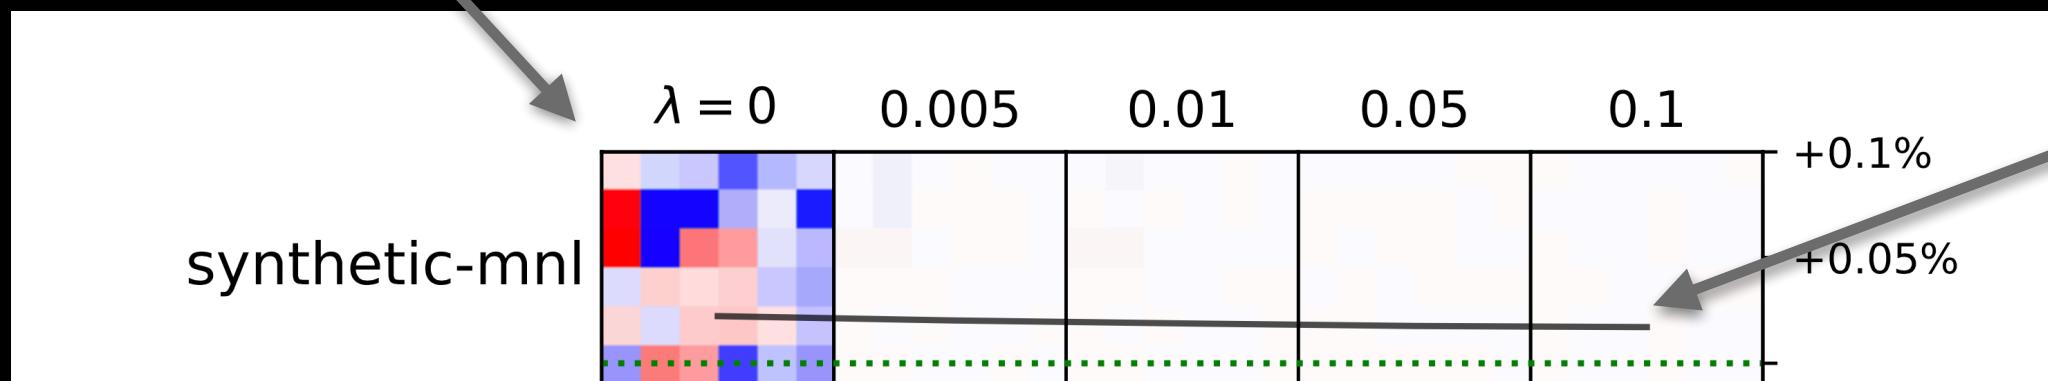
LCL reveals interpretable feature context effects

context effect matrix A

red: +, blue: -, white: 0

(column acts on row)

increasing L_1 regularization on A



NLL (lower = better fit)

Node features
(left-right, top-bottom)

1. in-degree
2. shared neighbors
3. reciprocal weight
4. send recency
5. receive recency
6. reciprocal recency

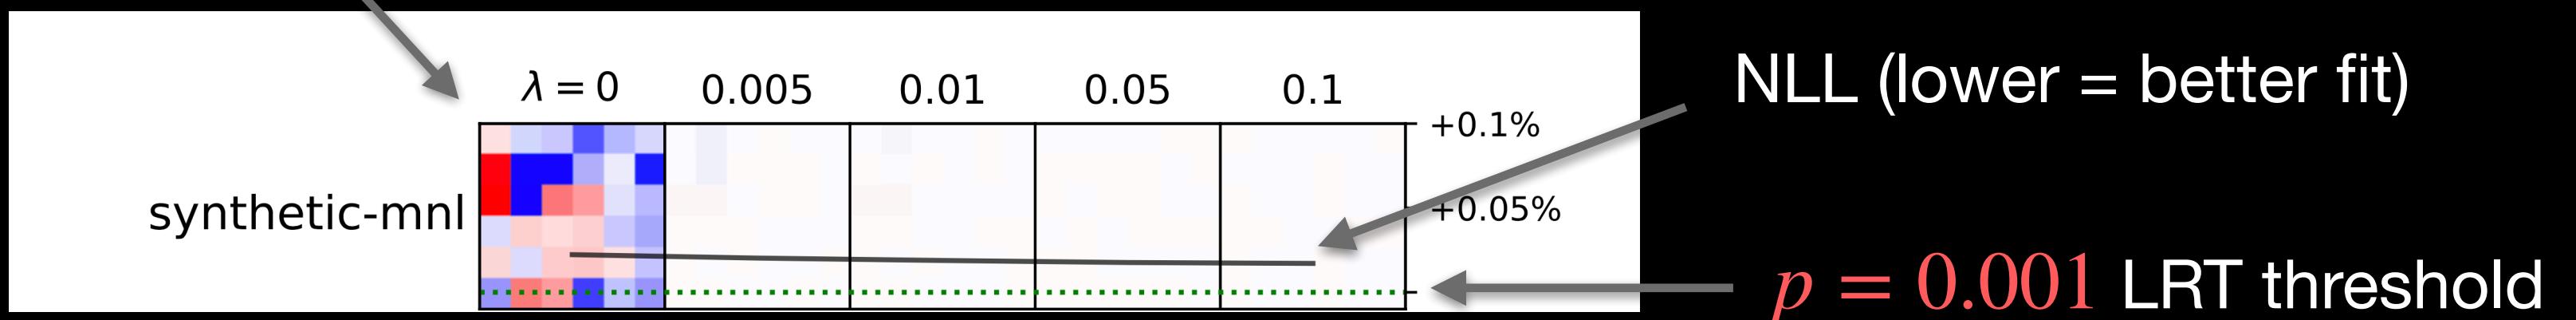
LCL reveals interpretable feature context effects

context effect matrix A

red: +, blue: -, white: 0

(column acts on row)

increasing L_1 regularization on A



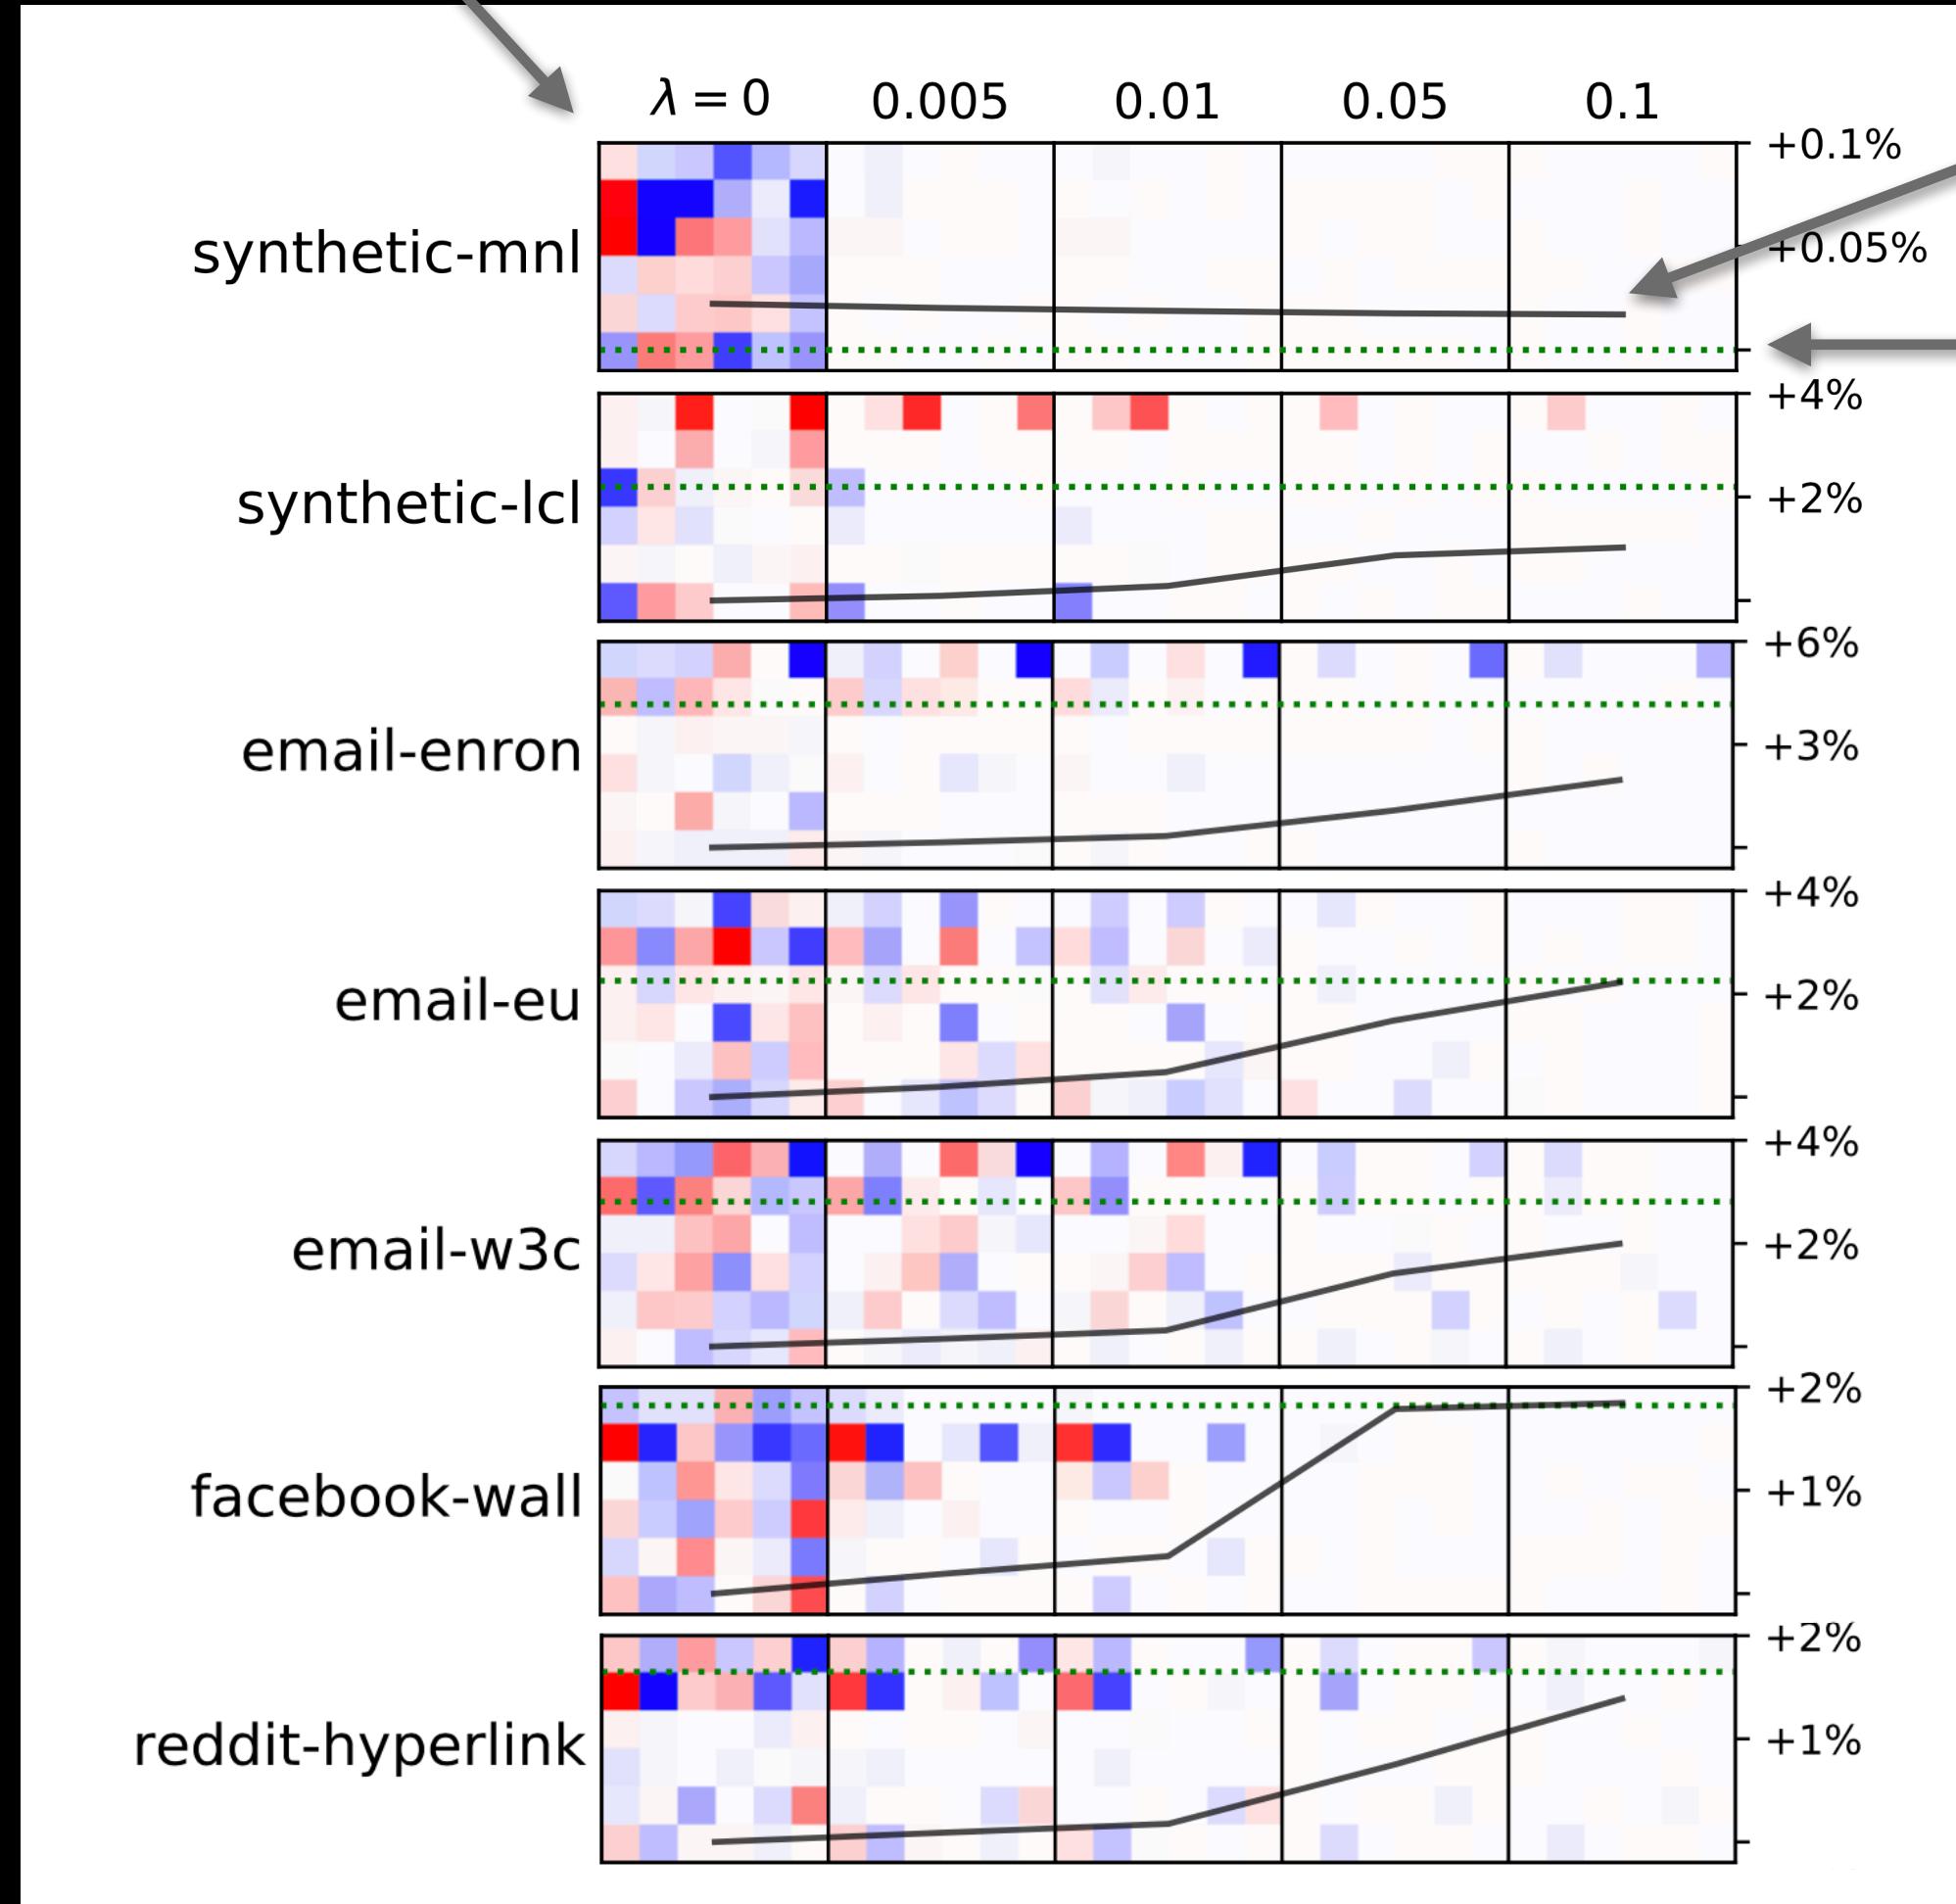
Node features
(left-right, top-bottom)

1. in-degree
2. shared neighbors
3. reciprocal weight
4. send recency
5. receive recency
6. reciprocal recency

LCL reveals interpretable feature context effects

context effect matrix A
red: +, blue: -, white: 0
(column acts on row)

increasing L_1 regularization on A



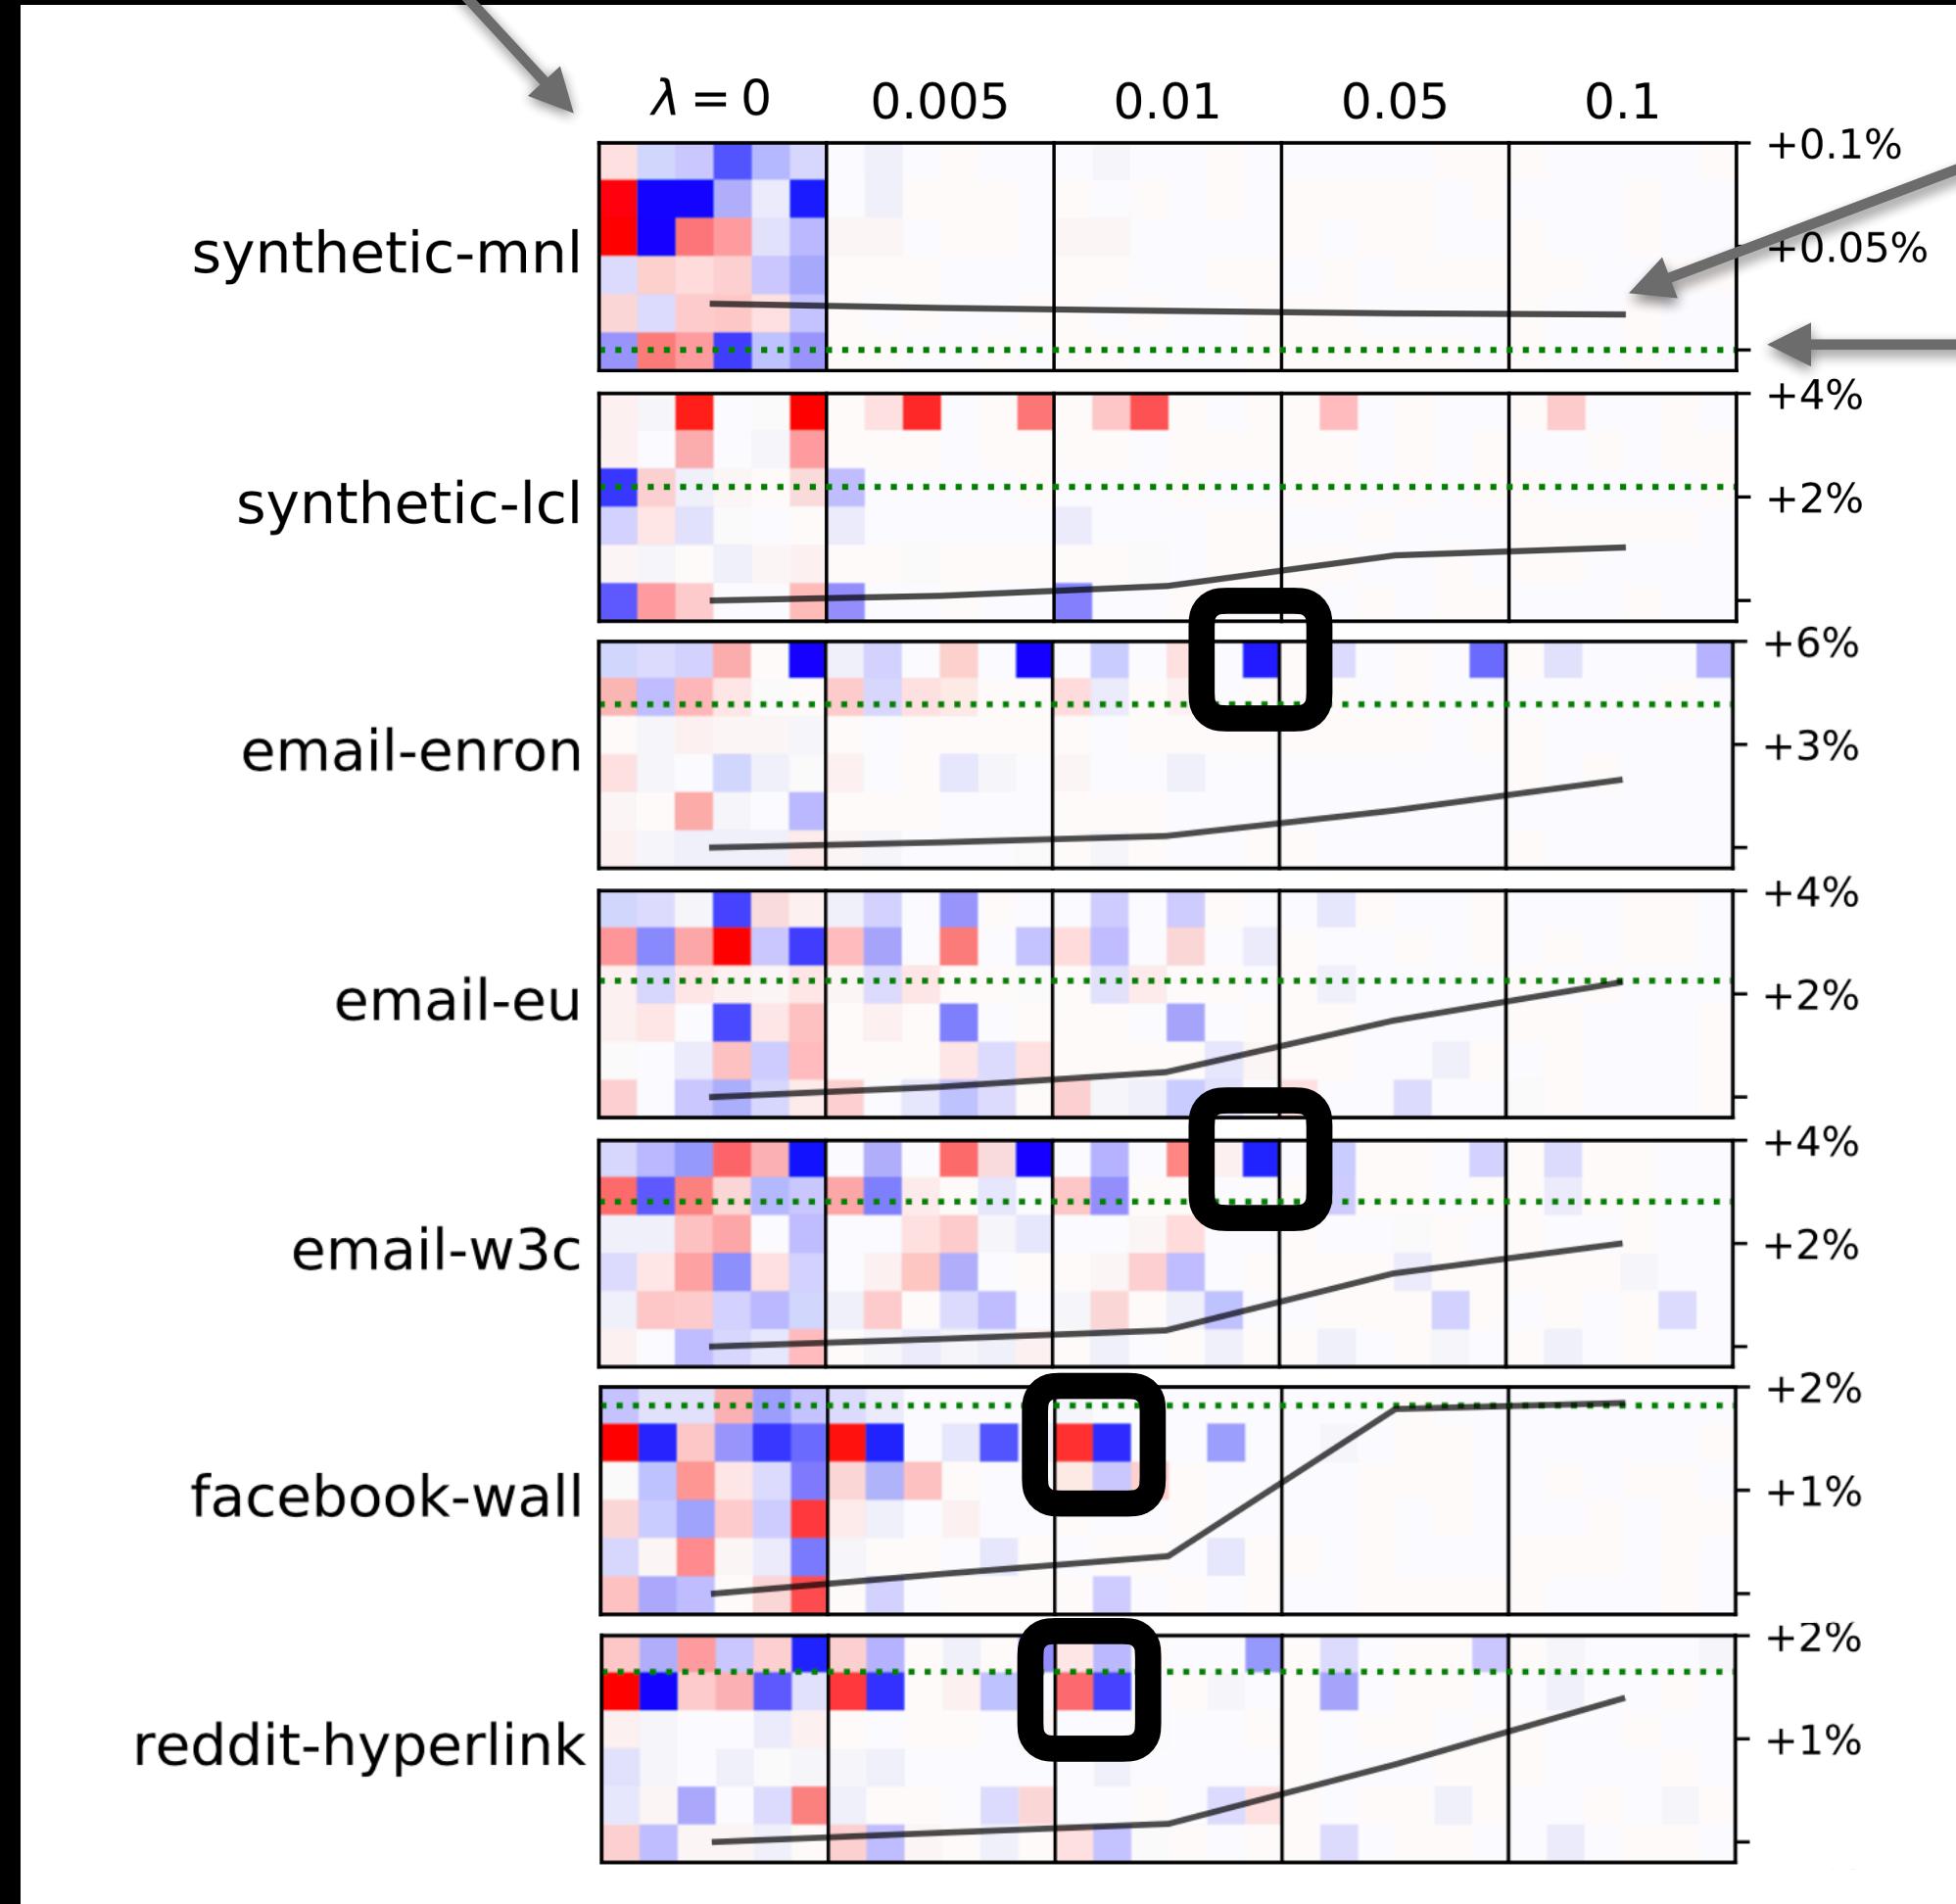
Node features
(left-right, top-bottom)

1. in-degree
2. shared neighbors
3. reciprocal weight
4. send recency
5. receive recency
6. reciprocal recency

LCL reveals interpretable feature context effects

context effect matrix A
red: +, blue: -, white: 0
(column acts on row)

increasing L_1 regularization on A



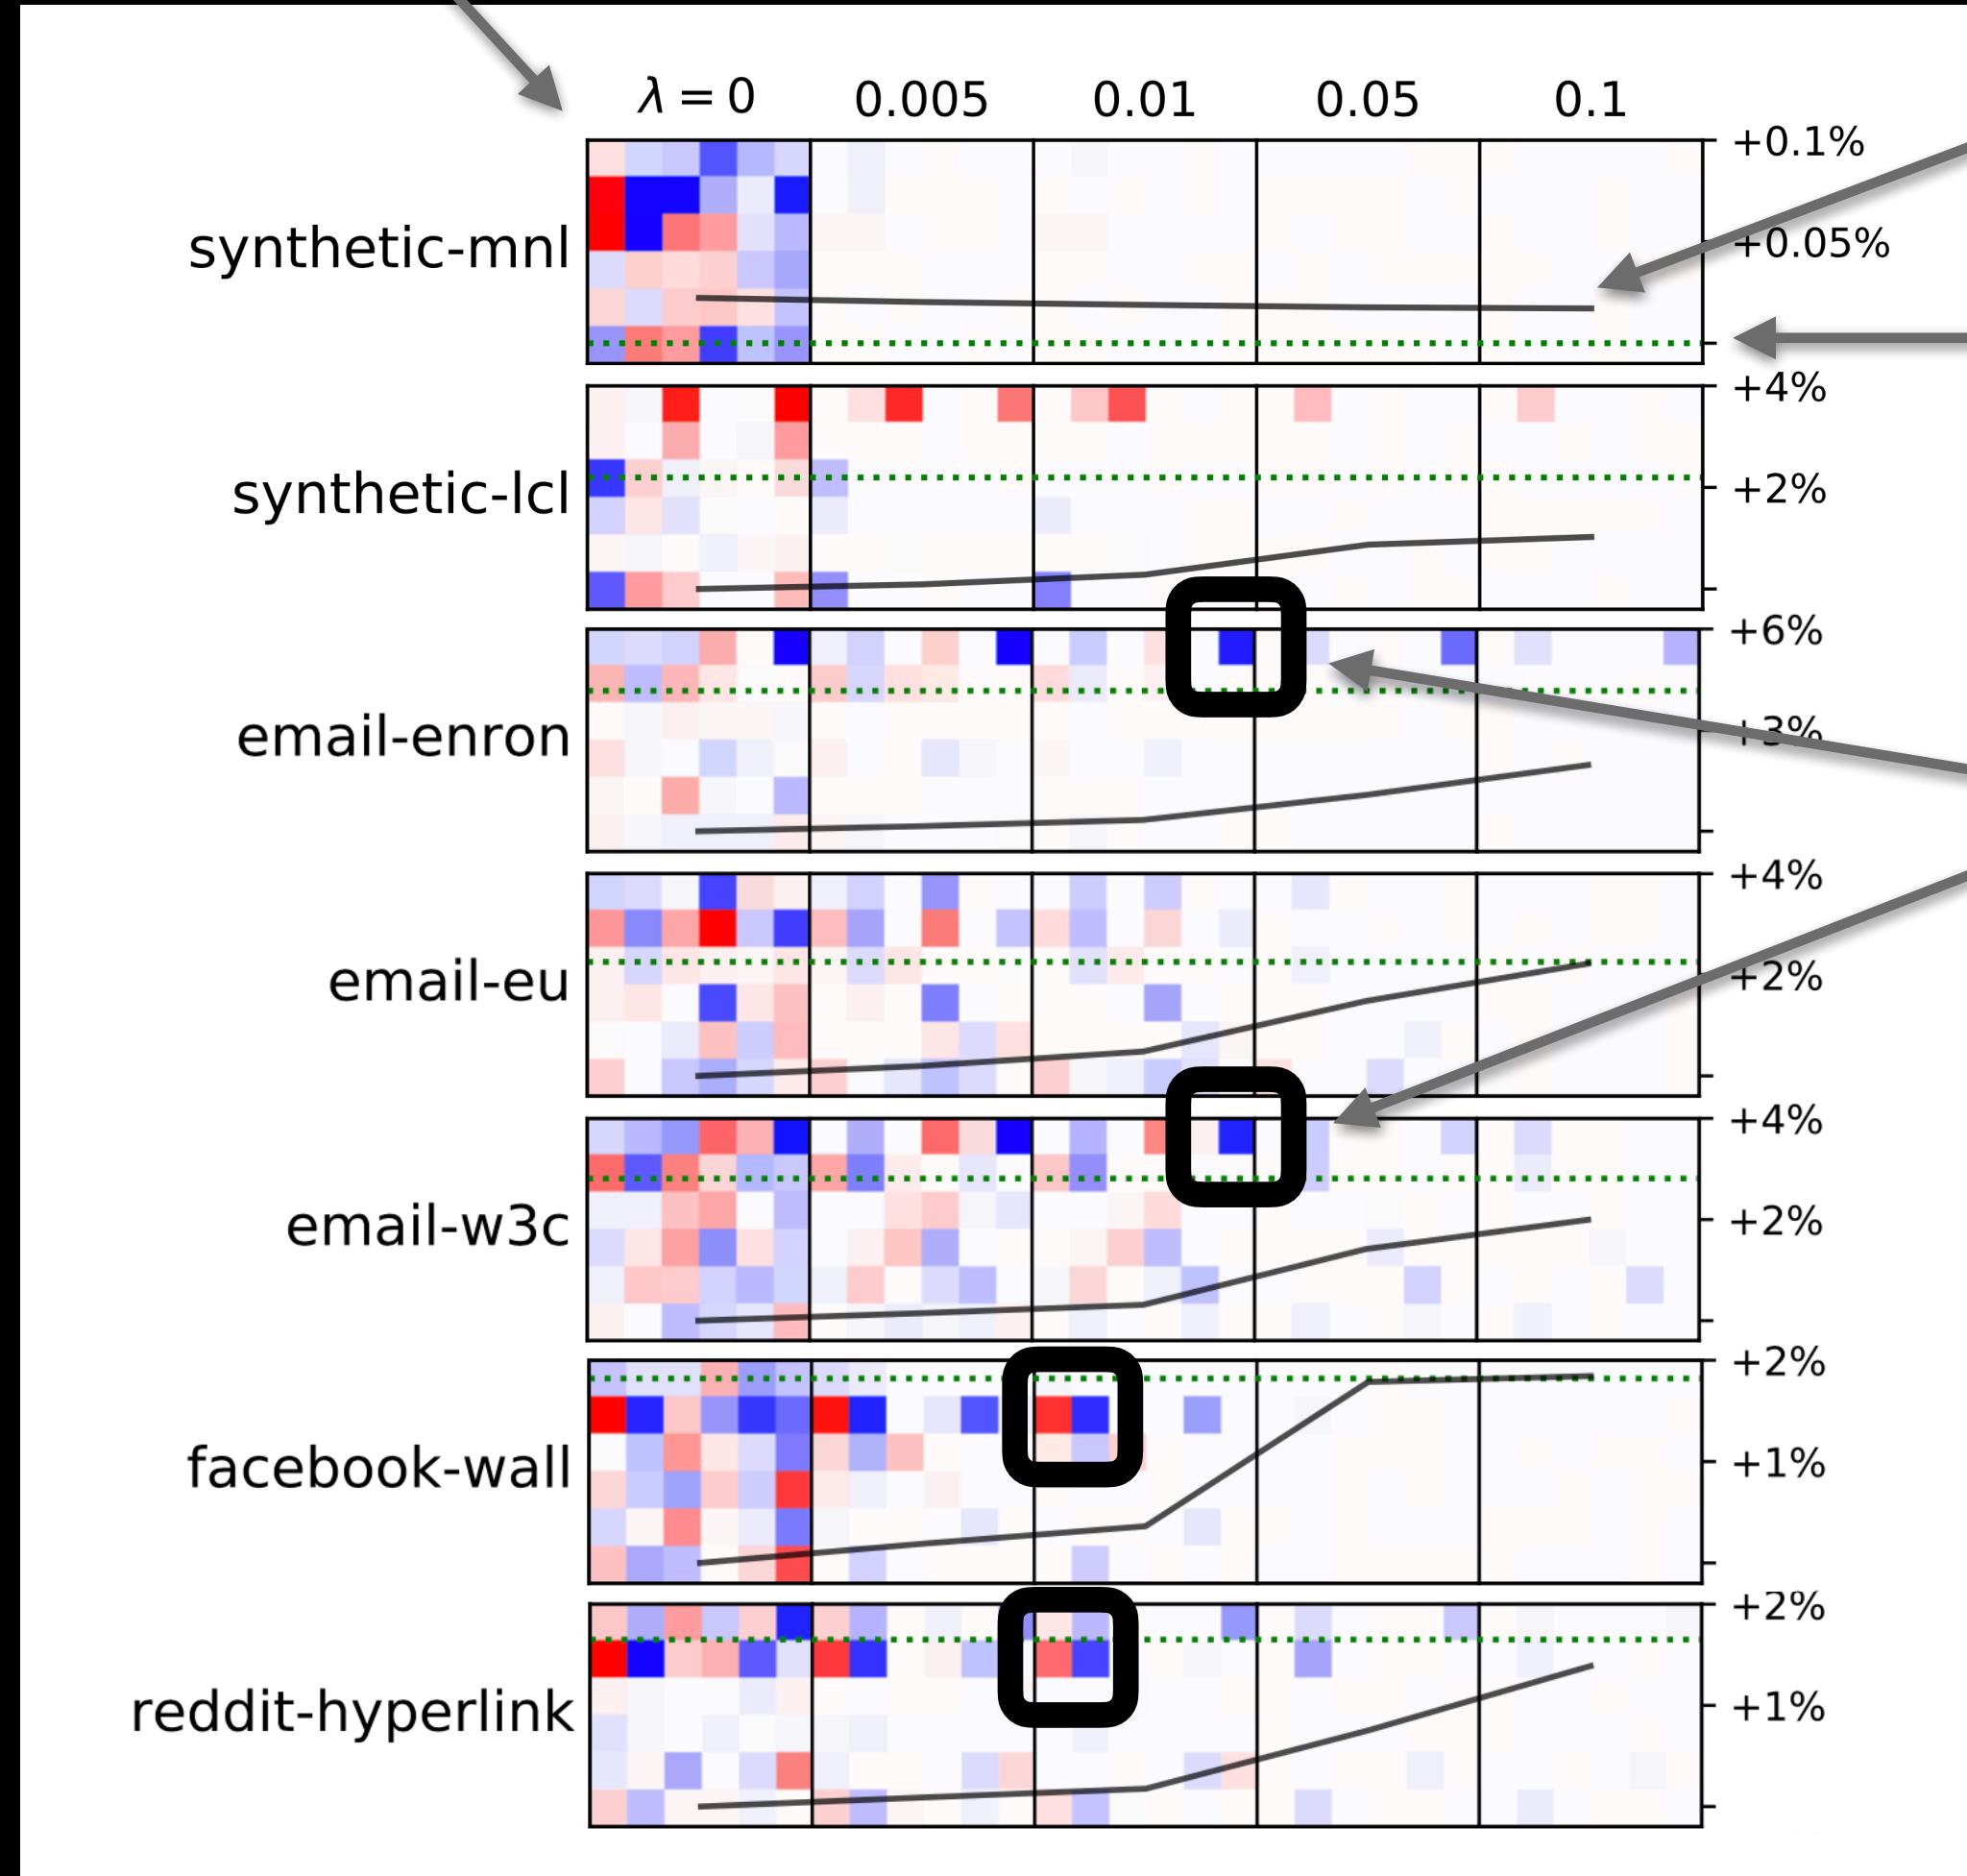
Node features
(left-right, top-bottom)

1. in-degree
2. shared neighbors
3. reciprocal weight
4. send recency
5. receive recency
6. reciprocal recency

LCL reveals interpretable feature context effects

context effect matrix A
red: +, blue: -, white: 0
(column acts on row)

increasing L_1 regularization on A



Node features (left-right, top-bottom)

1. in-degree
2. shared neighbors
3. reciprocal weight
4. send recency
5. receive recency
6. reciprocal recency

NLL (lower = better fit)

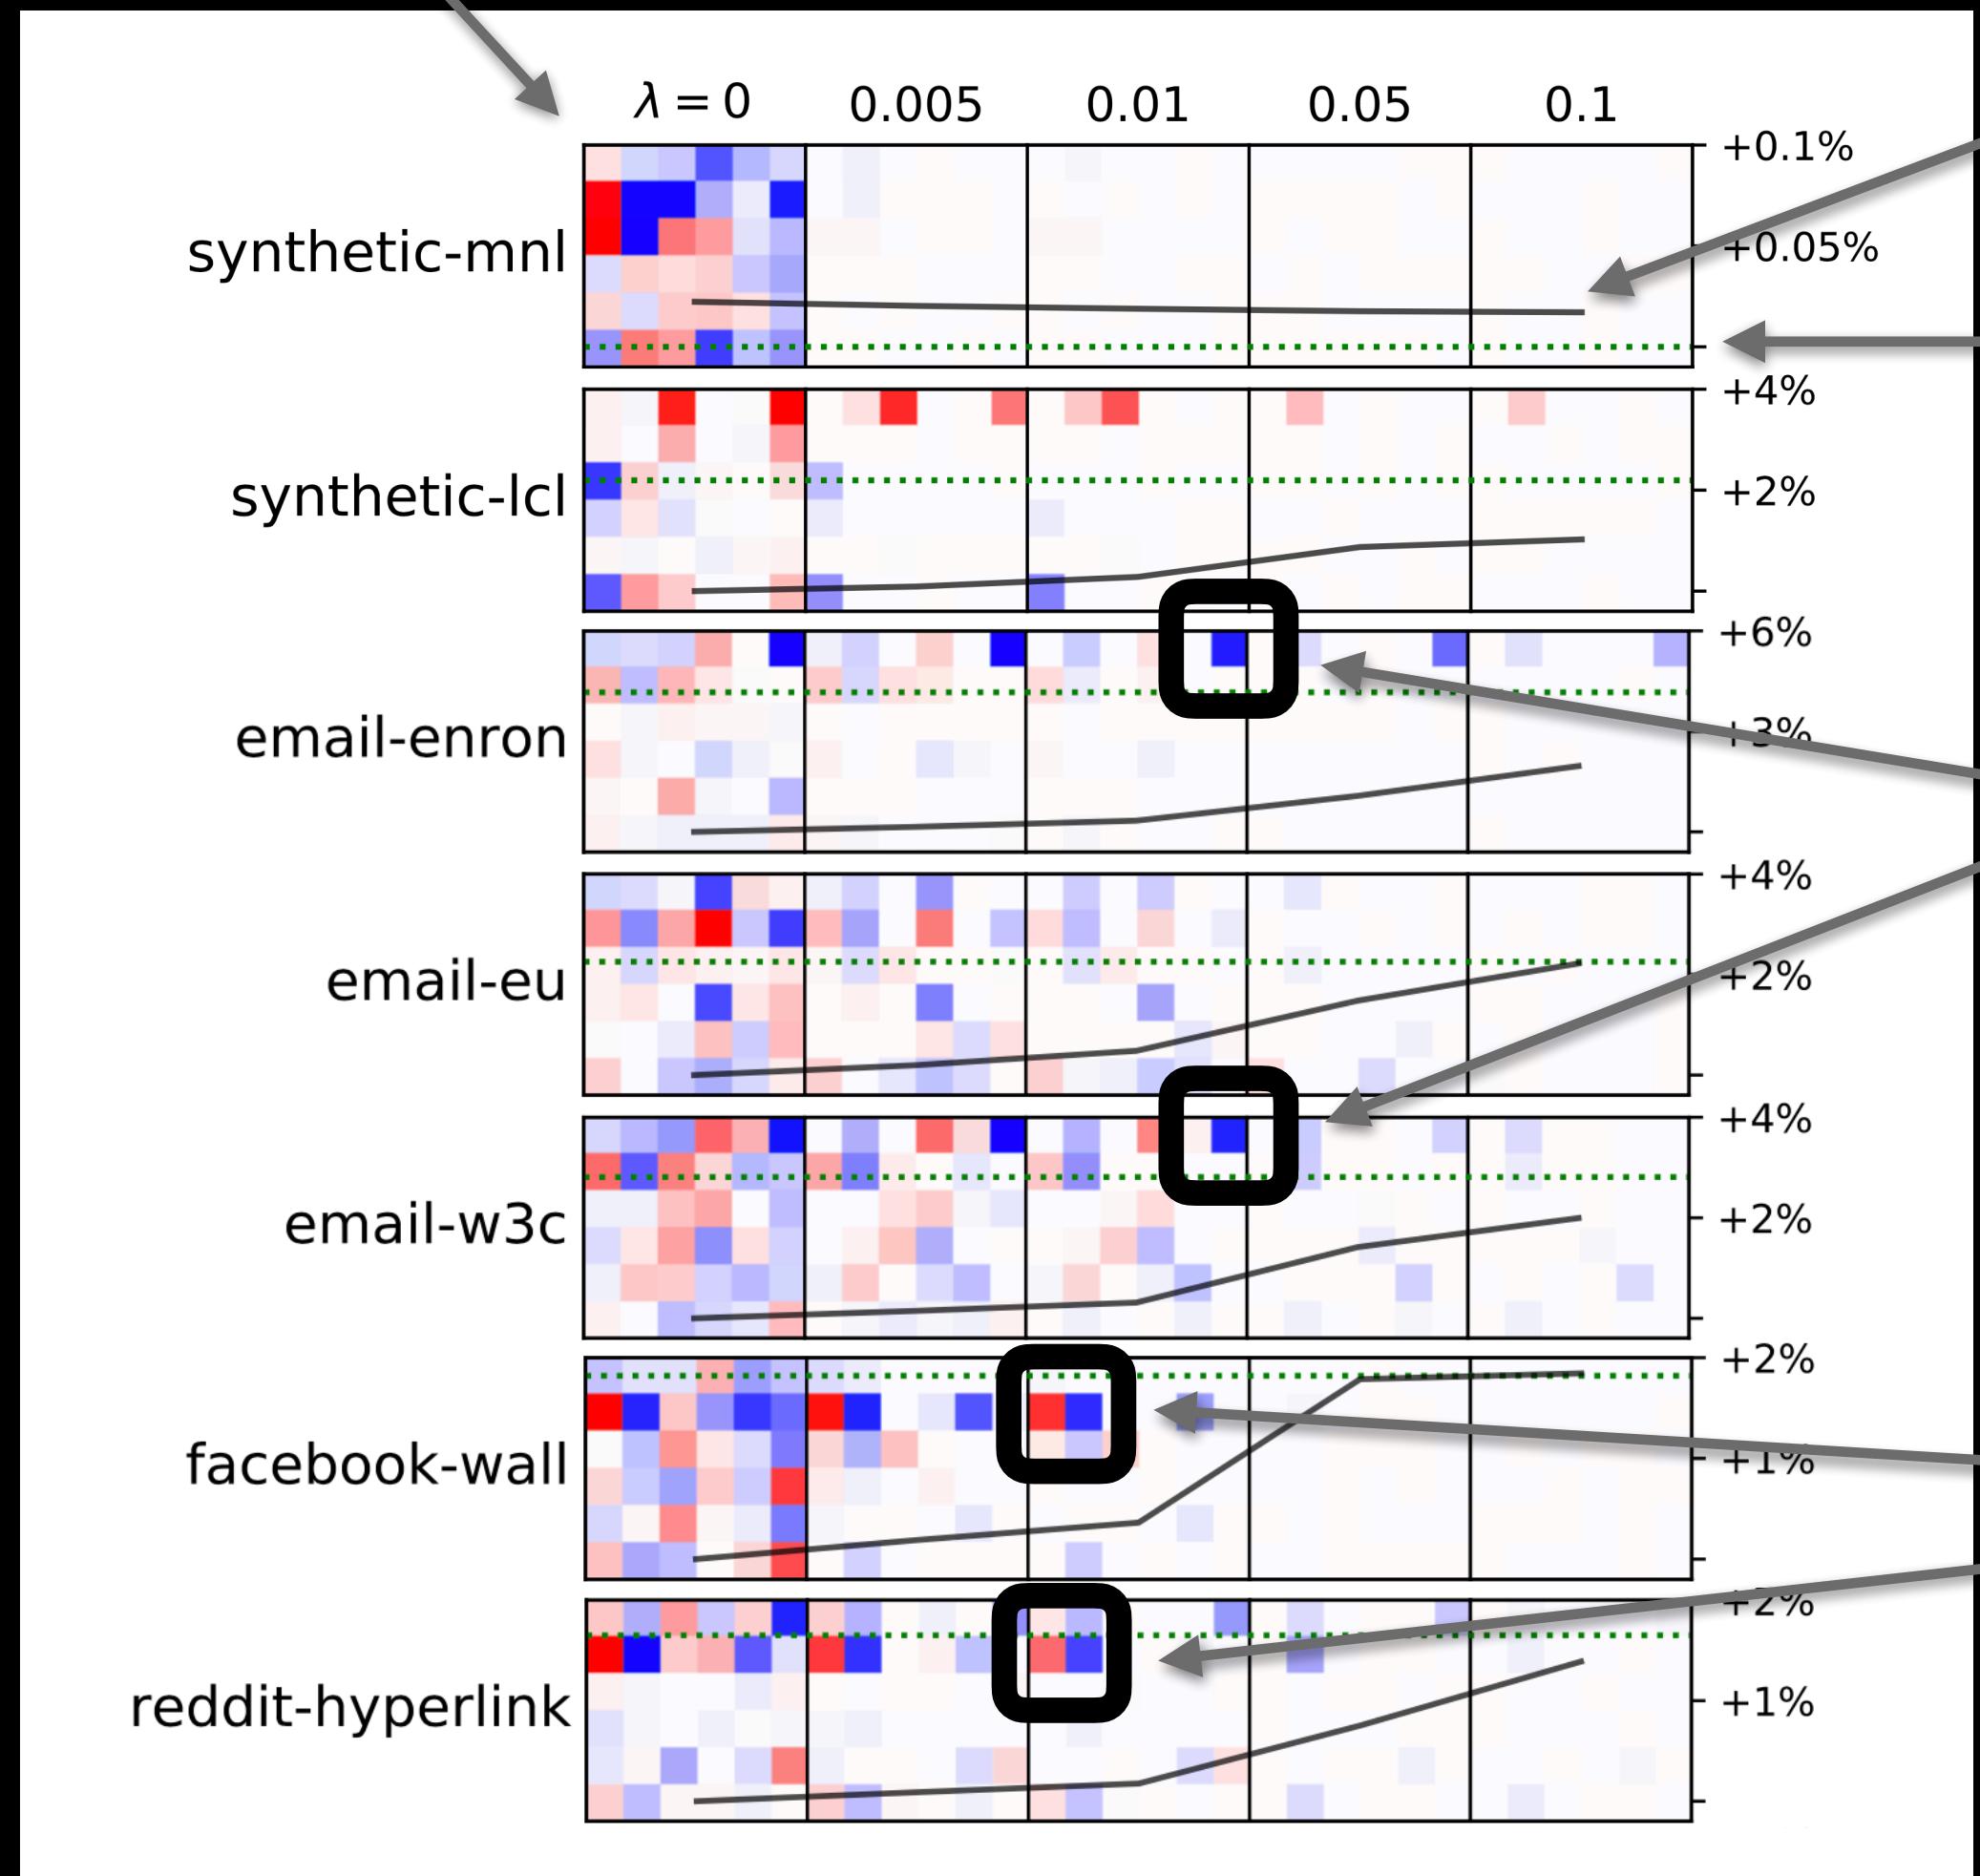
$p = 0.001$ LRT threshold

“cluttered inbox”
high choice set reciprocal recency
→ in-degree less important

LCL reveals interpretable feature context effects

context effect matrix A
 red: +, blue: -, white: 0
 (column acts on row)

increasing L_1 regularization on A



Node features (left-right, top-bottom)

1. in-degree
2. shared neighbors
3. reciprocal weight
4. send recency
5. receive recency
6. reciprocal recency

NLL (lower = better fit)

$p = 0.001$ LRT threshold

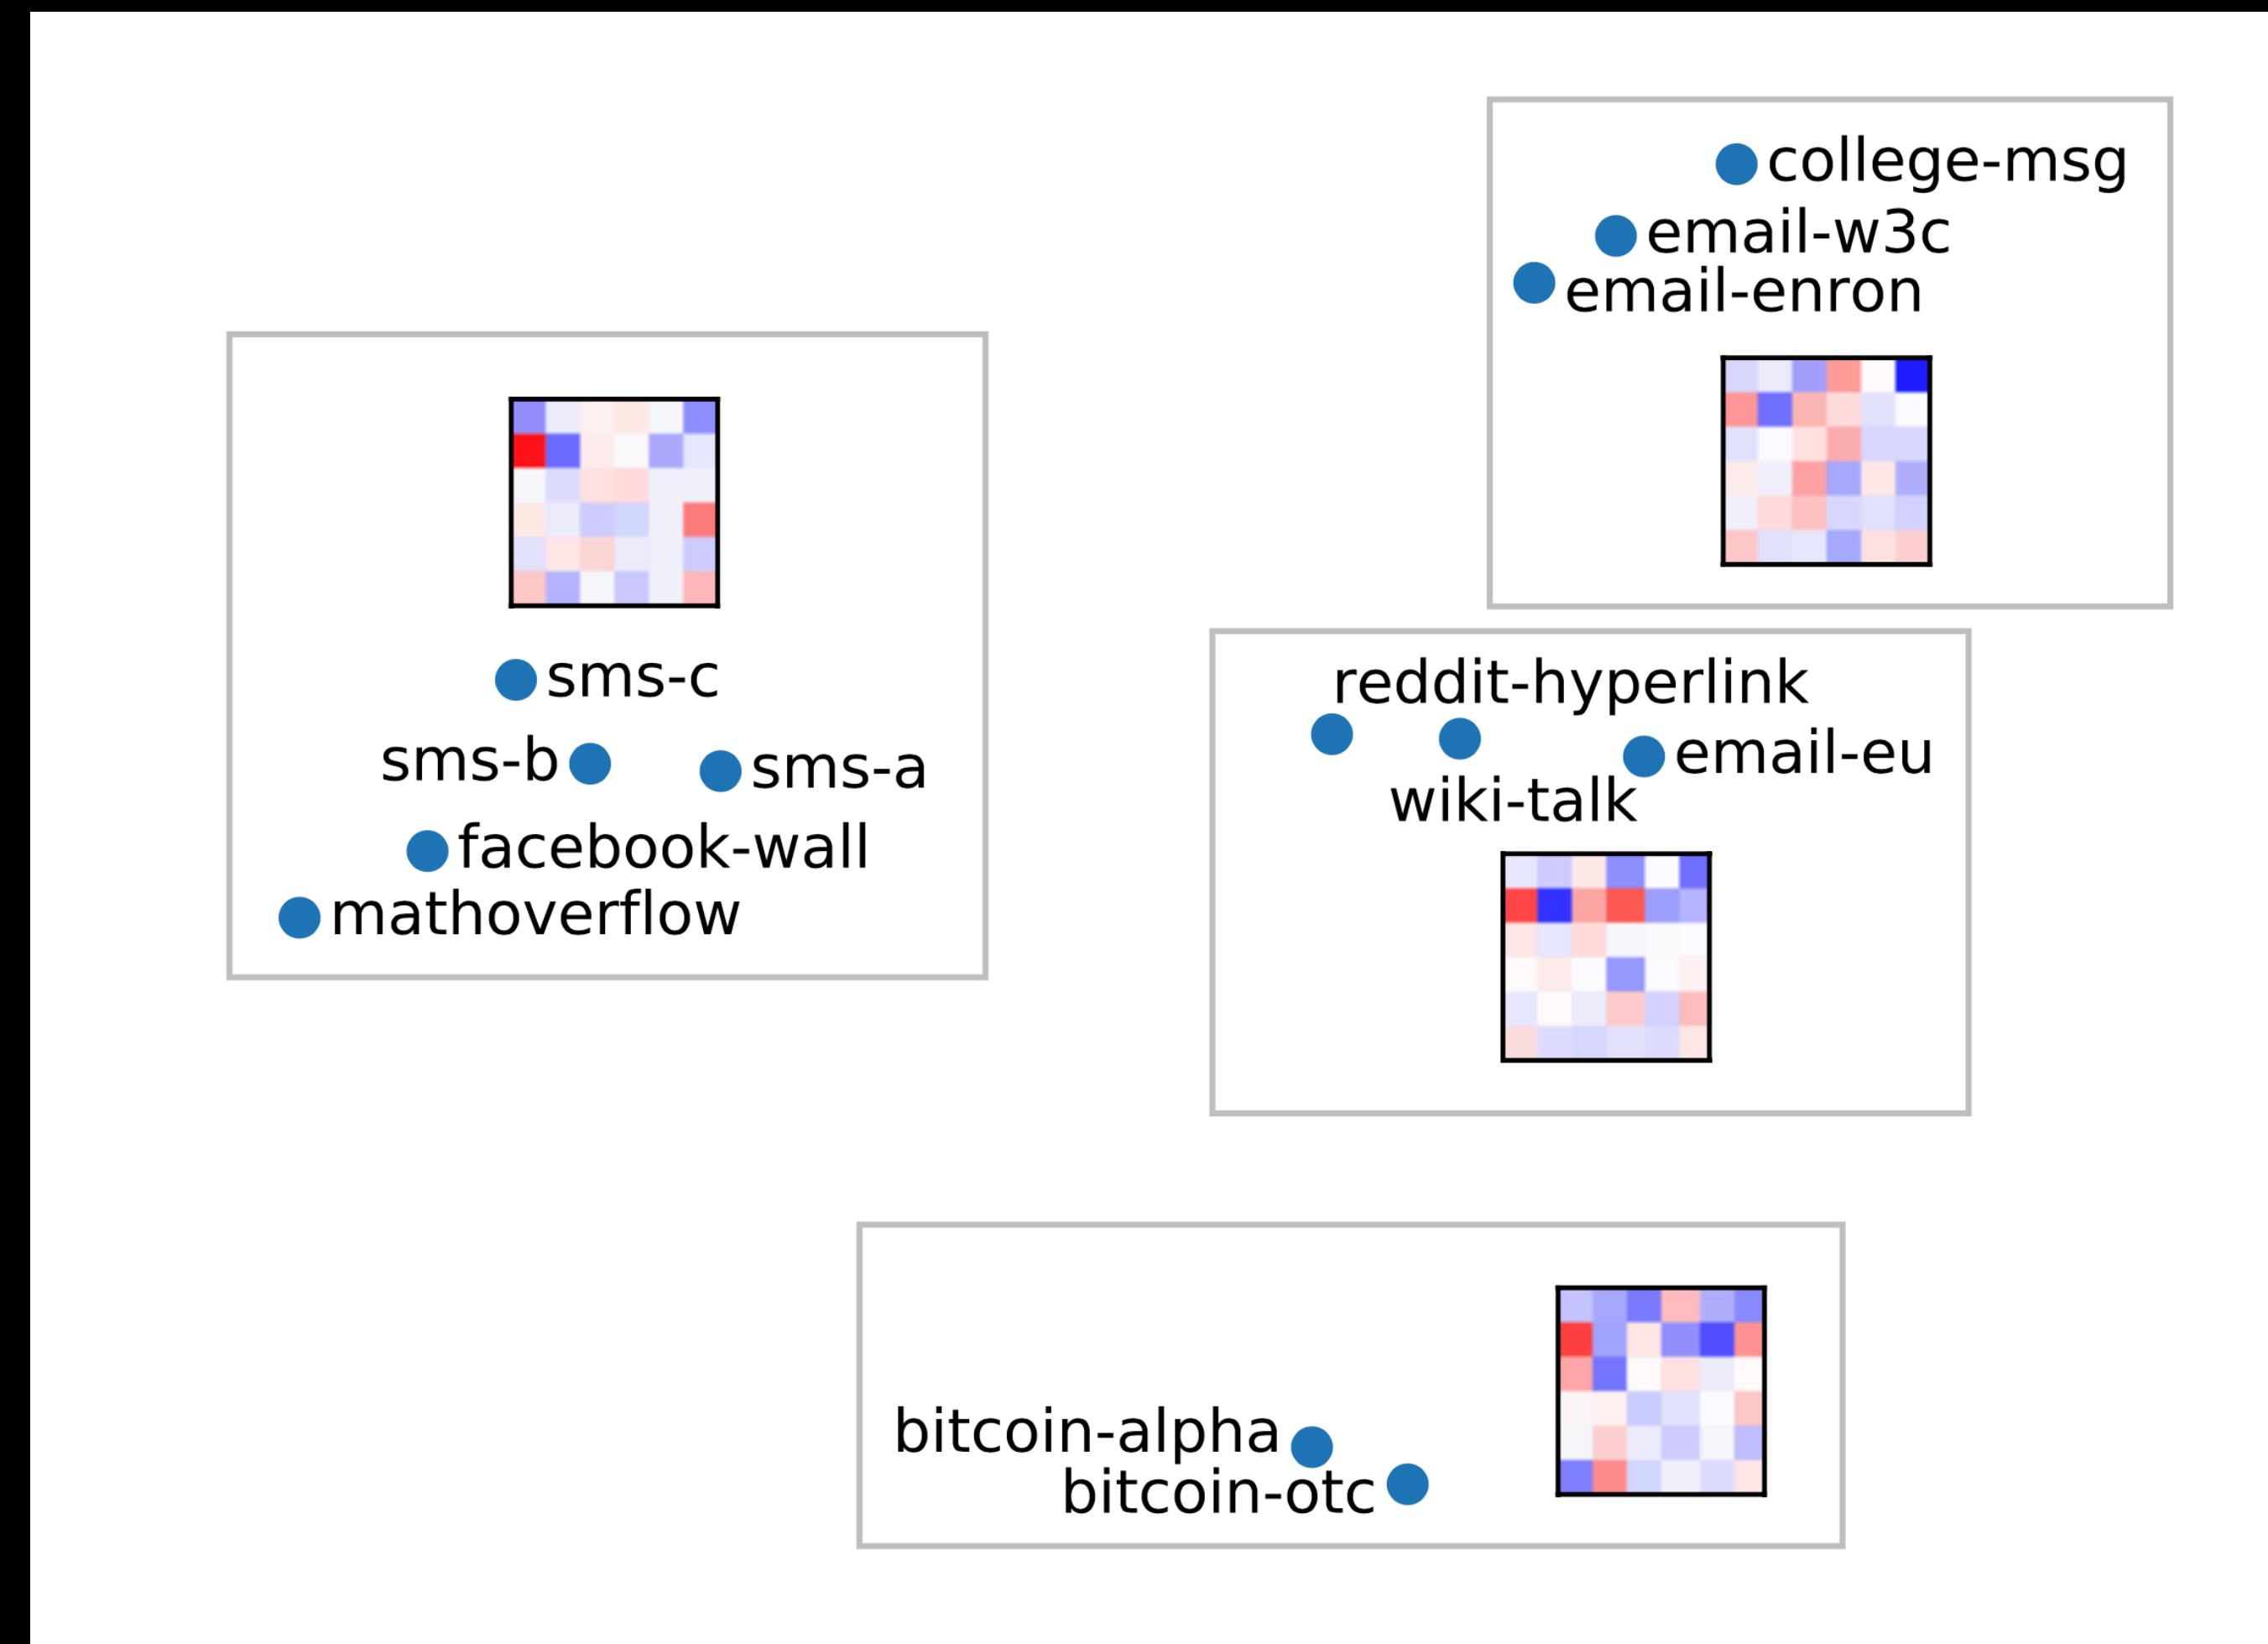
“cluttered inbox”
 high choice set reciprocal recency
 → in-degree less important

red: *“cocktail party introduction”*
 high choice set in-degree
 → shared neighbors more important

blue: *“familiarity saturation”*
 high choice set shared neighbors
 → shared neighbors less important

Similar datasets have similar feature context effects

T-SNE embedding of learned A matrices



Concluding thoughts

Concluding thoughts

Key takeaways

Feature context effects extend item-level effects

LCL offers an interpretable and tractable way to reveal them

Concluding thoughts

Key takeaways

Feature context effects extend item-level effects

LCL offers an interpretable and tractable way to reveal them

Challenges

Features correlate

Causal context effects?

Handling nonlinearity?

Concluding thoughts

Slides: bit.ly/lcl-slides-2
Preprint: bit.ly/lcl-paper
Code: bit.ly/lcl-code
Data: bit.ly/lcl-data

Key takeaways

Feature context effects extend item-level effects

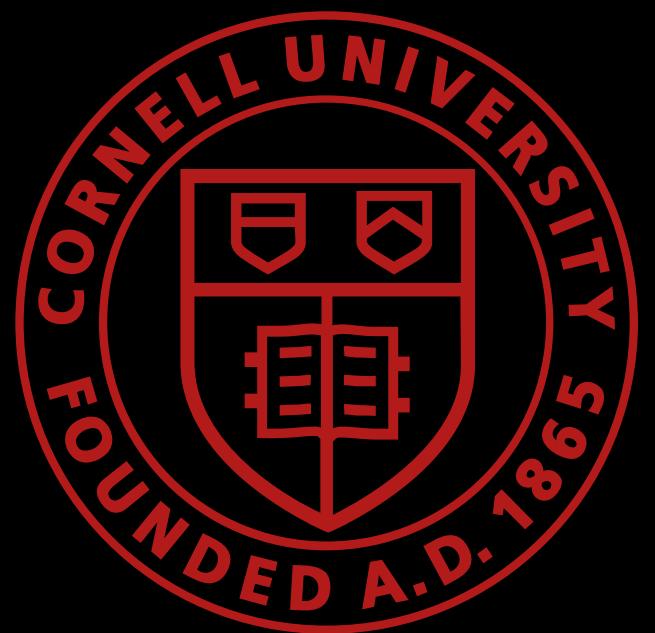
LCL offers an interpretable and tractable way to reveal them

Challenges

Features correlate

Causal context effects?

Handling nonlinearity?



Thank you!

More questions or ideas?

Email me: kt@cs.cornell.edu



@kiran_tomlinson

Acknowledgments

Funding from NSF, ARO

Thanks to Johan Ugander
and Jan Overgoor